



Universidade do Minho
Escola de Engenharia

Luís Carlos Sousa Magalhães

SDN support on V2X devices

Luís Magalhães **SDN support on V2X devices**



Universidade do Minho
Escola de Engenharia

Luís Carlos Sousa Magalhães

SDN support on V2X devices

Master Thesis
Master in Informatics Engineering

Work developed under the supervision of:
António Costa
Maria João Nicolau

COPYRIGHT AND TERMS OF USE OF THIS WORK BY A THIRD PARTY

This is academic work that can be used by third parties as long as internationally accepted rules and good practices regarding copyright and related rights are respected.

Accordingly, this work may be used under the license provided below.

If the user needs permission to make use of the work under conditions not provided for in the indicated licensing, they should contact the author through the RepositoriUM of Universidade do Minho.

License granted to the users of this work



Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International

CC BY-NC-SA 4.0

<https://creativecommons.org/licenses/by-nc-sa/4.0/deed.en>

STATEMENT OF INTEGRITY

I hereby declare having conducted this academic work with integrity. I confirm that I have not used plagiarism or any form of undue use of information or falsification of results along the process leading to its elaboration.

I further declare that I have fully acknowledged the Code of Ethical Conduct of the Universidade do Minho.

(Place)

(Date)

(Luís Carlos Sousa Magalhães)

”

*“- What do you do when you can't do nothing, but
there is nothing you can do? - You do what you can.”*

— **Grandad and Huey**, The Boondocks
(en)

Resumo

Suporte de SDN em dispositivos V2X

Nas últimas três décadas, a procura de integração da conectividade à Internet no domínio das redes veiculares, normalmente designadas por [VANET](#), tem vindo a aumentar. [SDN](#) representa um paradigma de rede que promete numerosos benefícios, e cuja concretização pode potencialmente melhorar a Internet. Uma abordagem promissora consiste na aplicação de [SDN](#) às redes veiculares, designada por [SDVN](#). No entanto, um dos principais obstáculos à sua utilização é a falta de plataformas experimentais que utilizem hardware real. Este problema representa um desafio significativo para a investigação em [SDVN](#), uma vez que é altamente complexo e exige uma compreensão abrangente dos domínios [SDN](#) e [VANET](#). Além disso, a integração da [P4](#) e da segunda geração de [SDN](#) com as [VANETs](#), é ainda área pouco explorada, abrindo novas oportunidades de investigação. Dada a natureza inovadora deste desafio e a complexidade destes dois domínios, a sua resolução constitui um objetivo altamente interessante e ambicioso.

Esta dissertação visa contribuir para este domínio de estudo, apresentando o desenvolvimento e a implementação de um [OBU](#)/sub-sistema [ITS](#) veicular que cumpra com as especificações de [SDN](#). O protótipo proposto baseia-se em hardware acessível e facilmente disponível e utiliza software de código aberto sempre que possível. Este é constituído por um sistema embebido PC Engines apu3d4, com placa de comunicação sem fios capaz de operar na banda de 5.9 GHz de acordo com a norma [IEEE 802.11p](#) para comunicações [V2X](#), e integrando um switch [P4 BMv2](#) para suporte deste. O sistema foi testado com sucesso num conjunto de cenários experimentais previamente definidos, cumprindo os requisitos estabelecidos para uso como plataforma de testes em ambientes reais. Esta solução estabelece as bases necessárias para o desenvolvimento de um ambiente [SDVN](#) completo.

Palavras-chave: Redes veiculares, Redes definidas por software, [Intelligent Transport System](#), Redes veiculares definidas por software, [On Board Unit](#), Sub-sistema [ITS](#) veicular, [IEEE 802.11p](#), [Programming Protocol-independent Packet Processors](#)

Abstract

SDN support on V2X devices

Over the past three decades, there has been a surge in demand for the integration of Internet connectivity within the domain of vehicular networks, commonly referred to as [VANETs](#). [SDN](#) represents a networking paradigm that promises numerous benefits, the realization of which could potentially enhance the Internet. A promising approach that has emerged involves the application of [SDN](#) to vehicular networks, a concept referred to as [SDVN](#). However, one of the main obstacles to its use is the lack of experimental platforms using real hardware. This problem represents a significant challenge for [SDVN](#) research, as it is highly intricate and necessitates a comprehensive grasp of both [SDN](#) and [VANET](#) domains. Furthermore, the integration of [P4](#) and the second generation of [SDN](#) with [VANETs](#) remains an under-explored field, thereby creating new research opportunities. Given the innovative nature of this challenge and the complexity of these two domains, addressing it is a highly interesting and ambitious goal.

This dissertation seeks to make a meaningful contribution to the field by presenting the development and implementation of a vehicle [OBU](#)/Vehicle [ITS](#) sub-system that complies with the specifications of [SDN](#). The proposed prototype is based on affordable and readily available hardware and utilizes open-source software wherever feasible. This system features an embedded PC Engines apu3d4 system, with a wireless communication card capable of operating in the 5.9 GHz band according to the [IEEE 802.11p](#) standard for [V2X](#) communications. It also features a [BMv2](#) switch, which enables support for [P4](#). The system successfully underwent testing in a series of previously defined experimental scenarios, thereby meeting the requirements established for its use as a test platform in real-world environments. Additionally, this paper provides a new perspective on the future of the [SDVN](#) field, based on a comprehensive global analysis. These findings present a potential trajectory for the evolution of [SDVN](#) technology. This paper presents the development of a prototype comprising a PC Engines apu3d4 that operates the [BMv2 P4](#) switch, communicating at 5.9GHz frequencies via the 802.11p standard.

This solution establishes the necessary framework for the assembly of a comprehensive [SDVN](#) environment.

Keywords: [Vehicular ad hoc Network](#), [Software Defined Networking](#), [Intelligent Transport System](#), [Software Defined Vehicular Network](#), [On Board Unit](#), [Vehicle ITS subsystem](#), [IEEE 802.11p](#), [Programming Protocol-independent Packet Processors](#)

Contents

List of Figures	x
List of Tables	xiii
Acronyms	xiv
1 Introduction	1
1.1 Context	1
1.2 Motivation	2
1.3 Objectives	2
1.4 Methodology	3
1.5 Main contributions	3
1.6 Document structure	4
2 Vehicular Ad hoc Networks	5
2.1 VANET characteristics	6
2.1.1 High mobility	6
2.1.2 Predicted mobility	7
2.1.3 Power and Computational ability	7
2.1.4 Congestion and Scalability issues	7
2.2 VANET components	8
2.2.1 C2C approach	8
2.2.2 ETSI approach	9
2.3 Communication Domains	12
2.3.1 C2C approach	12
2.3.2 ETSI approach	13
2.4 Communication categories	13
2.5 VANET architecture	15
2.5.1 Access layer	15

2.5.2	Networking & Transport layers	21
2.5.3	Facilities layer	22
2.5.4	Management Entity	24
2.5.5	Security Entity	25
2.6	Applications of VANET	27
2.7	Future trends and challenges in VANET research	28
3	Software Defined Networking	30
3.1	SDN definition	30
3.1.1	Main principles of SDN	31
3.1.2	Recent advancements	31
3.2	Motivation behind SDN	33
3.2.1	Network innovation	33
3.2.2	Network cost	34
3.2.3	Network control	35
3.3	SDN architecture	36
3.3.1	Data plane	36
3.3.2	Southbound API	44
3.3.3	Control plane	46
3.3.4	Northbound API	48
3.3.5	Management plane	49
4	Software Defined Vehicular Network	50
4.1	Benefits and Challenges	50
4.1.1	Mandated Interoperability	50
4.1.2	Network innovation	51
4.1.3	Network management and performance improvements	51
4.1.4	Issues with maintaining a global network view	52
4.1.5	Cheaper networks	53
4.1.6	Network security	53
4.2	SDVN architecture	53
4.2.1	Data plane	53
4.2.2	Control plane	53
4.2.3	Application plane	55
4.2.4	Communication interfaces	55
4.3	Related Work	55
4.4	Final considerations	58
4.4.1	The impact of interoperability	58
4.4.2	Controller placement	59

5 Device design and implementation	60
5.1 Architectural design	60
5.1.1 Data plane	60
5.1.2 Southbound API	64
5.1.3 Control plane	64
5.1.4 Northbound API and Application plane	64
5.2 Box implementation: version 1	65
5.3 Box implementation: version 2	67
5.4 P4/ONOS application implementation	68
5.5 Summary	70
6 Functional Tests	71
6.1 Experimental Setup	71
6.2 Experiment 1: testing the integration of OvS	72
6.2.1 Testing phase 1: standalone device with OvS and local controller	72
6.2.2 Testing phase 2: standalone device with OvS and remote controller	74
6.2.3 Testing phase 3: multiple devices with OvS and remote controller	76
6.2.4 Analysis and Interpretation	81
6.3 Experiment 2: testing the integration of BMv2 and P4	81
6.3.1 Testing phase 1: P4 dataplane on devices	81
6.3.2 Testing phase 2: P4 control over wireless interfaces	85
6.3.3 Analysis and Interpretation	89
7 Conclusions and Future work	91
7.1 Conclusions	91
7.2 Future work	92
Bibliography	93
Annexes	
I Bundles of services	103

List of Figures

1	Comparison between infrastructure networks [5]	6
2	A RSU provides Internet connectivity to one OBU through another OBU [9]	9
3	Personal ITS station in a Personal ITS sub-system [10]	10
4	Central ITS station in a Central ITS sub-system [10]	11
5	Vehicle ITS station in a Vehicle ITS sub-system [10]	12
6	Roadside ITS station in a Roadside ITS sub-system [10]	13
7	Communication domains in VANET as categorized by the C2C-CC [3]	14
8	ITS station reference architecture [10]	15
9	Channels allocated in the US by FCC for ITS [15]	18
10	European ITS channel allocation [12]	19
11	Power spectral limits on each ITS channel [15]	19
12	CAM structure [25]	23
13	DENM structure [25]	23
14	ITSC management entity as part of the ITS station reference architecture [10]	24
15	ITSC security entity as part of the ITS station reference architecture [10]	26
16	ITS Security Certificate Management System [40]	27
17	ITS strategy for the EU [43]	29
18	Traditional network view Vs SDN Network view [50]	32
19	Layered view of SDN Architecture [56]	37
20	Evolution of SDN [51]	38
21	High-Level schematic of a bare-metal switch [47]	42
22	Definition of a logical pipeline to provide a pipeline-agnostic view the control plane [47]	42
23	Protocol-Independent Switch Architecture [57]	43
24	East/Westbound interface between distributed controllers [53]	48
25	Conceptual view of SDVN architecture [72]	54
26	Back and inside view from device version 1	66

27	Back and inside view from device version 2	69
28	Setup from the first phase of the experiments with OvS	73
29	Diagram that represents the setup from the first phase of the experiments with OvS	73
30	Diagram that depicts the SDN perspective of the setup from the first phase of the experiments with OvS	73
31	Results from running the commands "ip a"and "ovs-vsctl show"in phase 1	74
32	Results from running the commands "ip a"on both VRFs and "ping"between them in phase 1	75
33	Diagram that represents the setup from the second phase of the experiments with OvS	75
34	Diagram that depicts the SDN perspective of the setup from the second phase of the experiments with OvS	76
35	Results from running the commands "ovs-vsctl show"and "ping"between both VRFs in phase 2	77
36	ONOS console with information relating to the devices, hosts and flows of the network in phase 2	77
37	Setup from the third phase of the experiments with OvS	78
38	Diagram that represents the setup from the third phase of the experiments with OvS	78
39	Diagram that depicts the SDN perspective of the setup from the third phase of the experiments with OvS	79
40	Results from running the commands "ovs-vsctl show"in the devices running OvS in phase 3	79
41	ONOS console with information relating to the devices and hosts of the network in phase 3	80
42	ONOS console with information relating to the flows of the network in phase 3	80
43	Results from running the commands "ping"in the VRF between the two devices on the edge of the network in phase 3	80
44	Setup from the first phase of the experiments with BMv2	82
45	Diagram that represents the setup from the first phase of the experiments with BMv2	82
46	Diagram that depicts the SDN perspective of the setup from the first phase of the experiments with BMv2 in phase 1	83
47	Consoles running BMv2 on both devices	83
48	ONOS console with information relating to the devices and hosts of the network in phase 1	84
49	ONOS console with information relating to the applications activated on the controller in phase 1	84
50	ONOS console with information relating to the flows of the network in phase 1	84
51	Results from running the commands "ping"in the VRF between the two devices on the network in phase 1	85
52	Results from running the commands "ip a"and "ping"between the two devices on the network using the IP address of the wireless interfaces	86
53	Setup from the second phase of the experiments with BMv2	87

54	Diagram that represents the setup from the second phase of the experiments with BMv2	87
55	Diagram that depicts the SDN perspective of the setup from the second phase of the experiments with BMv2	87
56	Consoles running BMv2 on both devices in phase 2	88
57	ONOS console with information relating to the devices and hosts of the network in phase 2	88
58	ONOS console with information relating to the flows of the network in phase 2	88
59	Results from running the commands "ping" in the VRF of the two devices on the network in phase 2	89
60	Diagram representing the SDN network as an Vehicle ITS sub-system	90

List of Tables

1	Fields recognized by the OpenFlow standard [58]	39
2	Summary of related literature	57
3	C-ITS service bundles for scenario building[16]	104

Acronyms

3GPP	3rd Generation Partnership Project (pp. 19, 20)
5GAA	5G Automotive Association (p. 20)
API	Application Programming Interface (pp. 31, 36, 38–40, 43–49, 52, 55, 56, 64, 81, 89)
AU	Autonomous Unit (pp. 8, 10, 13)
BIOS	Basic Input/Output System (pp. 44, 62)
BMv2	Behavioral Model version 2 (pp. v, vi, xi, xii, 63, 67, 68, 71, 72, 81–83, 85–91)
BSS	Basic Service Set (pp. xv, 17)
BTP	Basic Transport Protocol (pp. 21, 22)
C-V2X	Cellular V2X (pp. 19, 20, 28, 55, 62)
C2C-CC	Car-2-Car Communication Consortium (pp. x, 6, 8–14)
CAM	Cooperative Awareness Message (pp. x, 22, 23, 57)
CEN	European Committee for Standardization (p. 22)
CPU	Central processing unit (pp. 41, 55, 61, 65, 67, 68)
DENM	Decentralized Environmental Notification Message (pp. x, 22, 23)
DRAM	Dynamic Random-access memory (pp. 55, 65, 67)
DRAM	Static Random-access memory (pp. 41, 57)
ECC	Electronic Communications Committee (p. 18)
ECU	Electronic Control Unit (pp. 10, 11)
ETSI	European Telecommunications Standards Institute (pp. 6, 8–10, 12, 13, 17, 18, 20–22, 24, 26, 28, 56–58, 62, 68, 70, 90, 92)
EU	European Union (pp. x, 6, 16–18, 25, 28, 29, 50, 51)

FCC	Federal Communications Commission (<i>pp. x, 17, 18</i>)
GN6	GeoNetworking-IPv6 (<i>pp. 21, 22</i>)
GPS	Global Positioning System (<i>pp. 5, 56, 57</i>)
IBSS	Independent Basic Service Set (<i>p. 17</i>)
IEEE	Institute of Electrical and Electronics Engineers (<i>pp. v, vi, 16–20, 28, 55–57, 62, 63, 68, 70, 71, 85, 86, 90–92</i>)
IP	Internet Protocol (<i>pp. xi, xv, 9, 28, 76, 78, 81, 82, 85, 86, 88</i>)
IPv6	Internet Protocol version 6 (<i>pp. xv, 22, 28</i>)
ISO	International Organization for Standardization (<i>p. 22</i>)
ITS	Intelligent Transport System (<i>pp. v, vi, x, xii, xv, 6, 9–13, 15–20, 22–29, 51, 53–59, 64, 65, 71, 90, 91</i>)
ITS-S	ITS station (<i>pp. 9, 10, 90</i>)
ITSC	Architecture of Communications in ITS (<i>pp. x, 13, 24, 26</i>)
LTE	Long-Term Evolution (<i>pp. 19, 20, 56, 57, 62, 65, 68</i>)
MANET	Mobile ad hoc Network (<i>pp. 5–7, 21</i>)
mPCIe	Mini PCI Express (<i>pp. 55–57, 65, 67, 68</i>)
mSATA	Mini-Serial Advanced Technology Attachment (<i>pp. 65, 67, 68</i>)
NIC	Network Interface Controller (<i>pp. xvii, 61</i>)
NOS	Network Operating System (<i>pp. 46, 48, 49</i>)
NPU	Network Processing Unit (<i>pp. 41–43</i>)
OBU	On Board Unit (<i>pp. v, vi, x, 2, 3, 8, 9, 11, 13, 56–58, 61, 65, 71, 91, 92</i>)
OCB	Out of Context BSS (<i>pp. 17, 19</i>)
ODL	Open Day Light (<i>pp. 48, 64</i>)
ONF	Open Networking Foundation (<i>pp. 30, 39, 45, 70</i>)
ONL	Open Network Linux (<i>p. 62</i>)
ONOS	Open Network Operating System (<i>pp. xi, xii, 48, 64, 67, 68, 70–72, 74, 76–78, 80–86, 88–91</i>)
OS	Operating System (<i>pp. xv, xvi, 44, 46, 56, 60, 62, 65, 68</i>)
OSI	Open System Interconnection (<i>pp. 9, 15, 21, 22, 53</i>)
OvS	Open virtual Switch (<i>pp. xi, 63, 67, 71–76, 78, 79, 81, 91</i>)
P4	Programming Protocol-independent Packet Processors (<i>pp. v, vi, 4, 32, 39–41, 43, 45, 51, 58, 63, 64, 68, 70, 72, 81, 83, 85, 89–92</i>)

PC	Personal Computer (<i>pp. 40, 41, 71, 72</i>)
PCI	Peripheral Component Interconnect (<i>pp. xv, 61</i>)
PDU	Packet Data Unit (<i>pp. 22, 23</i>)
PISA	Protocol Independent Switching Architecture (<i>p. 43</i>)
POX	Pythonic Network <i>Operating System</i> (<i>p. 48</i>)
PSA	Portable Switch Architecture (<i>p. 43</i>)
 QoS	Quality of Service (<i>pp. 8, 55</i>)
 RAM	Random-access memory (<i>pp. xiv, 61, 65, 67</i>)
RSU	Roadside Unit (<i>pp. x, 8, 9, 12, 13, 56, 57</i>)
 SATA	Serial Advanced Technology Attachment (<i>pp. xv, 56</i>)
SD	Secure Digital (<i>pp. 65, 67</i>)
SD-VANET	Software Defined <i>Vehicular ad hoc Network</i> (<i>p. 50</i>)
SDN	Software Defined Networking (<i>pp. v, vi, x–xii, 1–4, 28, 30–41, 44–64, 67, 70–74, 76, 78, 79, 82, 83, 86, 87, 90–92</i>)
SDR	Software Defined Radio (<i>pp. 63, 92</i>)
SDVN	Software Defined <i>Vehicular Network</i> (<i>pp. v, vi, x, 1–4, 50–55, 57–59, 64, 91</i>)
SIM	Subscriber Identity Module (<i>pp. 65, 68</i>)
SNMP	Simple Network Management Protocol (<i>p. 45</i>)
SONiC	Software for Open Networking in the Cloud (<i>p. 62</i>)
SSD	Solid State Drive (<i>pp. 56, 65, 67, 68</i>)
 TNA	Tofino Native Architecture (<i>p. 43</i>)
 UDP	User Datagram Protocol (<i>p. 22</i>)
US	United States of America (<i>pp. x, 6, 17, 18, 22</i>)
 V2I	Vehicle to Infrastructure communication (<i>pp. 14, 17</i>)
V2V	Vehicle to Vehicle communication (<i>pp. 14, 17</i>)
V2X	Vehicle to Everything communication (<i>pp. v, vi, xiv, 4, 14, 16, 20, 51</i>)
VANET	Vehicular ad hoc Network (<i>pp. v, vi, x, xvi, 1–3, 5–9, 11, 12, 14–17, 19–21, 24–28, 50–53, 55–57, 62, 64, 71, 72, 81</i>)
VM	Virtual Machine (<i>p. 74</i>)
VRF	Virtual Routing and Forwarding (<i>pp. xi, xii, 72, 74–78, 80–82, 85, 86, 89</i>)
 WAVE	Wireless Access in Vehicular Environments (<i>pp. 16, 17</i>)

WNIC	Wireless Network Interface Controller (<i>pp. 55, 57</i>)
-------------	---

Introduction

This chapter introduces the work presented in this dissertation, beginning with the context and motivation, followed by the main objectives and the methodology used to achieve them. It concludes with a summary of the major contributions and an outline of the document structure.

1.1 Context

The Internet is universally acknowledged as one of the most significant technological advancements in human history, with its widespread adoption and the numerous benefits it has brought about having irrevocably transformed the manner in which societies in the 21st century live and communicate.

While this impact is indisputable, the exponential growth observed over the past two decades has revealed significant flaws in Internet design. Indeed, in its original state, the Internet was not fully capable of satisfying the demands of the modern digital landscape [1].

One such issue pertinent to this work arises from the fact that our society is becoming increasingly mobile. With the advent of widespread automobile use, researchers and engineers were quick to recognize the vast potential advantages of enabling these vehicles to communicate with one another. Such developments have the potential to enhance a number of aspects of modern transportation, from road security to the quality of life for drivers. Given the direct impact of this technology on driver wellbeing, it is crucial that rigorous real-world testing be conducted to guarantee the safety of the general public.

For the past three decades, universities, companies, and governments have been engaged in collaborative efforts to develop the technology known as [Vehicular ad hoc Networks \(VANETs\)](#), whose objective is to enable vehicles to communicate directly with each other on the road. This process has made enormous progress in adapting the Internet to vehicular environments.

[Software Defined Networking \(SDN\)](#), a new networking paradigm, promises to deliver increased innovation, reduced cost, simplified management, and improved performance. It thus promises to improve the Internet and facilitate further growth. In light of this potential, researchers have recently focused their attention on this promising field, treating it as a potential solution to many of the challenges currently faced in the current Internet landscape. Notably, this has occurred in the domain of [VANET](#), which has led to the creation of the [Software Defined Vehicular Network \(SDVN\)](#) field. The potential for [SDN](#) to enhance

the capabilities of this technology has prompted recent studies to apply **SDN** to vehicular networks with the goal of developing optimized solutions for the future. The capabilities of **SDVN** are not proven to be superior to regular **VANETs**, but this thesis assumes that the application of **SDN** principles in **VANETs** is worthwhile.

1.2 Motivation

In accordance with the indispensable role that real-world testing plays in the ultimate success of **VANET**, it is of the utmost importance to conduct comprehensive testing in realistic settings to ensure the successful deployment of **SDVN**.

The majority of **SDVN** projects are designed for use in the real world, yet they are predominantly subjected to simulation-based testing. Although the outcomes of simulations are of great importance, they must undergo a process of confirmation and validation to ensure their reliability. The most effective method for achieving this validation is to compare the simulation results with outcomes derived from relevant, real-world testing environments. This emphasizes the pressing need for the availability of real hardware testing platforms to achieve more accurate and reliable results in this critically important area.

The development of a functional hardware testing platform, utilizing free and open-source software components for **SDVN** projects, represents a substantial advancement in the evolution of both the **SDN** and **VANET** fields. This may demonstrate the value of **SDN** beyond the data center, establishing it as a dominant force in the field while also fostering improvements in **VANET**.

1.3 Objectives

This dissertation presents a proposal for the design and implementation of a functional **SDVN** device. In essence, the core objective is to develop a device that can be classified as both a **SDN** and a **VANET** device in a single entity.

As a further clarification, the device will be designed and assembled to fulfill the role of an **On Board Unit (OBU)** and thus be able to communicate in the vehicular ad hoc space. The device must also conform to the design principles of **SDN**.

This document will undertake a detailed and thorough analysis of the **SDN** and **VANET** fields in order to define the characteristics of this device. Nevertheless, it is possible to establish a preliminary set of guidelines to which the device must adhere even before such an evaluation is conducted. The following two criteria serve as the fundamental basis for the evaluation of the devices in question. First and foremost, the device must be affordable. Secondly, the software components utilized must be open-source, wherever feasible.

1.4 Methodology

The methodology employed in this work was designed with the objective of assembling the necessary theoretical basis for an assessment of the viability of the **SDVN** technology as a whole. The selected methodologies are explained and justified in the context of the research objectives.

The dissertation commences with a comprehensive analysis of the **SDVN** field. In order to achieve the proposed objective of developing an **SDN**-compatible **OBU**, it is essential to conduct such research. While a review of the existing software in the **SDVN** field is undoubtedly important, the goal of the research step of this document is to extensively analyze the **SDN** and **VANET** fields.

In regard to **SDN**, this research aims to provide a rationale for its integration with **VANET**. It is therefore essential to present a comprehensive analysis of the underlying motivations, advantages, and fundamental principles that shape this technology.

Conversely, research on **VANET** is primarily concerned with defining the characteristics that distinguish this field of study from other ad hoc technologies, thereby establishing a comprehensive understanding of this field. Moreover, it is of critical importance to gain insight into the current state of this technology and to identify the remaining steps required for its widespread adoption. In light of this, this study is primarily focused on investigating the advancements in the field within the European context.

The research component of this document concludes with a review of existing literature, the objective of which is to identify any studies that share similar motivations to this document.

The practical component of this dissertation is based on with a period of rigorous experimentation and familiarization with existing software in the field. Subsequently, a new **SDN OBU** is designed and assembled, utilizing upgraded hardware and adhering to the requisite standards for **SDN** compliance.

To conclude, a concrete testing scenario is defined and subsequently executed with the intention of illustrating and validating the aforementioned achievements.

1.5 Main contributions

The primary contributions of this dissertation are as follows:

- A comprehensive and insightful analysis of the regulations governing **VANETs**.
- A detailed examination of **SDN** and its underlying motivations and benefits.
- A well-founded and realistic assessment of the potential implications of **SDVN** for future developments.
- An evaluation of hardware and software alternatives for a **SDN** capable **OBU** prototype.
- Creation of a functional **OBU** prototype with **SDN** support.
- The evaluation of the prototye in predefined scenarios.

These achievements are significant as they demonstrate in practice how it is possible to build and test a [Vehicle to Everything communication \(V2X\)](#) device with [SDN](#) and [Programming Protocol-independent Packet Processors \(P4\)](#) support, laying the groundwork for new experiments and advancements in the field of [SDVN](#).

1.6 Document structure

This document is divided into eight other main chapters. Chapter [2](#) provides a broad context about vehicular networks, in order to understand the main challenges in this field. In chapter [3](#), a very detailed overview of [SDN](#) is made, with a higher emphasis on the data plane, to show the power and potential of this technology. Chapter [4](#) addresses the advantages and difficulties of using [SDN](#) in a vehicular context, while also presenting the [SDVN](#) architecture. It also presents the studies in this field that share a similar objective to the present study are reviewed and analyzed. Chapter [5](#) details the implementation phase of this dissertation by delineating the design of the prototype in layers and concluding with a description of the hardware that will be utilized. In chapter [6](#), the pertinent software and hardware are evaluated, showing that the complete prototype proposed operates as planned. Chapter [7](#) concludes this document and outlines future work.

Vehicular ad hoc Networks

The automobile was one of the most revolutionary inventions in human history, for most cities around the world are dependent on cars for the transportation of goods and people and, therefore, the prevalence of vehicles in our daily lives is greater than ever before. Even though it is impossible to predict if future vehicle usage will continue to be as dominant as today, it is reasonable to assume that this technology will never disappear completely. Even today, in cities dominated by alternative modes of transportation, the use of motorized vehicles is still present and, in some cases, a necessity.

It has been a long-time goal for car manufacturers and researchers to bring internet connectivity to automobiles, as the benefits are enormous. The Internet is an extremely powerful technology, not only for the countless services it provides but also for the ability of two endpoints to communicate almost instantly. Moreover, recent improvements to wireless communication technologies not only make this goal more realistic but also even more attractive, as integrating network interfaces with [Global Positioning System \(GPS\)](#) receivers and sensors only increases the potential of allowing cars to communicate with each other [2].

The efforts to bring connectivity to automobiles resulted in the emergence of the technology named [VANET](#), which is an adaptation of the [Mobile ad hoc Network \(MANET\)](#) technology, specialized for automobiles [3] [4]. In contrast to wired networks, where all devices are connected to the infrastructure at all times to communicate with each other, this technology allows devices to transmit data directly to other devices inside its range, as illustrated in Figure 1. As a result, networks are generated spontaneously between nodes that are in range of each other, which makes topologies unstable and requires each node in the network to act as both a transmitter and a receiver.

As the Internet was not designed with mobility in mind, most protocols and standards were not developed to support the high-speed scenarios of [VANET](#). A tremendous amount of work and research efforts has already been conducted in this area, amounting to decades of research, development, and standardization. However, the development of [VANET](#) technology began without any major centralized entity to standardize its implementation worldwide, resulting in various architectures emerging simultaneously all around the globe. It is common for different regions of the world to develop different standards for infrastructure-related technologies and these solutions tend to converge as they seek to solve the same problems in the most optimal way.

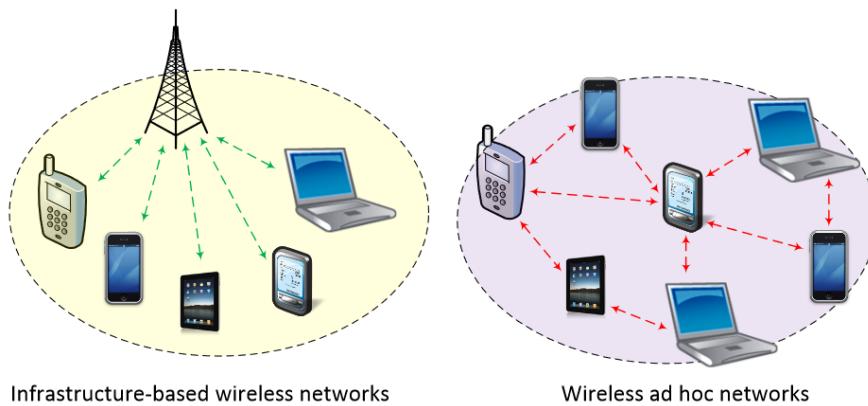


Figure 1: Comparison between infrastructure networks [5]

Numerous countries and institutions worldwide, including the [United States of America \(US\)](#), Japan, and the [European Union \(EU\)](#), have engaged in research initiatives concerning [VANETs](#). For the sake of clarity and relevance, this dissertation focuses on the European [VANETs](#) architecture through the work done by the [Car-2-Car Communication Consortium \(C2C-CC\)](#) and the [European Telecommunications Standards Institute \(ETSI\)](#), two prominent and influential institutions in the field.

The [C2C-CC](#) architecture proposal is designed with a high level of abstraction and is therefore fairly easy to understand. In contrast, the [ETSI](#) standards are significantly more detailed and technical, since they provide the official standard for the [EU](#). [ETSI](#) uses the term [Intelligent Transport System \(ITS\)](#) to refer to [VANETs](#), as their efforts to connect vehicles together goes beyond cars and aims to interconnect planes, trains, boats, etc.

There are other conventions in this field, such as the C-ROADS platform [6], which are important for the development of this technology. In fact, several conventions have influenced the [ETSI](#) standards and would be worth discussing. However, to keep this subject more concise and focused, only the [ETSI](#) and [C2C-CC](#) architectures will be covered.

2.1 **VANET characteristics**

The [VANET](#) technological domain was established as a specialized subset of the [MANET](#) research field based on some unique characteristics, and thus [VANET](#) shares most of its properties with [MANET](#). Some key differences require the development of novel solutions, and in this section, we will go through all the characteristics of [VANET](#) while shedding light on what sets this technology apart from its origin.

2.1.1 **High mobility**

[MANET](#) topologies are highly dynamic, driven by the mobility of nodes. High mobility leads to uncertain connectivity, which makes ad-hoc topologies extremely volatile and unpredictable, with a high chance of

partitions [7]. **VANETs** not only inherit this characteristic but also amplify its volatility since **VANET** nodes are vehicles capable of traveling at extremely high speeds.

For instance, two cars traveling in opposite directions on the same road will only be within range of each other for a few seconds. In this short time, the vehicles must attempt to establish a connection before attempting to exchange valuable information, which means some links may be disconnected before there is a chance to use them [4]. On the other hand, two cars traveling in the same direction can maintain a connection indefinitely as long as they keep on the same path at the same speed. These two scenarios are polar opposites, making the reliability of connections inconsistent.

2.1.2 Predicted mobility

Beyond being more dynamic, **MANET** topologies are also random, as nodes can move arbitrarily and at unpredictable times, with virtually no restraint. **VANET** topologies diverge from **MANET** in this matter, as it is possible to reliably predict the path of nodes in an ad-hoc network. Vehicles are primarily used on pre-built infrastructure [4], like roads and motorways, and have to follow predetermined paths to reach their destination, which makes it possible to predict their future location. Vehicles are also forced to obey road signs and traffic lights and have to respond to other vehicles' movements [3], making topologies even more predictable.

On top of all of that, if drivers disclose their desired destination it becomes possible to predict the complete route of a node with near 100% precision. It is important, however, to keep in mind that this is not always true, and in some instances, drivers may adjust their path in reaction to the information received by the network [4].

A negative consequence of road use can be seen in more linear topologies when compared with **MANETs**, which inevitably undermines path redundancy [7].

2.1.3 Power and Computational ability

Another relevant change from **MANETs** is the computational capabilities and power constraints of the nodes. **MANET** was designed with small devices in mind, which have limited battery capacity and computational power constraints. In contrast, vehicles have access to a reliable power source [4] [2], so there are essentially no concerns in this regard. However, cars should not be thought of as an infinite power source, and normal precautions should be taken when optimizing routing algorithms, etc.

2.1.4 Congestion and Scalability issues

VANET suffers from an absurdly variable node density. A vehicle can be on a road in a remote location with the nearest internet connection of other vehicle kilometers away from their location, or in the middle of a traffic jam at an intersection surrounded by hundreds of vehicles all trying to communicate within its range.

Congestion scenarios create a bandwidth problem that is exacerbated by the presence of nearby obstacles like other cars and buildings, which are common in an urban environment [7]. Wireless links are significantly more fragile than their wired counterparts because they are subject to fading, noise and interference [8]. This also forces researchers to come up with effective measures for sharing physical media, as another issue of concern is ensuring [Quality of Service \(QoS\)](#) measures for critical messages, like road hazard warnings, which presents a core issue [7].

Variable network density presents a necessity for the protocols developed to be capable of dealing with any situation thrown at them in the real world.

2.2 [VANET](#) components

The first step in the understanding [VANET](#) scenarios is the analysis of the basic elements involved. With this in mind, the atomic components of both the [C2C-CC](#) and [ETSI](#) are explained in this section.

2.2.1 [C2C-CC](#) approach

According to the [C2C-CC](#) architecture, three main components can be identified, these being the [Roadside Unit \(RSU\)](#), [OBU](#) and [Autonomous Unit \(AU\)](#).

1. [Roadside Units](#) are stationary devices placed alongside roads, usually in high-density zones such as junctions, intersections and traffic lights.

These devices have a reliable connection to the Internet and communicate with all vehicles within their effective range, bridging the gap between vehicle ad hoc networks and the rest of the Internet. [RSUs](#) can also be equipped with multiple communication devices for optimal wireless connectivity. In addition, [RSUs](#) can also take advantage of the ad hoc characteristics of on-board units by using them as transmitters, thereby extending their effective range, as depicted on Figure 2.

[RSUs](#) can run safety applications like fixed road hazard warnings (work-zone warnings, bridge warnings, etc.) by broadcasting safety messages to nearby vehicles, making them more than mere connection points for vehicles.

2. [On Board Unit](#) is a small device integrated into vehicles that operates in the ad hoc space as they receive, send, and forward messages between each other and with [RSUs](#) in an ad hoc manner. [OBUs](#) provide it's wireless communication capabilities to any connected device used within the vehicle.
3. The term [Autonomous Units](#) refers to all the various devices that run a single or a set of applications that connect to the [OBU](#) and harness its wireless capabilities. The [AUs](#) can be physically integrated or permanently wired to the [OBU](#) being an integrated part of the vehicle. These types of [AUs](#) are typically safety-related applications. [AUs](#) can also be wireless devices like laptops or smartphones, wirelessly connecting to an [OBU](#) and utilizing its capabilities for leisure applications.

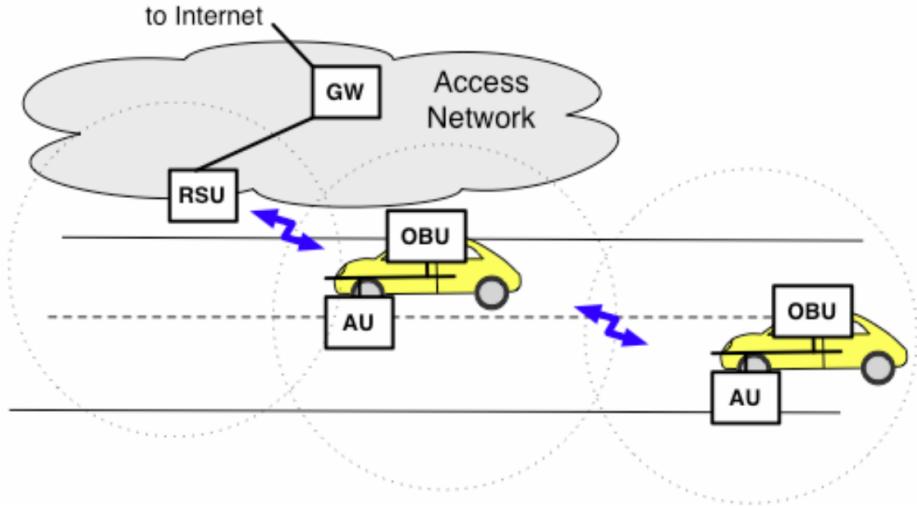


Figure 2: A RSU provides Internet connectivity to one OBU through another OBU [9]

4. In addition to these three main components, this consortium acknowledges the potential presence of other network entities, like public hotspots, instances of a Mobile Internet Protocol (IP) infrastructure, application servers, control centers, etc. However, these are not categorized and are therefore outside of the scope of this Consortium.

2.2.2 ETSI approach

The **ETSI** standard diverges from the simple and straightforward definition of the **C2C-CC** by specifying two categories to define the communication entities of **VANETs**: the **ITS** sub-systems and the functional components. The former, comparable to the **C2C-CC** components, represent similar system categories that exist in the **VANET** space. The latter are the building blocks that make up the **ITS** sub-systems and are therefore responsible for specific functions or tasks within the overall systems.

2.2.2.1 Functional components

The functional components described in this architecture follow the layered conceptual structure of the **Open System Interconnection (OSI)** model [10], comprising five such components: the **ITS station (ITS-S)** host, the **ITSS gateway**, the **ITSS router**, the **ITSS border router** and the **ITSS interceptor**.

The most significant component of this architecture is the **ITS-S** host, since it provides the minimum level of functionality required to run **ITS** applications. It represents the basic reference architecture for any device wishing to communicate in the **VANET** space, and is present on all **ITS** sub-systems. This architecture is described in depth in section 2.5

In addition, all other components are a variant of the **ITS-S** host, adding communication capabilities to it. The **ITSS gateway** converts messages between Internet and the **ITS** domain at layers 5 to 7 of the **OSI** model. The **ITSS router** converts various **ITS** protocols at layer 3, and the **ITSS border router** performs a similar task, doing the same as the router, but with a more generic network, without management or

security awareness. The [ITS-S](#) interceptor is a generic term equivalent to any of the previous three components, plus it can represent an implementation-specific connection of the [ITS](#) station-internal network to another network.

2.2.2.2 Sub-systems

As previously introduced, all functional components come together to form various sub-systems, of which there are four. These sub-systems represent a specific context in which the functional components combine to attempt to provide wireless communication capabilities, either in the form of direct contact with the ad hoc environment or indirectly.

All [ITS](#) sub-systems consist of the unique context they support and a specific [ITS](#) station to meet the needs of that unique situation. These stations can also be implemented in a single physical unit or in multiple physical units.

1. Personal [ITS](#) sub-system:

The Personal [ITS](#) sub-system represents a mobile or other personal device that can connect to the vehicle's internal network and use its ad-hoc wireless capabilities. It includes the user's device and a Personal [ITS](#) Station. This station is the simplest of the four, consisting of a single [ITS](#) host. This can be observed in Figure 3.

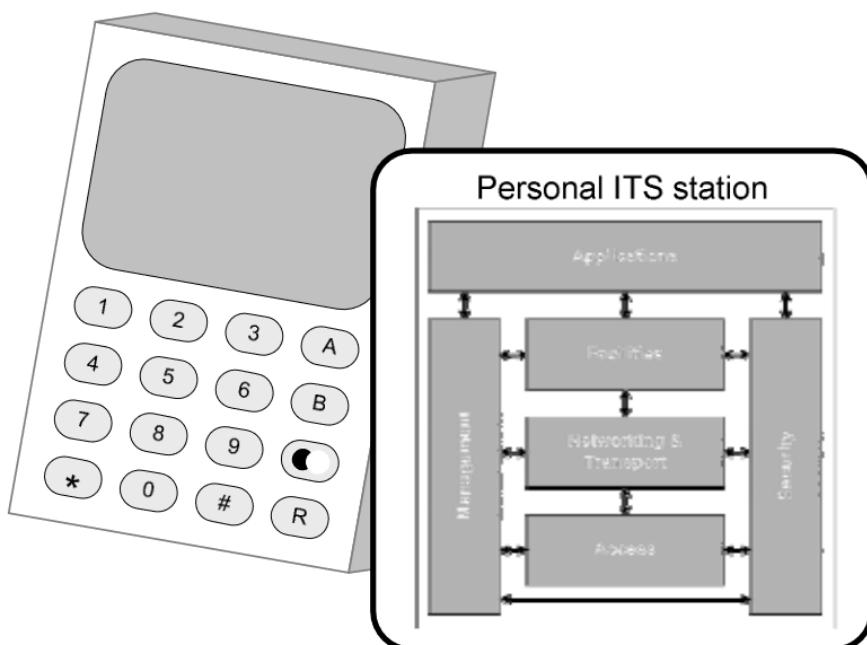


Figure 3: Personal [ITS](#) station in a Personal [ITS](#) sub-system [10]

This sub-system is comparable to an [AU](#), with a narrower definition. The [ETSI](#) architecture defines the existence of an [Electronic Control Unit \(ECU\)](#), which also falls under the [AU](#) definition created by the [C2C-CC](#). The [ECU](#) will be described along with the vehicle [ITS](#) sub-system.

- Central ITS sub-system; The purpose of the Central ITS sub-system is to provide centralized ITS applications, like traffic operations or content delivery, to VANET users. It is composed of a central ITS station that can be located anywhere. Examples of a Central ITS sub-system are traffic management centers and road operator centers.

The Central ITS station contains an ITS host and two optional ITS interceptors, being an ITS gateway and border router. The Central ITS sub-system is represented in Figure 4.

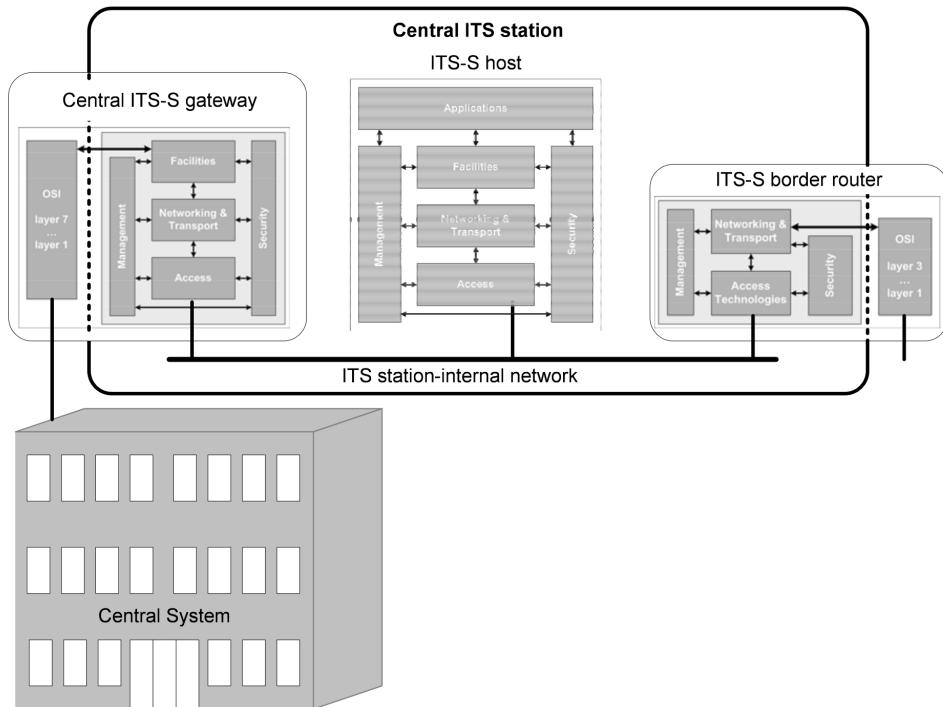


Figure 4: Central ITS station in a Central ITS sub-system [10]

The border router communicates with any other ITS system via any external network, while the gateway provides a more standardized connection.

3. Vehicle ITS sub-system:

The Vehicle ITS sub-system represents the cars, buses, trains, trucks, airplanes, and any other vehicle or means of transportation that can communicate wirelessly in an ad hoc manner and belongs to the VANET space. This definition is very similar to the OBU definition made by the C2C-CC, but there are some important minor differences that make it more rigorous.

This sub-system is made up of the Vehicle ITS station and the internal vehicle network, which is connected to a variety of ECUs. ECUs represent any kind of system or application that requires a connection to the Internet or to any other vehicle.

Like all stations, the Vehicle ITS station is made up of an ITS host with two optional ITS interceptors, nominally an ITS gateway and an ITS router. This sub-system is depicted in Figure 5.

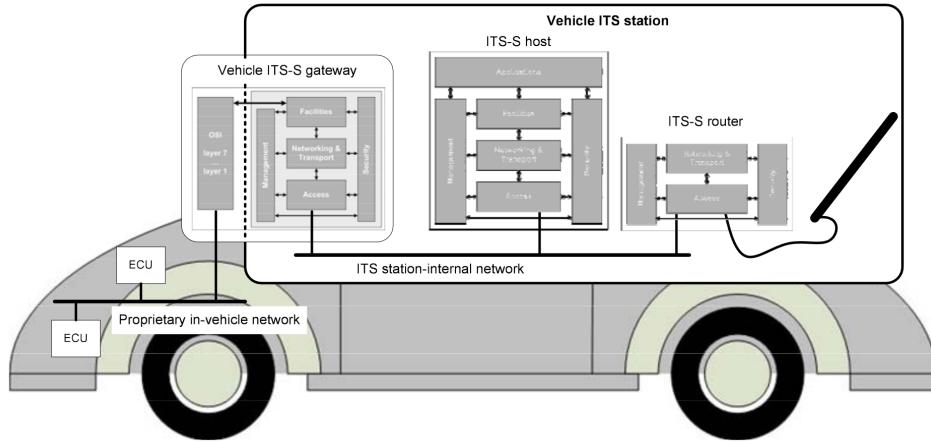


Figure 5: Vehicle ITS station in a Vehicle ITS sub-system [10]

The ITS gateway connects the internal vehicle network, which is proprietary in most cases, to the ITS station. Meanwhile, the ITS router connects to other Vehicle ITS stations and, most importantly, to the Roadside ITS station.

4. Roadside ITS sub-system:

The final element is the Roadside ITS sub-system. It is the connection point between the vehicle ad hoc network and the rest of the Internet. Like the RSU, it is placed in locations of high vehicle volume next to roads and intersections.

This sub-system comprises a roadside ITS station and any proprietary roadside network that may exist. It connects to any vehicle ITS station and can provide ITS applications to the vehicle users. This station is also connected to a central ITS station, being the bridge that allows the central ITS station to provide ITS applications to vehicle users.

The roadside station is made up of an ITS host with optional ITS interceptors. These interceptors can be any of the other types of functional components beyond the ITS host. The Roadside ITS sub-system is better depicted in Figure 6.

The ITS router connects the station to the ITS vehicle station, the border router connects to an ITS central station and the gateway connects to any existing proprietary existing networks.

2.3 Communication Domains

With the main building blocks of both architectures laid out, we can divide communications into various domains. Both the C2C-CC and ETSI have different ways to tackle this question.

2.3.1 C2C-CC approach

The C2C-CC focuses on the VANET space and divides it into three areas: the in-vehicle domain, the ad hoc domain, and the infrastructure domain. The first represents the network inside a vehicle, composed of an

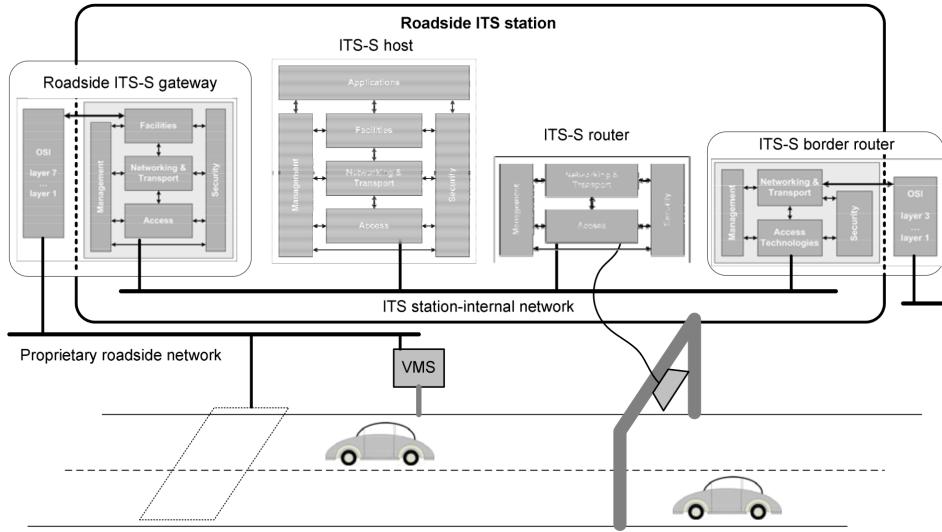


Figure 6: Roadside ITS station in a Roadside ITS sub-system [10]

OBU and the multitude of **AUs** connected to it. It encompasses all communications between the different **AUs** and the **OBU**. The ad hoc domain involves all ad hoc communications between different **OBUs** and between vehicles' **OBUs** with **RSUs**. The remaining components of the Internet that don't directly affect it, but help provide connectivity to vehicles, are included in the Infrastructure domain.

All of these domains, as well as the functional components of the **C2C-CC** architecture can be observed on Figure 7.

2.3.2 ETSI approach

The **ETSI** architecture proposes a broader categorization by establishing two domains. The first and most important is the **ITS** domain, since it contains all the elements and communication engaged by the **ITS/Architecture of Communications in ITS (ITSC)** standards. Everything else is encompassed by the generic domain, including everything in the rest of the Internet that might interact and communicate with the **ITS** domain.

2.4 Communication categories

In addition to communication domains, most authors use a classification based on communication types [9] [2] [7] [11] [12] [3] [13] [14] [15] [4] [16] [17] [18] [19] [20] [21]. This communication-oriented view of architecture aims to classify the multitude of types of communication that can occur between all the different entities that exist.

Different categories have emerged and evolved throughout the years as they have been used by most of the relevant researchers in the field. It is therefore difficult to find an origin for these concepts, and the following are most commonly used terms.

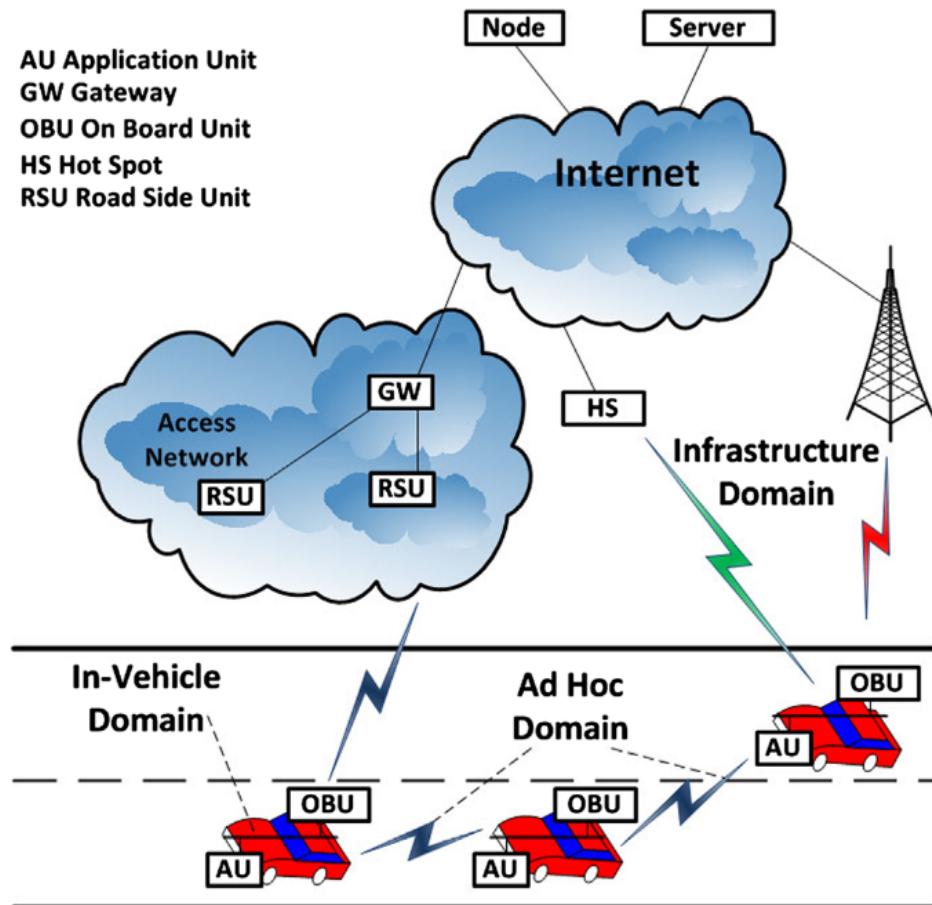


Figure 7: Communication domains in [VANET](#) as categorized by the [C2C-CC](#) [3]

1. **Intra-Vehicle communication** This term encompasses all communications that occur inside the vehicle between different vehicle components.
2. **Vehicle to Vehicle communication (V2V)** communication [V2V](#) communication refers to the ad hoc communication between vehicles.
3. **Vehicle to Infrastructure communication (V2I)** communication [V2I](#) communication refers to ad hoc communication between vehicles and roadside infrastructure. These communications do not encompass all messages that get exchanged between vehicles and the roadside infrastructure, and only include messages whose source and destination are a vehicle and a roadside station.
4. **V2X** The last type of communication is the most encompassing. [V2X](#) communication includes all communication that can occur in the [VANET](#) space between a vehicle and anything else. It is generally used to refer to the previous two types of communication.

2.5 VANET architecture

The architecture of an ITS host can be visualized on Figure 8. This section will dissect this architecture layer by layer in order to explain the functionalities and capabilities expected from it. The last layer, called Applications, will not be described as it represent ITS applications that use the services provided by the rest of the stack to connect one or more applications together and provide services to users [10].

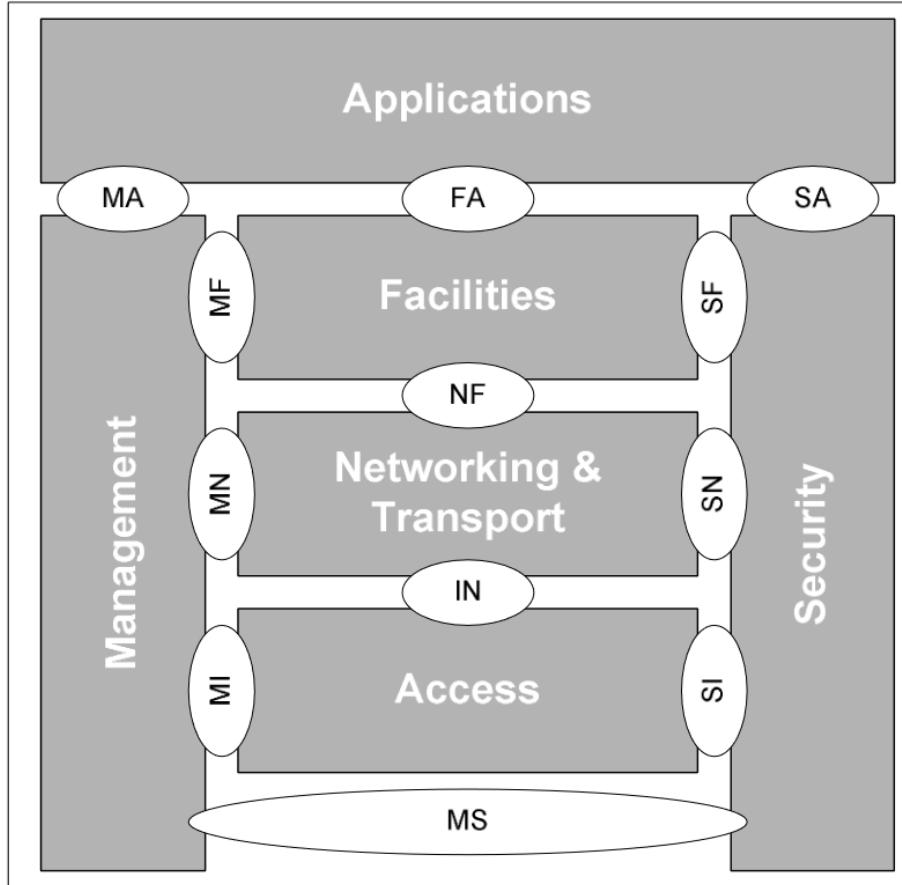


Figure 8: ITS station reference architecture [10]

The functionalities specified in each layer do not necessarily have to be implemented in every instance of the architecture, as it depends on the specific implementation of the ITS station [10].

2.5.1 Access layer

The Access layer, which is illustrated in the architecture of an ITS host in Figure 8, corresponds to the first two layers of the OSI model, the Physical layer and the Data link layer [22].

Among all the existing Access layer technologies, the best suited to meet the necessary requirements for use in VANETs were the existing wireless access technologies [3]. These technologies ranged from long range signaling technologies such as Microwave and WiMAX to short range technologies such as Infrared or Bluetooth [13].

Over the years, most of these technologies have been studied and tested as possible solutions for use in **VANET** communications, and in most cases they were deemed unsuitable for the **VANET** environment. For instance, short-range communication protocols such as the examples above were discarded due to lack of range, throughput, and reliability and as a result, medium and long range communications took center stage.

While unable to fully address the demanding scenarios of **VANETs**, a few of these technologies showed promise as possible solutions. Over the years since **VANET** research began, new standards and protocols have been developed based on these existing technologies, focusing on solving the problems posed by the characteristics of **VANETs**. In addition, these technologies themselves have evolved to meet the high data rates required by today's Internet.

Today, the **IEEE** 802.11 WiFi and Cellular technologies are considered as the two primary solutions for **VANETs** worldwide.

2.5.1.1 IEEE-based technologies

The **Institute of Electrical and Electronics Engineers (IEEE)** is an American organization dedicated to promoting innovation and technological excellence for the benefit of humanity. Founded in 1963, it is the world's largest professional society of engineers, scientists, and other technical professionals. **IEEE**'s primary areas of contribution are the electrical, electronic, and computing fields [23].

The **IEEE** 802.11 family of standards, an Access layer technology, was considered a promising solution for **VANETs** from the start. Not only would the existing **IEEE** 802.11 standards be cheaper to adopt due to being already established technologies, but WiFi also already met many **ITS** requirements due to its compatibility with mobile environments [17].

Despite these significant advantages, the existing **IEEE** 802.11 standards were not able to meet all the requirements imposed by the unique characteristics of **VANETs**, necessitating the creation of a new variant that would be better suited for such scenarios. Thus, **IEEE** 802.11p was created as an adaptation of the existing **IEEE** 802.11 standards. Given its American origin and contribution to the development of **V2X** communication technologies, it is important to mention that **IEEE** 802.11p was conceived in the context of **Wireless Access in Vehicular Environments (WAVE)** [24]. The **EU**, while initially open to many different access technologies, eventually converged on a **IEEE** 802.11p based solution dubbed **ITS-G5**.

IEEE's main motivation behind the development of **IEEE** 802.11p was to enable effective communication capable of handling the rapidly changing environments of **VANETs** [24], while making the minimum necessary changes to the **IEEE** 802.11 Physical layer. Modifying the Data Link layer is akin to a software update, while a Physical level amendment would require the design of an entirely new wireless airlink technology [24], which would be expensive and time consuming.

With this in mind, the **IEEE** began work on a variant of the **IEEE** 802.11 standard that would be feasible for **V2X**, capable of operating at speeds up to 200 km/h and ranges up to 1 km [2]. Experimental work began with the full deck of **IEEE** 802.11 technologies [7], but soon narrowed down to **IEEE** 802.11a because it operates at 5 GHz, which is closest to 5.9 GHz, the desired frequency for **ITS** operations in the

US. This made it easier to configure devices to operate at the desired frequency [24].

IEEE 802.11p uses orthogonal frequency division multiplexing. The IEEE 802.11 family of standards offers three channel bandwidths of 5, 10, and 20 MHz, and while IEEE 802.11a uses the full 20 MHz bandwidth, the 10 MHz bandwidth was chosen for ITS scenarios. This is achieved either by halving the clock/sampling rate or by simply doubling all the standard orthogonal frequency division multiplexing timing parameters [17].

Reducing the channel bandwidth was done in an effort to increase the root mean square delay spread [7]. Root mean square delay spread measures the time difference in seconds between the first and last signal components. These components represent different versions of the same transmitted signal, with the first to arrive coming from the shortest path which is the light of sight and the remaining coming from reflections and other interactions with the environment.

The standard 20MHz guard used by 11a proved sufficient to stop intersymbol interference between a radio and itself, so measures had to be taken to reduce the interference caused by multipath delays and by the Doppler effect. Halving channels width was the most simple and sensible solution to this problem [24]. Beyond reducing the signal bandwidth, the data throughput ranges were also reduced to 3 to 27 Mb/s instead of the original 6 to 54 Mb/s [7].

With the goal of enabling effective communications that can cope with rapidly changing environments, IEEE 802.11p communications needed to get faster in order to better utilize the sometimes narrow windows available for communication. Such a feat demanded a reduction in the overhead required before actual data transmission, which was achieved by simplifying the **Basic Service Set (BSS)** operations present in IEEE 802.11a [24].

The BSS represents a group of stations that are wirelessly connected to the same access point. It is created by an access point, and any device requesting to join it can exchange information with it after proper authorization. A BSS variant for ad hoc environments exists, called **Independent Basic Service Set (IBSS)**, which works without the need for infrastructure. However, it also proved to be insufficient to deal with the peculiarities of **VANETs**.

As a solution, IEEE 802.11p uses an ad hoc mode called **Out of Context BSS (OCB)**. OCB simplifies setup operations compared to its counterparts in other IEEE 802.11 standards by eliminating management procedures such as channel scanning, authentication, and association. This means that IEEE 802.11p has no setup time allowing stations to communicate with each other directly and immediately [25].

Removing all of these processes makes the network less secure. However, these security vulnerabilities are addressed by other standards in the ITS domain.

In Europe, ETSI created a solution for the access layer called **ITS-G5**. It not only defines the technologies and protocols to be used in communications, but also includes the frequency allocated exclusively for **V2V** and **V2I** communications for the whole of the EU [16]. ITS-G5 uses the existing IEEE 802.11p standards created for the American **WAVE**, adapting it for the European region.

In the US, the **Federal Communications Commission (FCC)** was responsible for the allocation of the frequency from 5.850 to 5.925 in 1999. In 2003, the FCC divided the American band into seven non-overlapping 10 MHz channels with a 5 MHz unused band at the lower end 9. These channels were

numbered from 172 to 184. This division also allows for operations in the 20 MHz channels 175 and 181, which are the overlap of two 10 MHz channels [15].

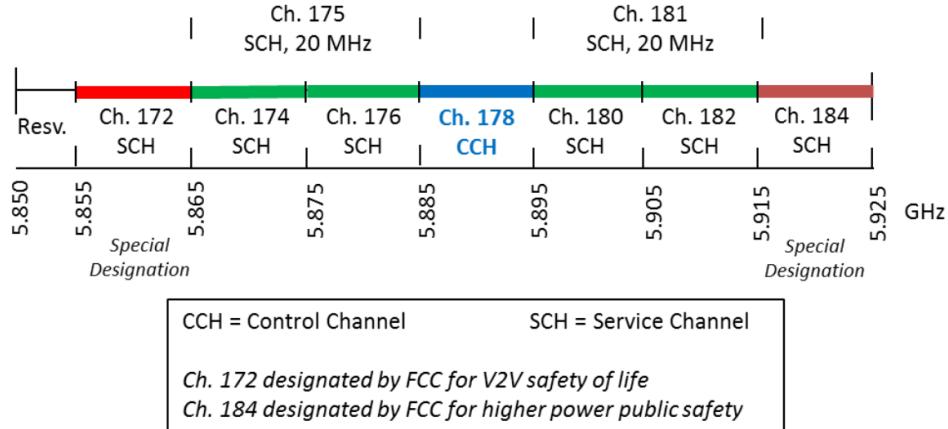


Figure 9: Channels allocated in the **US** by **FCC** for **ITS** [15]

In the **EU**, the **Electronic Communications Committee (ECC)** is responsible for spectrum regulation and the European Commission is responsible for enforcing the regulated spectrum in all member states. The **ECC** is made up of radio and telecommunications regulatory authorities from all member countries [15] [16].

The spectrum allocation for **ITS-G5** began in 2005, with a recommendation from **ETSI** to the **ECC**. In its TR 102 492-1 [26], **ETSI** recommended the allocation of 30 MHz in the 5,875 GHz to 5,905 GHz band for safety applications, which would align with the **US** control channel frequency of 5,885 GHz to 5,895 GHz. This recommendation also expected the allocation of a future 20 MHz band from 5,905 GHz to 5,925 GHz for future **ITS** extensions [15] [16].

ETSI based this recommendation on the allocation in the **US** frequency band, and also took into account the 5.8GHz toll collection band used in **EU** countries [15] [16].

In 2008, the **ECC** accepted this recommendation and allocated the requested frequency range plus another 20 MHz frequency band from 9,855GHz to 9.875GHz intended to be used by non-safety applications [15] [16]. All of the spectrum involved for present and future **ITS** operations can be easily observed in Figure 10.

Even though the American spectrum was taken into consideration, in the end the control channels of both the **EU** and **US** solutions were not in the same frequency, with the European channel being 10 MHz higher. This is not a total loss, as the same hardware can still be used in both countries requiring just a software change to work [15] [16].

The different allocated frequencies have different standardized transmission power limits based on use cases and interferences on adjacent bands and on the more important control bands [25]. The specific restrictions for each 10 MHz channel can be observed in figure 11.

In recent years, the **IEEE** began work on **IEEE 802.11bd**, a successor to the **IEEE 802.11p** standard, which was approved in October 2023 [27]. As a new standard, it is expected to take several years before

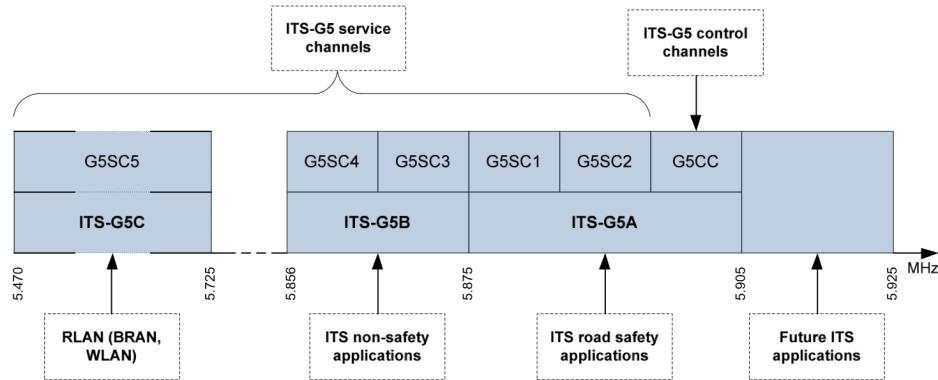


Figure 10: European ITS channel allocation [12]



Figure 11: Power spectral limits on each ITS channel [15]

it completely replaces 11p, and most automakers will continue to release cars with IEEE 802.11p for the foreseeable future.

This new version promises twice the throughput and longer range, while ensuring interoperability, coexistence, backward compatibility and fairness with existing OCB terminals. In addition, it operates beyond the 5.9 GHz band by utilizing the 60 GHz band.

2.5.1.2 Cellular-based technologies

Cellular systems have been used since the 1970s to transmit data over long distances via radio waves [13]. The 3rd Generation Partnership Project (3GPP) is an international consortium established in 1980 with the goal of developing technical specifications and technical reports for 3G mobile systems and is now the leading global reference for mobile standards and is responsible for the development of Long-Term Evolution (LTE) and 5G [28].

Although most work on VANETs has been based on IEEE 802.11p, Cellular V2X (C-V2X) has gained renewed support from academia and industry as this technology has matured [19]. Rival performance to the dominant IEEE 802.11p has led to Cellular technologies being considered by some organizations as a replacement for existing IEEE standards, so its status as the second most important VANET technology is not surprising.

In September 2016, the [5G Automotive Association \(5GAA\)](#) was established as a global, cross-industry organization of companies from the automotive, technology, and telecommunications industries with its main objective to promote the use of cellular [3GPP](#) standards in [ITS](#) scenarios [29].

Most researchers dismissed this new technology because its shortcomings made it seem inadequate compared to the strengths of [IEEE 802.11](#). [C-V2X](#)'s main drawback is its centralized nature. Cellular technology relies on a centralized infrastructure, which introduces latency by forcing all data to pass through a base station.

In addition, the large coverage area of cellular antennas can also pose an issue, as the coverage area of an antenna is typically larger than the zone of relevance of safety messages. This makes broadcast an inadequate solution, as many vehicles would receive warnings not intended for them. Multicast can be introduced in an attempt to reduce this problem, but it introduces new delays from the overhead required to create these multicast groups [19].

On the other hand, Cellular technologies have several technical and economic advantages for their application in [V2X](#) communications. On the technical side, cellular technologies provide wide coverage, support high vehicle speeds, support a large number of connections to a single cell tower and can deliver high data rates. Economically, the wide use of these technologies allows the reuse of already deployed hardware for use in [V2X](#) [19].

These benefits, along with the renewed support from the industry created from [5GAA](#), motivated the [3GPP](#) to begin studying the feasibility of [LTE](#) technology for supporting [V2X](#) communications [19]. In 2019, [3GPP](#) began work on [C-V2X](#) efforts with Release 14, which made great strides in reducing existing drawbacks.

In order to work as a wireless access technology capable of serving all [VANET](#) test cases, this technology needed to ditch the necessity of using infrastructure and become ad hoc. Therefore, Release 14 established a new mode of communication, called direct communication [18]. This resulted in [LTE](#) and 5G supporting two relevant interfaces for communication in [VANETs](#), the Uu interface and the PC5 interface.

The Uu interface is used for long range communication with cellular infrastructure in the commercial cellular spectrum. It uses existing [LTE](#)'s large coverage to provide vehicles with all kinds of services [18].

The direct communication mode, also referred to as "Mode 4", utilizes the short range PC5 interface in the [ITS](#) 5.9GHz band and allows for direct exchange of information between vehicles without passing through a cell tower.

[ETSI](#) initially considered Cellular technologies as possible solutions, and at that time included 2G and 3G as possible access technologies in the initial [ITS](#) station reference architecture. As stated earlier, [ETSI](#) centralized the Access layer of its architecture on the [IEEE 802.11p-based](#) [ITS-G5](#). Recently, however, reflecting on the efforts of [3GPP](#), [5GAA](#), and [ITS](#) implementation in other countries like China, [ETSI](#) made amendments to existing standards in order to achieve compatibility with [C-V2X](#) technologies [18]. This indicates that in the [C-V2X](#) architecture, the layers above the Access layer consists of the pre-established standards of different countries, demonstrating significant technological compatibility.

2.5.2 Networking & Transport layers

The Networking and Transport layer, as implied by its name, corresponds to the third and fourth layers of the [OSI](#) model.

layer 3 algorithms commonly deployed in traditional networks are unsuitable for use in [VANETs](#) [7], and even existing [MANETs](#) algorithms have proven to be inadequate for use in [VANETs](#) [4]. The characteristics of [VANETs](#), as described in section 2.1, present several unique challenges not only to existing layer 3 algorithms, but also to Layer 4 protocols, the most important of which are high node mobility and variable node density.

The high mobility of nodes leads to frequent topology partitions, which results in highly unstable routes. Therefore, pre-determining routes in these conditions is often irrelevant and extremely complicated. Additionally, the algorithm must operate under conditions of both extreme congestion and isolation due to variable node density, further increasing its complexity.

Researchers have recognized the adoption of broadcast as the primary communication mode as a great solution to mitigate interference in high congestion scenarios. Nevertheless, this approach alone falls short of the desired outcome. For instance, in case of a crash, it is essential to rapidly disseminate security messages as the number of affected vehicles can rapidly rise. However, if every vehicle broadcasts simultaneously, it will lead to channel congestion, making it more difficult for the event to spread [7]. Therefore, novel approaches are required to curb message duplicates.

Although the challenges presented by these aspects of [VANETs](#) affect algorithm reliability, researchers have leveraged other characteristics of [VANETs](#) to develop novel solutions to the aforementioned problems. Besides the lack of power constraints, the predictable mobility of nodes allows algorithms to perform more effective link selection, which facilitates network management and opens opportunities to optimize the flow of information, ultimately improving routing protocols.

In Europe, researchers at [ETSI](#) developed new algorithms and defined new protocols in an attempt to find an adequate solution to the aforementioned challenges while taking advantage of [VANETs](#) newfound opportunities. As part of its architecture, [ETSI](#) introduced a new layer 3 protocol called GeoNetworking [30] [31] [32] [33] [34] [35], a new layer 4 protocol called [Basic Transport Protocol \(BTP\)](#) [34] and a sublayer 3 extension called [GeoNetworking-IPv6 \(GN6\)](#) [35].

GeoNetworking is an ad hoc routing protocol that was developed solely for the transmission of safety related messages and therefore it provides the ability to forward packets without the need to exchange any signaling messages beforehand.

The GeoNetworking protocol uses the geographical coordinates of nodes as their address and utilizes location-based data to help make packet forwarding decisions.

By using a device's geographic location as its address, this protocol enables packets to be sent to all nodes within a specific geographical area. Nodes can effortlessly broadcast messages to an entire geometrically shaped area without congesting the entire network. This approach proves to be more efficient than broadcasting because it curbs congestion by minimizing transmissions to unintended destinations. Moreover, it makes the packet relevance area independent of the sender's range.

This protocol was built and optimized for multi-hop communication with geo-addressing, providing more technical features in application support than the alternative in the [US](#). These capabilities come to a cost to protocol complexity and message overhead [14].

The GeoNetworking protocol was designed to be integrated with the [BTP](#) protocol. [BTP](#) is a transport layer protocol that is similar in function to [User Datagram Protocol \(UDP\)](#), functioning as a connectionless protocol [25].

[ETSI](#) has established the GeoNetworking extension [GN6](#) as an alternative to routing packets through [BTP](#). By using the [GN6](#) sub-layer, [Internet Protocol version 6 \(IPv6\)](#) packets can be transmitted over the GeoNetworking protocol without the need to modify the [IPv6](#) protocol [25].

2.5.3 Facilities layer

The Facilities layer encompasses layers 5, 6, and 7 of the [OSI](#) model, and its standardization defines application-specific functionalities [14]. The development of Facility layer messages is the responsibility of European institutions such as [ETSI](#), [European Committee for Standardization \(CEN\)](#), and [International Organization for Standardization \(ISO\)](#).

All the messages defined by European standard makers serve safety purposes, and therefore run over the GeoNetworking protocol [14]. Multiple different messages have been defined over the years [25], which makes it difficult to list them all. However, [CAMs](#) and [DENMs](#) are universally acknowledged as the most significant in the context of [ITS](#).

A [Cooperative Awareness Message \(CAM\)](#) [36] is a periodic safety message that contains vital status information about the originating vehicle, with the main goal of enabling other vehicles to take appropriate preemptive measures to avoid potentially dangerous situations [3].

This objective is accomplished by exchanging data, including speed, location, direction, and additional non-safety application data with nearby [ITS](#) stations, allowing vehicles to monitor each other's movements [14].

[CAMs](#) begin transmitting as soon as the vehicle enters a safety-relevant context, which is considered to be anytime the vehicle is in operation [25]. The transmission rate of [CAMs](#) is subject to specific rules, including both maximum and minimum transmission times between [CAMs](#), relevant changes in position, direction, or velocity, and congestion in the wireless channel [25].

[CAM](#) fields are shown in figure 12. It includes an obligatory [ITS Packet Data Unit \(PDU\)](#) header, Basic container, and High frequency container, along with optional Low frequency and Special vehicle containers.

This structure allows for a high degree of flexibility in message format, allowing messages to adapt more effectively to their environment, thereby minimizing congestion on wireless channels [25].

A [Decentralized Environmental Notification Message \(DENM\)](#) [37] is an event-driven safety message that can be triggered by an [ITS](#) station in the event that hazardous conditions are detected. [ITS](#) stations are capable of detecting a wide range of events and this message type is employed to describe such events to other [ITS](#) applications.

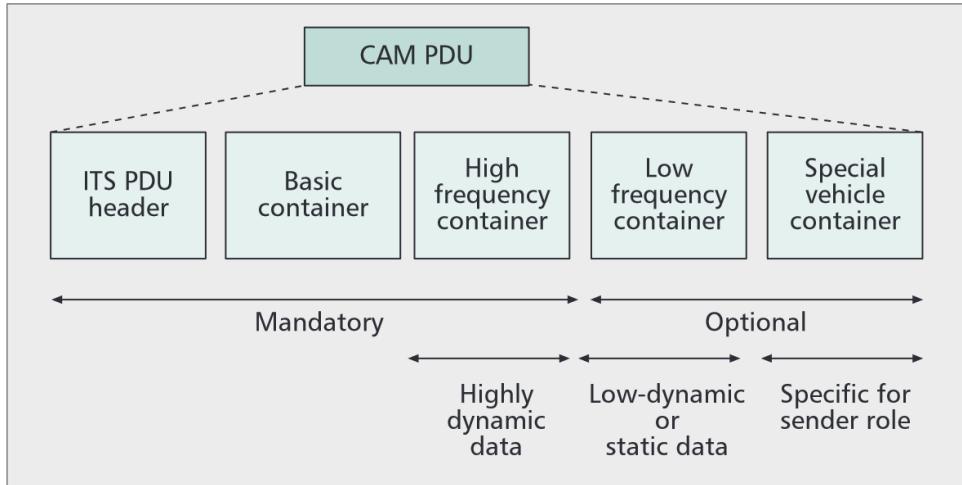


Figure 12: CAM structure [25]

DENMs serve the purpose of warning ITS applications to a detected event within a designated geographical area [25]. This type of messages are only relevant in a specific location and therefore are only disseminated in that specific geographical area. These geographical areas are the ones defined by the GeoNetworking protocol.

As an event-triggered message, DENMs are considered high-priority messages, as ensuring a quick delivery is crucial to diminishing the consequences of the events that triggered their generation [3].

Throughout the lifespan of the DENM, a number of techniques are employed to ensure the dissemination of the message in its relevant area. DENMs are repeated, generally at a lower frequency than CAMs, with the intent to allow vehicles entering the relevant area to receive said information. Similarly, if the originator ceases to transmit the DENM message for any reason, another ITS station can replace the originator and continue to transmit the message. DENM messages can also be canceled by the ITS station that created them [25].

DENM fields can be observed on figure 13. The ITS PDU header and the management container are mandatory, while the situation, location and a la carte container are all optional. These message fields mainly contain information about the relevance area of the message, the type of event and the time in which the message remains relevant [3].

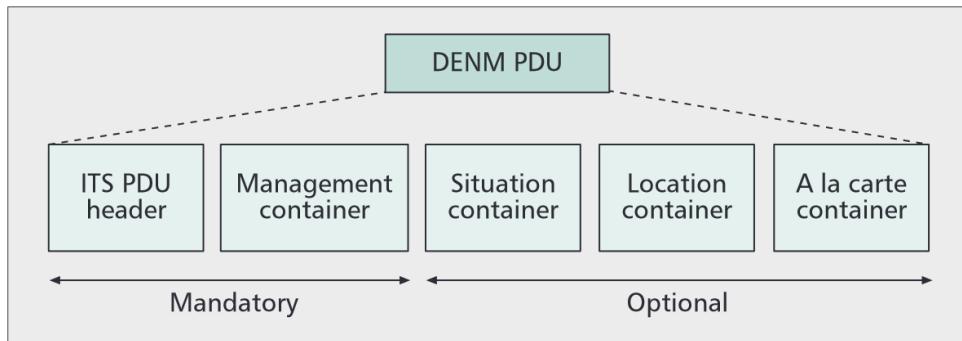


Figure 13: DENM structure [25]

2.5.4 Management Entity

Network management is a critical component of network maintenance as it ensures that a network operates as intended. The unique challenges presented by [VANETs](#) contribute to a significantly more unstable network environment than usual, making management even more essential. With this motivation in mind, [ETSI](#) incorporated a management entity into its architecture [14](#).

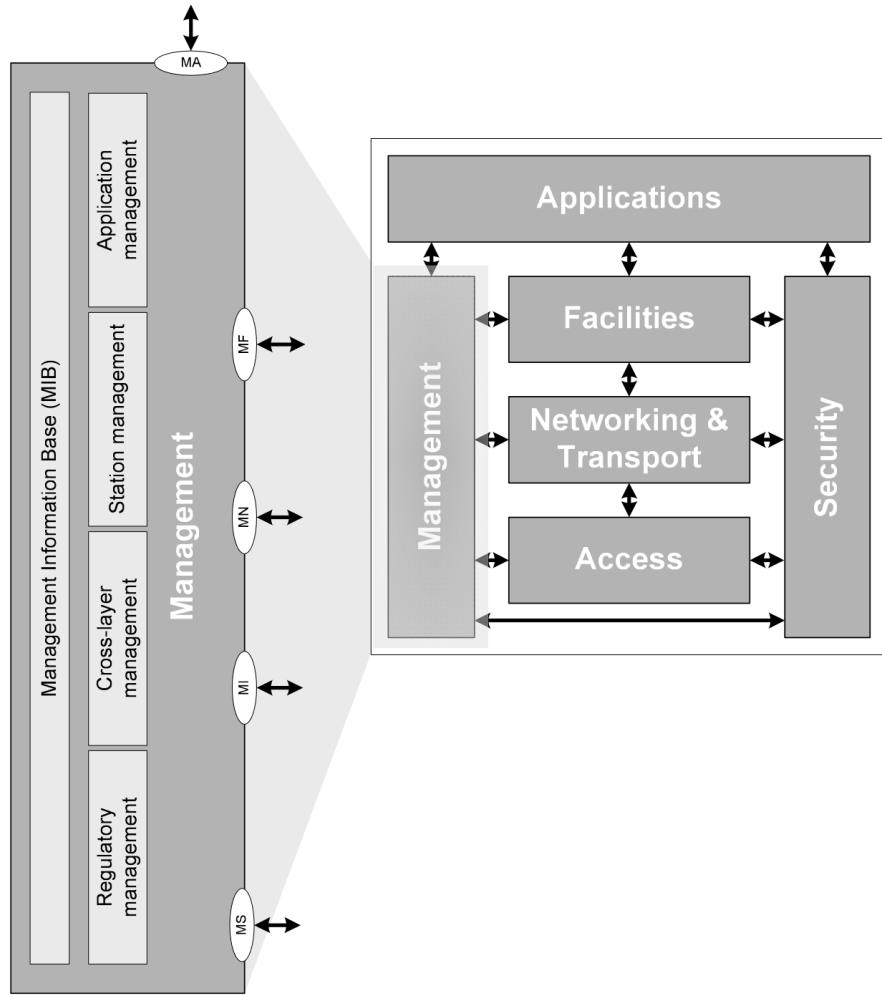


Figure 14: [ITS](#) management entity as part of the [ITS](#) station reference architecture [10]

The [ITS](#) management entity is responsible for both configuring and operating its [ITS](#) station, while also overseeing cross-layer information exchange between multiple layers. It contains interfaces to every other component in the [ITS](#) station and a management information base [32].

Decentralized congestion control is a cross-layer [ITS](#) functionality that is coordinated by the management entity and is one of its most important responsibilities. Its main purpose is to control congestion in the channel by managing the amount of messages exchanged, the transmission power and any other useful parameter. The changes to these parameters, being orchestrated by the management entity, are made based on various information extracted from different layers [25].

2.5.5 Security Entity

Security in [VANETs](#) has been one of the biggest concerns and one of the biggest obstacles to its deployment. It is crucial to protect the [ITS](#) domain from malicious attacks and network abuse, and as such it has been a high priority for researchers and developers to address security threats prior to deployment. Vehicle messages can contain a large amount of sensitive information, such as vehicle trajectory and location data. This information can be used to deduce the driver's identity, activities, habits, and so on. This type of information must not only be kept private under [EU](#) law, but could also be used by bad actors for extortion [4] [38]. Attackers could also exploit safety messages for their own benefit or to the detriment of others. For instance, greedy drivers could broadcast false traffic alerts to reduce traffic on their own routes. In a more extreme case, robbers or terrorists could abuse traffic alerts to clog roads and thereby delay emergency vehicles [38]. With this motivation in mind, researchers have laid out several security requirements [VANETs](#) must meet. The following list contains the most important ones, based on the works of [39] [38]:

1. Authentication: this security requirement ensures that a recipient can identify the sender of all messages received, thereby ensuring that each message is generated by an authenticated user.
2. Non-repudiation: sometimes called auditability, it is a tricky to enforce but essential security requirement. Non-repudiation ensures that once a message is sent, the sender can't deny ownership of the message. In [VANETs](#), this requirement is essential for identifying compromised users.
3. Integrity: this security requirement ensures that the message remains unchanged during transmission.
4. Confidentiality: a security requirement for any type of network, and [VANET](#) is no different. Some messages contain important information and therefore should only be accessible to the sender and receiver to prevent eavesdropping.
5. Availability: one of the most important concerns in [VANETs](#), this requirement seeks to ensure the availability of the wireless channel, as a message that is delayed by seconds becomes useless.
6. Access control: this property creates different levels of access for different entities. Access control bars users from accessing any information or sending message types they are not allowed to. Distinguishing multiple levels of access allows for a higher degree of network control and is essential to stop known bad actors from exploiting network dynamics.
7. Privacy: Privacy is the ability of users to hide their personal information from the rest of network users. Implementing mechanisms to protect the privacy of all drivers is a top priority to ensure driver privacy, with the main goal to provide location privacy. Anonymity is the process of hiding one's identity, which is extremely important to achieve privacy.
8. Data verification: this property is essential to avoid false messaging as it allows all messages to be tested by their time relevance. This is necessary to avoid replay attacks in the network.

9. Physical Security: as vehicles are a widely spread technology that will be distributed indiscriminately, it is important to implement hardware security in order to avoid tampering and compromising of the vehicle.

Authentication and privacy are desirable properties of [VANETs](#) but ensuring anonymity while enforcing network liability in [VANETs](#) is a somewhat conflicting property in securing vehicular networks. This comes from the necessity to protect drivers' privacy while being able to track down attackers [24] [38]. Another important consideration to take is that security measures can create performance problems, as they add overhead time in communications and can significantly degrade message quality [7].

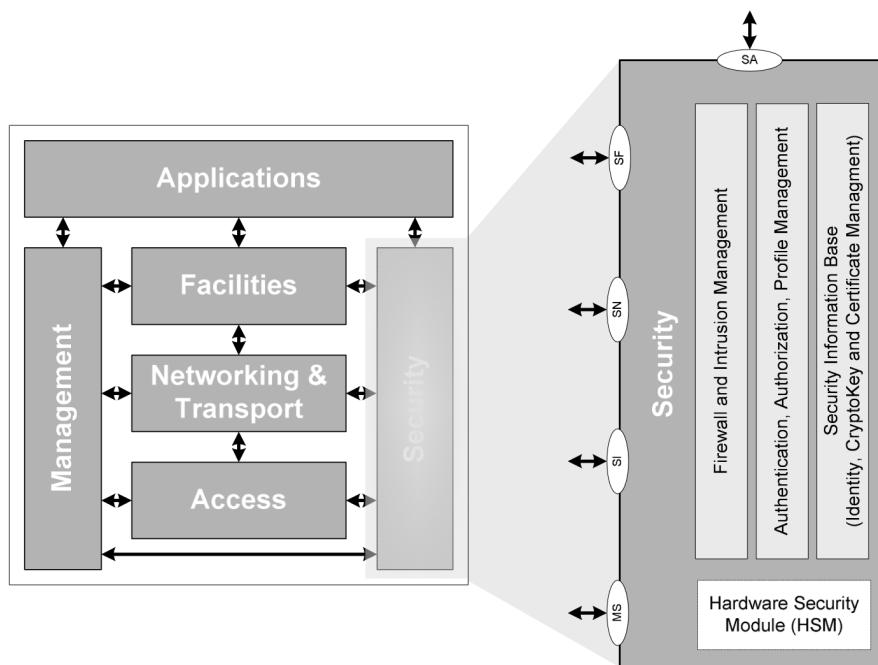


Figure 15: [ITSC](#) security entity as part of the [ITS](#) station reference architecture [10]

The Security Entity provides a plethora of services to ensure security and privacy. These include a multitude of secure messages at different layers of the communication stack, management of identities and security credentials, and other aspects relevant to secure platforms such as firewalls, security gateways, and tamper-proof hardware [32].

[ETSI](#) developed an [ITS](#) communication system that relies on indirect trust relationships, which are built upon trusted third party certificates. It is a solution to the privacy problem and results in the implementation of a so-called authority hierarchy. This hierarchy is composed of manufacturers, enrolment authorities, authorizations authorities and the [ITS](#) station.

When a car wishes to begin communicating, it must use its canonical credentials given by the manufacturer to request valid enrolment credentials to an enrolment authority. Then, using the enrolment credentials it must request an authorization authority for authorization credentials. These last ones are the ones that will be used for communications, and as such in order to protect a driver's privacy are only valid in a time frame. After these expire, the [ITS](#) host must request new authorization credentials to the

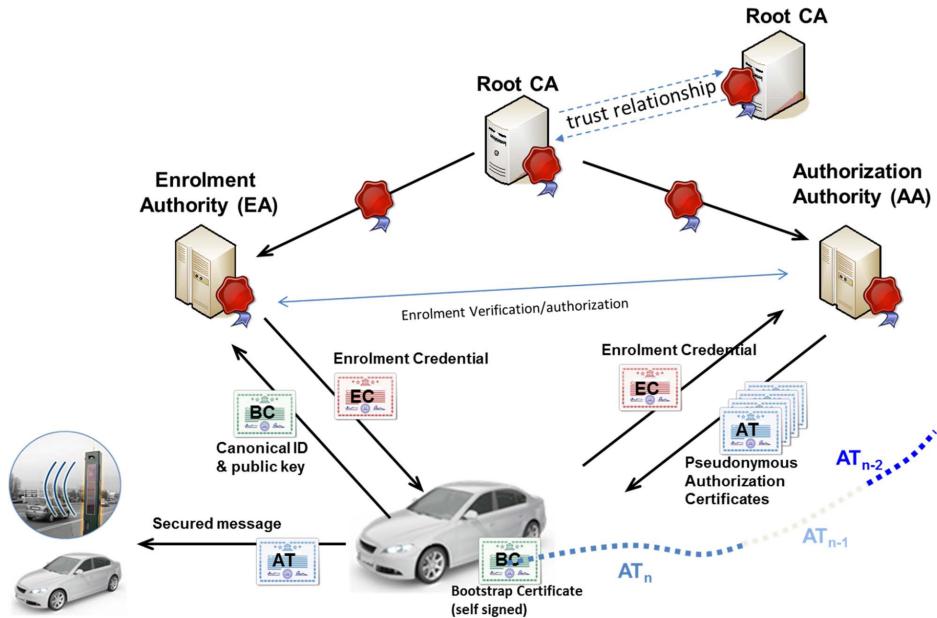


Figure 16: ITS Security Certificate Management System [40]

authorization authority. From this example, we can observe that the enrollment authority validates that the vehicle has a valid ITS station and that the authorization authority uses the authorizations given by the enrollment authority to allow the ITS station access to the ITS domain and to utilize a specific type of applications, service or privilege. This process can be observed on Figure 16 [41].

2.6 Applications of VANET

The motivation behind VANET research and development is to provide wireless services to both drivers and vehicles. The goal of these service applications is to improve driving safety, passenger comfort, and traffic efficiency. Hence, VANET applications can be classified into safety and non-safety applications according to their intended goal. Safety applications are considered the main driving force behind the development of VANETs [4] as the main goal of these types of applications is to decrease both road fatalities and road pollution [3]. Safety applications can be divided into two categories: Cooperative road safety and Cooperative traffic efficiency. Certain applications have the potential to improve both of them. Cooperative road safety applications harness the wireless communication capabilities of vehicles to provide useful information to the vehicle and driver, with the ultimate goal of mitigating the likelihood and severity of accidents. Therefore, any application intended to enhance the safety of individuals is classified as an Cooperative road safety application. Examples of such applications are notifications for upcoming hazardous locations and approaching emergency vehicles. Cooperative traffic efficiency applications aim to make road operations more efficient. These efforts can both greatly reduce unnecessary carbon emissions and save the time of drivers. Examples of such applications are off street parking information and green light optimal speed advisory. VANETs could also become a crucial step in the journey to fully autonomous vehicles, by exploiting expected advanced cooperation systems to exchange sensor information and status

information among vehicles. These future applications also fall under the cooperative traffic efficiency umbrella. Some of these envisioned safety [VANET](#) applications require a certain percentage of road-wide deployment to become useful [2]. Therefore, [ETSI](#) has defined a basic set of applications to be deployed as [ITS](#) systems mature, with the aim of deploying them simultaneously at that time. These applications are commonly referred to as Day 1 applications and aim to provide societal and economic benefits to both the private and public road transportation sectors. Remaining safety applications have been grouped into Day 1.5 deployments, with the goal of improving and extending Day 1 applications. These applications can be further divided into bundles, which can be viewed on Annex I. The remaining applications fall under the category of non-safety applications. Such applications, also known as co-operative local services and global internet services, are those designed to enhance the comfort of drivers and passengers by enabling access to internet services from their vehicles [3]. Based on this, any application that provides value-added services is considered non-safety applications [7]. All types of entertainment and information-based applications, such as music, movies, podcasts, online games, or instant messaging platforms, are examples of the applications covered under this category. Essentially, any application that is hosted on the World Wide Web and is accessible via the [IP](#) stack falls into this category. Non-safety applications should not interfere with safety applications, and thus they use different physical media and protocols [2]. In detail, non-security applications are typically delivered using [IPv6](#) over [C-V2X](#), while safety applications exclusively use [GeoNetworking](#).

2.7 Future trends and challenges in [VANET](#) research

Predicting a date for when the full potential of the [VANET](#) technology will be unleashed is incredibly difficult but it is a case of when, not if, as the benefits of [VANETs](#) remain attractive and a lot of work has already been done.

Current and future research seeks to update current technologies in use, such as the aforementioned development of [IEEE 802.11bd](#) 2.5.1. Additionally, efforts are being made to augment this technology with other technologies such as [SDN](#), Edge Computing, and Artificial Intelligence to take it to the next level [42].

The deployment of [VANET](#) infrastructure in Europe is currently underway, with projects such as Trans-European Transport Network. In line with the European Green Deal and the renewed focus on the climate crisis, [VANETs](#) are expected to help reduce traffic emissions by increasing traffic efficiency. Besides, the goal to reduce traffic fatalities to 0 by 2050 remains a top priority [43]. Figure 17 provides a comprehensive representation of the plans for [VANETs](#) in the EU.

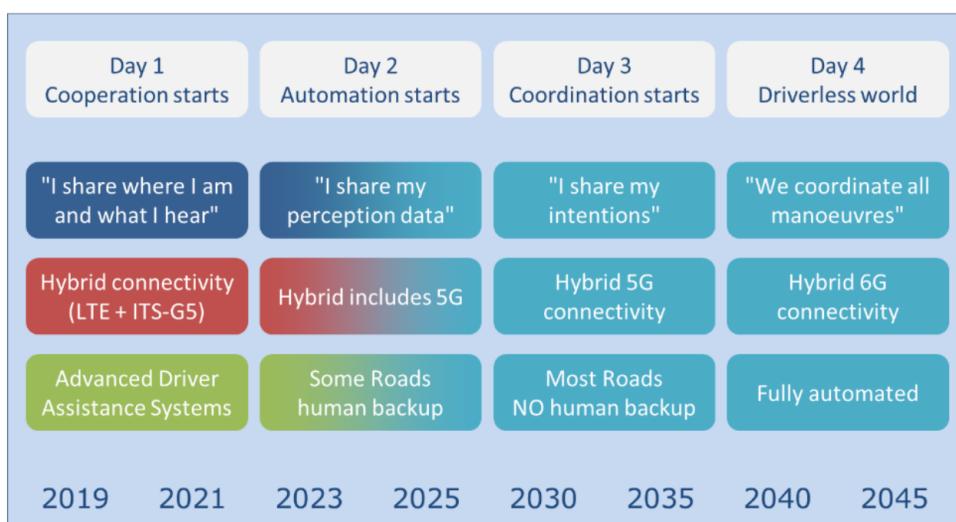


Figure 17: ITS strategy for the EU [43]

Software Defined Networking

[SDN](#) is a networking paradigm that simplifies network management and promotes innovation through software-based approaches, rather than traditional hardware-based configurations. [SDN](#) has the potential to shape the future of the Internet and is therefore often referred to as a "radical new idea in networking" [44].

[SDN](#) dates back to the 2000s, and its concepts can be traced as far back as the 1990s [1], so this technology is not exactly new. The core ideas and motivations behind [SDN](#) have been worked on and evolved over the past thirty years, with several previous efforts already being undertaken to promote network programmability [1]. Even though it is not the first, [SDN](#) is considered the most promising as its uniqueness stands from providing programmability while offering to greatly simplify network devices [45].

In 2011, the [Open Networking Foundation \(ONF\)](#) [46] consortium was founded with support from major players in the technology industry, including Google, Facebook, and Microsoft. [ONF](#) develops open standards for [SDN](#) technologies to promote the [SDN](#) vision. It is currently the leading consortium in this field.

The information in this dissertation on [SDN](#) is largely supported by the book "Software-Defined Networks: A Systems Approach" [47], authored by leading researchers in the field and cited by the [ONF](#) itself. Before delving deeper into this topic, it is worth emphasizing that [SDN](#) is an approach to networking, rather than a point solution, so while some of the most important and widely adopted technologies will be mentioned in this chapter, they should be viewed as an attempt to achieve the [SDN](#) vision.

3.1 SDN definition

Traditional networks have a highly decentralized structure, which was essential for the early Internet to achieve high network resiliency. As a consequence, Internet devices are vertically integrated, meaning the network control logic is directly coupled to the underlying forwarding hardware. This leads to routing decisions being made by complex programs running on each device through the definition of low-level policies [48].

For a long time, researchers and developers have acknowledged that changes in the way the Internet works, nominally in its structure, were necessary to simplify and automate network management[47].

These changes came to take place as the technology known as **SDN**, which breaks down the vertical integration of devices [49].

3.1.1 Main principles of **SDN**

SDN technology nowadays is rather an umbrella term, as many different researchers have their own nuanced definition arising from different levels of implementation of the same **SDN** vision. Such discrepancies in the **SDN** definition are acknowledged by [47] and can be visualized in [1], [49], [44], and [48]. The early **SDN** proposal stood on three main principles, which are still relevant to this day and are as such:

- First and foremost this technology is characterized by the decoupling of the control plane from the data plane.

The control plane can be considered as the “brain” of the network. It is responsible for commanding forwarding operations in the whole network, making decisions on how to handle network traffic.

The data plane is the remaining forwarding hardware that follows the commands set by the control plane. By removing the decision making components from network devices, it makes them simple forwarding devices.

Network intelligence is thus removed from individual devices, transforming all networking equipment into simple packet-forwarding switches.

- The second principle of **SDN** requires a logically centralized control plane that manages the data plane through well-defined **Application Programming Interfaces (APIs)**. This means that a single entity, commonly known as the controller, oversees all of the network.
- Lastly, the controller must have the power to control network behavior through the programming of network functions. **SDN** must allow software developers to harness and utilize network resources similarly to computing resources and storage.

To sum up, and as can be seen in Figure 18, **SDN** requires the decision making elements of a network are decoupled from forwarding devices and logically centralized in a single entity, which is a programmable controller.

As a consequence from these principles, an **SDN** switch can function as any traditional Internet device, such as a router, switch, firewall, network address translator, or even a combination of these.

3.1.2 Recent advancements

These three principles remain relevant today and form the de facto definition of **SDN**. However, in recent years, new efforts have been made beyond this initial definition to further fulfill **SDN**’s original goal. These efforts aim to increase programmability to the data plane, allowing network operators to fully customize all traffic in the network.

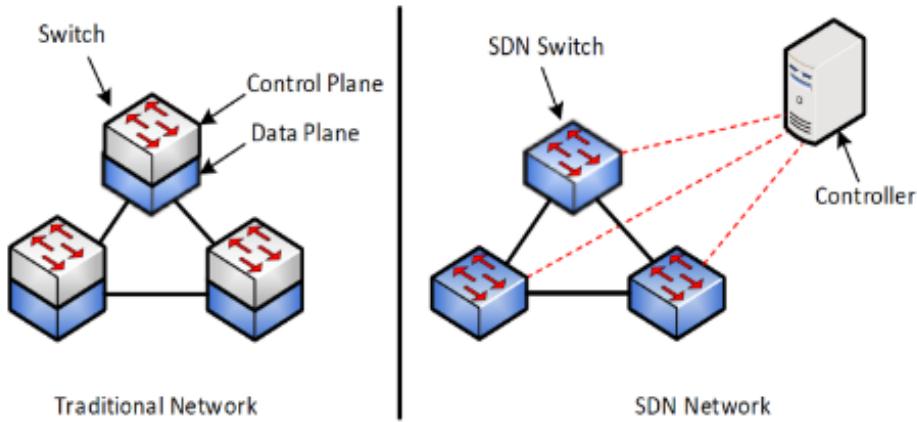


Figure 18: Traditional network view Vs [SDN](#) Network view [50]

Some articles refer to these recent efforts as next generation [SDN](#) [51] [52], while predicting further and further coupling of [SDN](#) with network virtualization.

Being one of the first use cases found for [SDN](#), Network virtualization traces its roots to the 90's, around the same time the idea of programmable networks was emerging [53] [1].

Network virtualization enables the creation of network topologies decoupled from the physical topology. This allows multiple virtual networks to share the underlying physical infrastructure, functioning seamlessly over the same physical network [49]. The advantage of network virtualization is that it allows each individual virtual network to be simpler than the real network topology, granting increased flexibility.

Network virtualization and [SDN](#) are two distinct technologies that do not require each other to exist. However, [SDN](#) technology enables network virtualization, effectively creating a mutually beneficial relationship between the two that has grown much closer over the years [1].

Another way to interpret these recent [SDN](#) developments is to divide efforts into two phases. In the first phase the control plane is centralized and opened for operators, for more network control, and in the second phase the data plane is opened, allowing network operators to better control packet processing in the network.

In reality, [SDN](#) implementations may mean more or less than the original definition depending on the involved stakeholder that is deploying the technology, as each network operator has the freedom to fulfill this view however they wish [47].

From this, and from the fact that [SDN](#) is a disruptive technology, hybrid solutions that aim to harness some of the [SDN](#) benefits while maintaining the status quo in the industry have appeared [47]. These solutions mainly appear as attempts from manufacturers to maintain their established business models.

Lastly, it should also be noted phase 1 [SDN](#) was largely impelled forward by Openflow, while phase 2 is marked by the development of P4. Both these technologies will be further discussed in Section 3.3.

3.2 Motivation behind SDN

The Internet currently faces numerous challenges, ranging from expensive, complex network structures that are difficult to manage and optimize, to significant barriers to the development and deployment of new network design ideas. The networking research community and industry have recognized that these challenges require a different approach to network structure. [SDN](#) promises to solve these problems by providing programmability, flexibility, improved performance, and support for implementing and testing new network ideas.

This section explores the benefits of [SDN](#) to demonstrate its usefulness and explain its motivations. As this document will not delve into the [SDN](#) disadvantages, it is important to keep in mind that [SDN](#) is not a universal solution for all networking issues. For example, in high performance networks, where each millisecond in packet processing is precious, the benefits of [SDN](#) might not be worth its small drawbacks, as a vertically integrated device is theoretically faster [44].

3.2.1 Network innovation

New ideas in networking have had increasingly less impact in the real world. This phenomenon is commonly referred to as "Internet ossification" and it means that the Internet has become static, like a fossil [44]. Most researchers consider this to be the greatest challenge to networking in the last three decades.

Networks are dominated by proprietary equipment. The distributed nature of the Internet already introduces a great deal of complexity, but these components must also support an extensive set of features while maintaining network flexibility, dynamicity, performance, and efficiency [1]. This results in extremely complex network systems that require a tremendous amount of effort from network engineers while being extremely dispendious.

This status quo causes new networking features to require a long and arduous process with immense implementation, experimentation, and deployment problems. Current networking devices and their maintenance also doesn't come cheap, with long monetary returns for manufacturers in investment cycles. These facts hamper innovation by inducing network developers to delay the introduction of new network features until they are widely requested [48] [53].

Another important factor is the reluctance to test new ideas for fear of disrupting existing traffic. The internet has a massive deployment base and has long been considered as critical infrastructure [44], so network ideas are not tested for fear of disrupting current services. Consequently, experimentation scenarios are rarely possible and conducted in simplistic scenarios, which diminishes confidence in their results [45]. Realistic testing scenarios are necessary to gain the confidence needed for widespread acceptance [54].

The perfect example of internet ossification can be seen in the transition from IPv4 to IPv6. This transition started more than two decades ago and still remains incomplete, while being a simple protocol update [53].

Network architecture should encourage network innovation, rather than stand in its way. Use cases for the internet in the form of internet applications are continuing to transform and evolve and it is impossible to predict future application requirements and perfectly design future networks to meet those requirements. Instead, networks should evolve and adapt in the face of surging problems [45].

The primary motivation behind programmable networks has been to enable innovation and minimize the barrier of adoption to new network services [1], making this an integral goal of [SDN](#). As a methodology, [SDN](#) provides a great platform for development with new ideas, techniques, and designs, encouraging experimentation.

The separation of the control plane from the data plane abstracts the definition of network policies from the underlying implementation in the switch hardware. This abstraction separates network problems into smaller traceable pieces, which enables innovation of individual components [53]. As an example, technological developments in the control plane can be made independently of the data plane, enabling faster technological development in both planes [49]. [SDN](#) is also based on open standards for communication between planes, which further stands to enable greater innovation through interoperability.

Through network programmability, [SDN](#) enables the development of new ideas using these hardware abstractions, which allows the programmer to avoid dealing with low-level hardware details. This enhances convenience and flexibility, which in turn makes network application programming more accessible to a wider audience.

Lastly, the programmable controller also benefits from a network-wide view and control, meaning it can instantly and seamlessly deploy new network applications, such as protocols or policies, across the entire network with minimal network downtime. This ability to easily update devices enables devices to be incrementally tested and refined as feedback is received. In traditional networks, these processes required the installation of expensive hardware and manual updates to individual devices [45].

3.2.2 Network cost

Routing devices are extremely costly to purchase and maintain, especially given a handful of companies have tight control over this market. The complexity of the devices prevents new competitors from entering, which is the main reason for this dominance. [SDN](#) promotes a more affordable alternative to these traditional devices through its open-source approach, which allows any compatible hardware device to be used as a networking device, commoditizing networking hardware [44].

In an attempt to address the lack of functionality in current network devices, traditional networks are filled with a myriad of specialized components and middleboxes performing a variety of different functions [1] [53], further driving up network costs. Examples of these middleboxes include firewalls, load balancers, network access controllers, and network address translators [44]. In [SDN](#), these devices are replaced by software programs which are implemented by the network controller as network-wide policies, thus reducing network costs.

[SDN](#) can also reduce costs by optimizing energy consumption. In large-scale network scenarios, such as data centers, the cost of energy consumption from networking equipment may seem small, but it still

constitutes a significant expense. **SDN** allows networks to dynamically power on and off devices based on expected network traffic, reducing energy consumption and costs [44].

3.2.3 Network control

Computer networks are often too large, complex, and heterogeneous for conventional methods of configuration, optimization, and troubleshooting to be effective [45] [53]. Network management becomes very difficult due to this complexity, which in turn hinders efforts to optimize network performance.

Network operators are required to configure their networks for a wide range of applications through manual translation of desired high-level policies into low-level commands, while coping with changing network conditions. This task is further complicated by vendor-specific configuration interfaces, which are often complex, limited, and can vary from model to model, even within the same vendor [1].

Beyond being extremely tedious, manual network configuration is extremely error-prone. Configuration changes can therefore cause network instability, leading to outages, security flaws and performance disruptions [44]. Additionally, configuration errors require a significant amount of effort to troubleshoot.

Efforts to improve network performance are also severely limited, with current optimization efforts typically confined to small regions of the network or specific network services, efforts that are further constrained by the lack of global network information [45].

Another major motivation for **SDN** is to shift network control from vendors to operators, and ultimately from operators to users [47]. In today's Internet, the aforementioned complexity of networks impedes this goal so **SDN** is designed to make networks more flexible, open, and simple to configure.

Networks are made simpler first by breaking down networking into smaller pieces through the abstraction gained from the decoupled control and data planes. This allows network managers to implement the desired network outcome in the control plane without having to deal with complex lower-level instructions.

Secondly, the complete and centralized view of the entire network helps simplify network management [48], given that current network policies are extremely complex due to the need to cope with the distributed nature of the Internet. **SDN** enables a more consistent application of network policies throughout the entire network, rather than on individual devices, making policies more uniform and easier to implement [44].

Finally, **SDN** networks are also designed to be highly programmable, granting unprecedented control that can leverage the aforementioned benefits to enable simple, flexible and agile management of network resources. This in turn simplifies the design, deployment, and management of complex network environments and allows for more effective use of network monitoring through the ability for the entire network to homogeneously adapt to any necessary reconfiguration.

These evolutions in network structure enable new powerful solutions to classical networking problems such as end-to-end congestion control, energy efficient operations, load balanced packet routing, quality of service support, and data traffic scheduling [45]. In contrast to traditional networks, where these would be implemented by deploying specialized middleboxes at precise locations in the network, **SDN** allows them

to be developed in software as network applications. This capability is what makes physical middleboxes obsolete, as noted above, and contributes to simplifying the deployment of new network services.

Network applications can be tailored for their purpose, ensuring that hardware resources are not wasted on irrelevant functions [48]. They can be easily shared and deployed as third-party “apps” on any network [44] [53], making more efficient implementations readily available and thus improving network performance.

Implementing automated network management is another major [SDN](#) use case [45] [55]. Intelligent network control can reduce reliance on network administrators by enabling network applications to quickly and easily reconfigure the network to meet changing requirements and make real-time adjustments to automatically optimize network performance. [SDN](#) enables network performance and traffic flow monitoring, enabling automatic and dynamic network configuration with centralized validation that automatically optimizes network performance based on desired high-level network policies.

[SDN](#) enables networks to be more application-aware by allowing applications to communicate with the controller and request network services [49]. An example of this is transport with quality of service insurances.

3.3 [SDN](#) architecture

The [SDN](#) architecture divides networks into abstracted layers, inspired by computer systems, each with its own unique purpose [53]. It consists of three main layers that are illustrated in Figure 19: the data plane, the control plane, and the application plane. These layers are linked by the Southbound API, and the Northbound API. Variations of this [SDN](#) view are numerous and will be explained when relevant.

A fundamental design principle of [SDN](#), which is a cornerstone of the architecture, is to emphasize standard, open interfaces between the different planes. The purpose of this structure is to promote compatibility and interoperability between different data, control and application plane projects so that all components can work together interchangeably [53]. A result of this open approach to development is that [SDN](#) relies on the aforementioned open APIs and a collection of software components that support those APIs, rather than any specific protocol stack [47].

When surveying the [SDN](#) architecture, it is best to approach it and each of its components from the bottom up. This section will therefore cover all possible layers and components in [SDN](#), starting from the data plane and ending on the application plane. Specific technologies will also be mentioned as they represent the current state of the art. However, it is important to keep in mind that there exist several alternatives, and these technologies are constantly evolving, with different implementation approaches that may become more prevalent in the future.

3.3.1 Data plane

The data plane is the first and the lowest layer of the [SDN](#) architecture and is made up of networking infrastructure which in traditional networks are switches, routers and any other middlebox appliance.

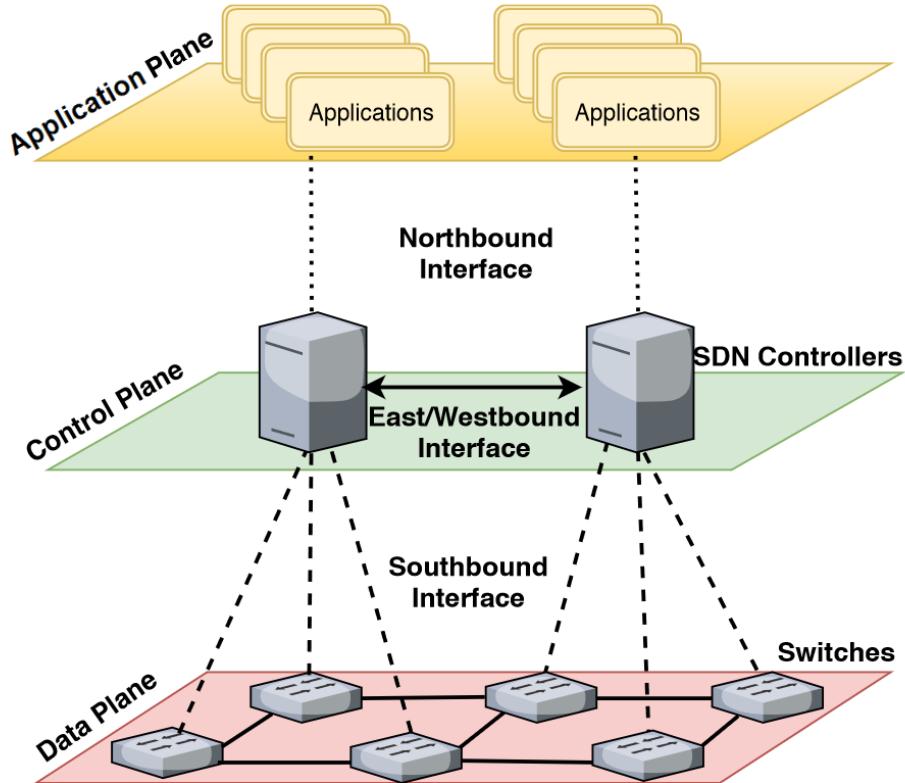


Figure 19: Layered view of SDN Architecture [56]

These distributed forwarding devices are interconnected with each other through traditional means like wireless radio channels or wired cables. Unlike in traditional networks, however, these devices don't have the autonomy to make networking decisions and instead follow the instructions given by the controller [53].

This last property means that in SDN, all network equipment is referred to as "switches" or "bare metal switches" [47] because without embedded intelligence in them, forwarding hardware becomes interchangeable and different names for equipment become redundant. On top of the hardware, SDN switches also have software running on every individual device. This is distinct from the control plane software and its purpose is to abstract the switch hardware, communicate with the control plane and implement forwarding decisions.

The data plane is responsible for receiving, processing and forwarding packets based on rules set by the controller, implementing the packet processing policies defined by the control plane [54] [45]. These devices may also gather and store network status information, including network topology, traffic statistics, and network utilization, to later transmit to upper layers [45]. This data can then be used to verify network behavior, monitor network performance and handle failures, all of which improve overall network health [48].

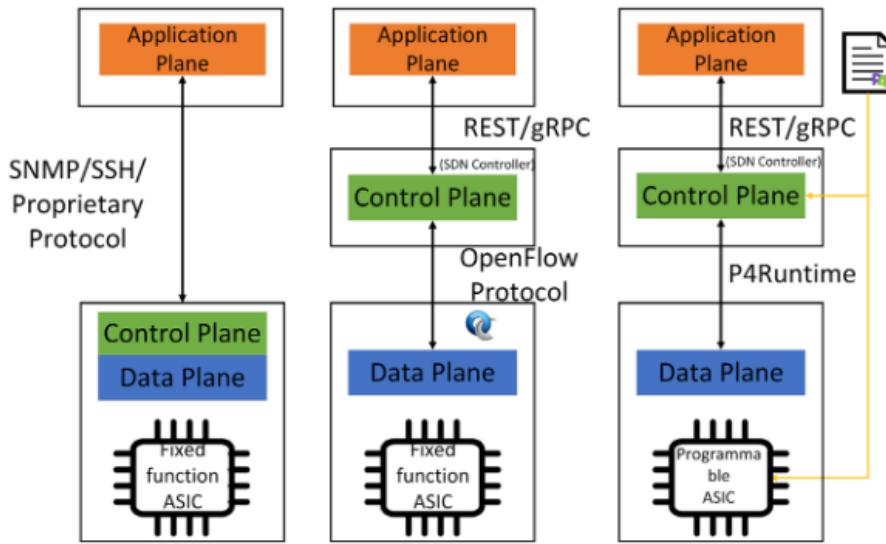


Figure 20: Evolution of SDN [51]

3.3.1.1 Data plane programmability

SDN switches can be divided into two categories based on whether they are programmable or static [47]. This classification is driven by changes not only in behavior, but also in hardware. Static devices use fixed-function chips that process headers based on fixed protocols that cannot be changed. Programmable devices, on the other hand, use programmable chips for dynamic processing.

The programming of data plane devices is a relatively recent development, which has led to the two phases of SDN mentioned in section 1. Figure 20 depicts these two phases, accompanied by a comparison to traditional networks, for a better understanding of this transition.

These efforts ultimately allow for the behavior of the data plane to be modified. Data plane programmability refers to the ability of a network operator to systematically and extensively reconfigure the processing logic of a switch as desired [48] [57]. Simply put, data plane programming enables users to implement their own data plane algorithms on any switch.

Most importantly, data plane programming can unlock complete control over network packet processing. This allows customization of network protocols, meaning they can be added, modified, or removed by the programmer. In traditional networks, this capability can only be achieved by a manufacturer [57].

The main benefit of data plane programmability is therefore its increased flexibility [48]. This feature allows for the optimization of packet processing functions to meet specific use cases or requirements, tailoring the data plane to the network's specific requirements. Additionally, it facilitates the deployment of brand new network protocols.

The concept of data plane programming emerged in response to an issue found with then standard SDN solutions. SDN's ambitious goals naturally lead to initial approaches falling short. The first phase of SDN is defined by the Southbound API OpenFlow. The soundbound API is an interoperable interface that is used by the control plane to communicate routing decisions to devices, dictating how the control plane

Version	Date	Header Fields
OF 1.0	Dec 2009	12 fields
OF 1.1	Feb 2011	15 fields
OF 1.2	Dec 2011	36 fields
OF 1.3	Jun 2012	40 fields
OF 1.4	Oct 2013	41 fields
OF 1.5	Dec 2014	44 fields
OF 1.6	Sep 2016	Restricted to ONF members

Table 1: Fields recognized by the OpenFlow standard [58]

communicates with switches [55]. There are several other Southbound APIs available, with OpenFlow being the most popular in this first phase. OpenFlow is a protocol-dependent API that defines a specific set of headers for controllers to use in defining rules in the data plane. Openflow was standardized by [ONF](#), an organization established with the objective of advancing [SDN](#).

As OpenFlow evolved, researchers realized that its static nature was extremely limiting [58] [57] [55]. Protocol-dependent Southbound APIs assume a specific data plane functionality that cannot be altered. When working with OpenFlow, network policies must adhere to the existing matching fields and actions of OpenFlow, which undermines the controller's power over the protocols used on the network.

To work around its static nature, and as a necessity to support many different protocols, the number of actions and matching fields increased with each version of OpenFlow, which has a price in complexity. The first version of OpenFlow standardized 12 matching fields and 10 actions. As new versions were released, this number skyrocketed to 44 matching fields and 19 actions in version 1.5 [55] [58], as can be seen on Table 1. This remedy can be thought of as a band-aid solution, as it increases complexity without solving the underlying problem of inflexibility.

Another important issue to consider is the reactive nature of definition of these headers. OpenFlow solely defines new headers when they become widely requested. Therefore, any new useful header fields must wait for standardization in a new version of OpenFlow before they can be integrated into networks. A proactive and open approach would be more appropriate, as this approach contradicts the [SDN](#) principle of enabling innovation [55].

OpenFlow attempted to fix this problem in version 1.2 through an OpenFlow Extensible Match, which allows for adding custom matching fields [53]. However, the fundamental problem remained and a brand new approach was preferable. Many different technologies aimed to further enhance data plane programmability emerged. This dissertation will focus on the most popular, [P4](#). As final observation, the protocol-dependent nature of OpenFlow also limits the protocols that can be used to standardized ones, meaning the defined matching fields of flow tables are limited to traditional network protocols.

3.3.1.2 [P4](#)

[Programming Protocol-independent Packet Processors](#) is a high-level programming language designed to specify switch packet processing pipelines. Similarly to other data plane programming efforts, it aims to

generalize the OpenFlow match-action framework and it has emerged as the most popular technology in this race, with strong support from both industry and academia [57].

The P4 project began with three main objectives [58]. The first was to enable device reconfigurability, so that the way a switch processes packets could be easily modified even after deployment. The second objective was to promote protocol independence by allowing switches to run any arbitrary protocol. Lastly, to achieve the last two goals while abstracting the underlying hardware, so that P4 code can run anywhere.

P4 achieves these goals by a new revolutionary process that, rather than keeping track of the vast number of standardized protocols used, defines an open, high-level and flexible approach for defining custom packet parsing and header matching mechanisms. This vision argues that instead of having to repeatedly extend the OpenFlow specification, hardware manufacturers should support flexible mechanisms for parsing packets and matching header fields so that controllers can leverage these capabilities through a common, open interface [58].

P4 can actually be used with both programmable and fixed-function devices [47] and it works in two distinct phases [58] [48].

During the initial phase, operators utilize the P4 language to define the parser and deparser, determine the order of match+action stages, and specify the header fields processed by each stage. For programmable devices, P4 prescribes the desired behavior in the pipeline, while in fixed-function devices, P4 merely describes the chip.

The second step occurs at runtime and consists of populating the configured pipeline with rules. This process is device-independent, with the control plane adding or removing entries in the match+action tables of the switch pipeline using the Southbound API. This phase determines the usage of defined protocols and the application of policies at a particular moment, such as adding entries to flow tables and associating actions with flows.

3.3.1.3 Hardware

In the creation of SDN infrastructure, any available hardware can be considered for use in the switches, with each having its own set of advantages and disadvantages. Following the approach of [45], hardware can be classified into three categories: general purpose computers, vendor specific hardware, and open networking hardware.

1. Open networking hardware: Open networking hardware refers to switch specific hardware that is based on open standards and can be used with a variety of software. This provides a programmable and vendor-independent solution for building networks. The most appropriate analogy is to compare a switch to a Personal Computer (PC) built from commodity, off-the-shelf components. This allows for the assembly of a high-performance switch [47] using open-source software. For example, the Open Compute Project takes advantage of existing commodity switching components and offers full architectural specification for switches online [59].

2. Vendor specific hardware: Continuing the analogy with the [PC](#) world, in that ecosystem pre-built computers are available for purchase from several vendors, such as Dell or HP. Similarly, bare-metal switch vendors like EdgeCore and Delta offer pre-built switches [47]. This hardware is specifically designed and manufactured by a particular vendor for networking purposes in [SDN](#) and uses proprietary software. Moreover, it is of interest to note that an initiative designated Indigo [60] has been established with the objective of facilitating the integration of [SDN](#) capabilities into conventional hardware.
3. General purpose computers: These are regular computers, not specifically designed for networking, that are used as a platform to implement the data plane by running a software switch that performs packet forwarding functions. Hardware virtualization is used to run virtual switches on standard operating systems running on regular hardware. The widespread availability of general-purpose hardware makes it inexpensive, and easy to reuse for different tasks. However, this hardware is not only not optimized for networking tasks, but also relies on computer network interface cards, which typically have a limited number of ports and slower speeds.

Both open network and vendor-specific hardware exhibit similar high performance, while using general hardware can result in decreased performance as it is not specialized for networking. General purpose hardware is usually the least expensive, followed by open network hardware, with vendor-specific devices being the most expensive. This price difference can be attributed to the scalability characteristics, ease of use, and performance of each type of solution. All and all, each approach has its benefits and drawbacks, so the best choice should be based on the specific needs and constraints of the target network.

Figure 21 presents a high-level schematic of a bare-metal switch. This abstraction better represents the open networking approach because Vendor-specific solutions are proprietary and therefore harder to generalize and the general hardware only has a [Central processing unit \(CPU\)](#), virtualizing to operate in a similar manner to the depicted.

The diagram depicts three main components. The [Network Processing Unit \(NPU\)](#), represented as a combination of [Static Random-access memory \(DRAM\)](#)-based memory and an application-specific integrated circuit based forwarding pipeline, is a chip optimized for parsing packet headers and making forwarding decisions. This representation ensures packets are buffered in memory while being processed in the pipeline. A general-purpose processor is connected to and controls the [NPU](#). This chip runs the software of the switch and deals with the off-switch communications with the control plane. The last component is the device ports, which interconnects devices.

To achieve its goal of interoperability between different hardware, [P4](#) programs must be portable across different devices without requiring modification. Achieving this is a problem, however, because different vendors implement different physical pipelines. To address this issue, [P4](#) introduces the concept of an intermediate layer between the core [P4](#) language and the targets [57]. This layer abstracts the details of the physical pipeline into a common logical pipeline that provides a programming model that easily describes how packets are processed. As a result, [P4](#) programs can be developed for a logical architecture that can be deployed on multiple targets. Currently, there is no consensus on a standard

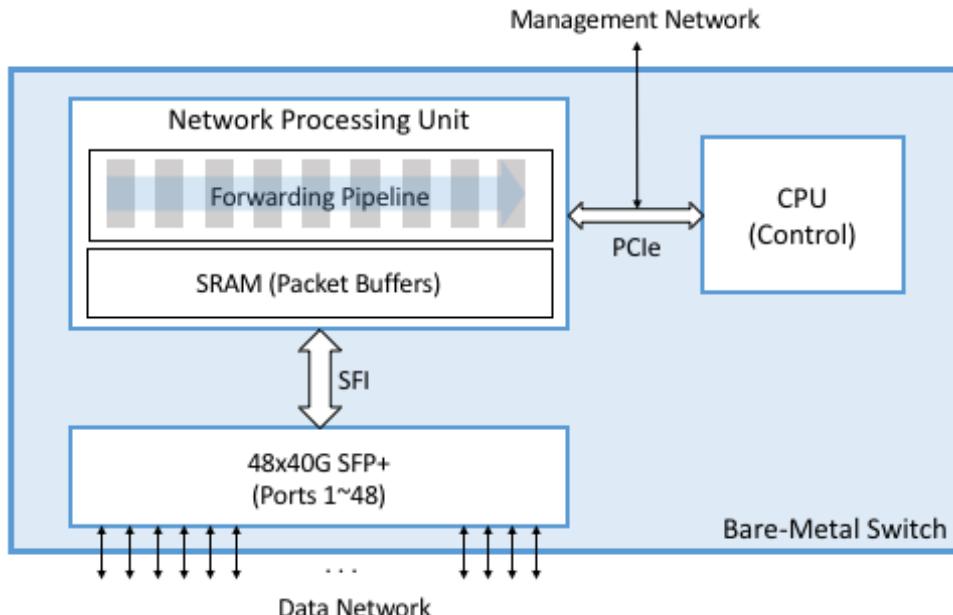


Figure 21: High-Level schematic of a bare-metal switch [47]

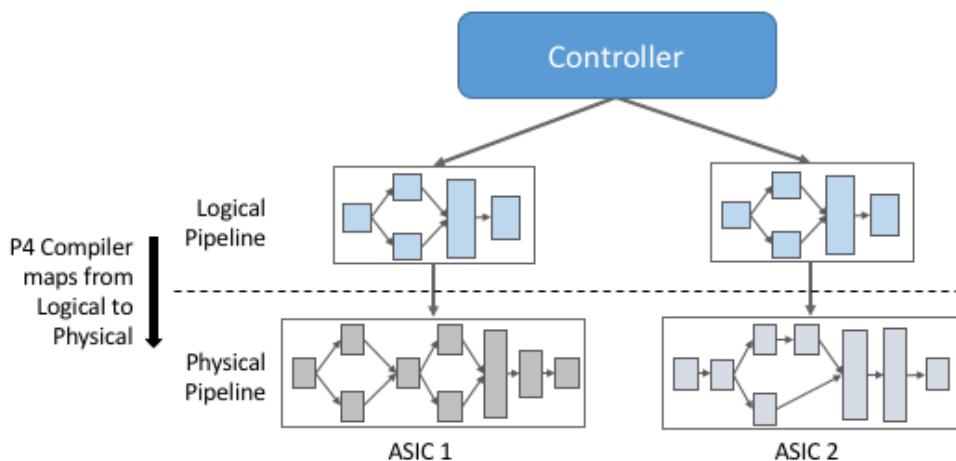


Figure 22: Definition of a logical pipeline to provide a pipeline-agnostic view the control plane [47]

logical pipeline for all devices. Nevertheless, this approach enables pipeline-agnostic controllers as can be seen on Figure 22.

The logical pipelines can be represented using different abstractions. In Figure 22, the pipeline is described using the data match-action pipeline. This abstraction describes the packet processing functions of a network device as a pipeline of lookup tables. Each table is responsible for matching specific fields in the packet header and performing an action based on the result [48] [53]. Thus, the data match-action abstraction describes the **NPU** as a sequence of lookup tables.

Other abstractions exist for describing the switch pipeline, but this is the most popular by far. OpenFlow is largely credited with popularizing the match-action abstraction [48]. However, the original proposal for OpenFlow took advantage of the fact that most routers and switches already contained flow tables used to implement additional network functions such as firewalls, network address translation, quality of service, and statistics collection, to define a common set of functions and provided an open protocol for

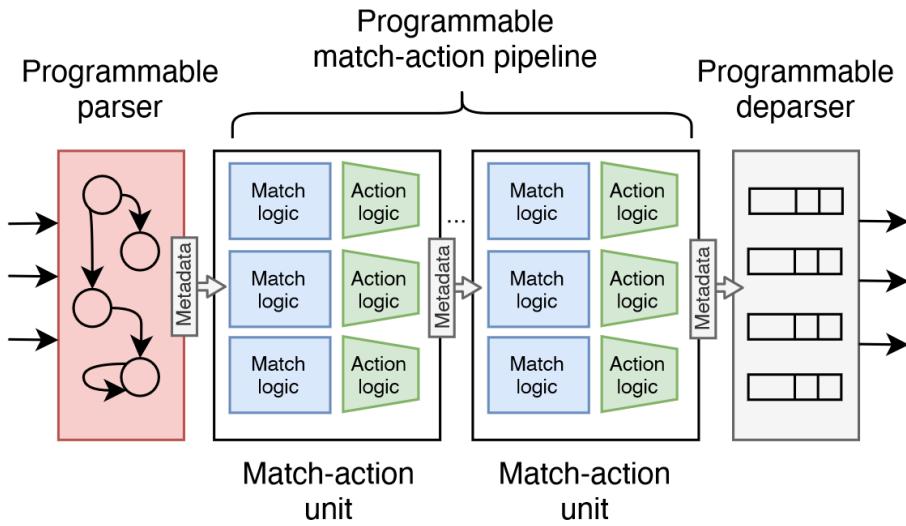


Figure 23: Protocol-Independent Switch Architecture [57]

programming these flow tables [54]. This implies that the architecture of the [NPU](#) and the [Southbound API](#) strongly influenced each other.

This abstraction depicts multiple tables because high speed switches require a multi-stage pipeline for packet processing. It is the optimal approach for packet processing because this process involves checking multiple header fields, and with multiple stages, different stages can perform different tasks simultaneously [47].

There are various [P4](#) architectures, including [Protocol Independent Switching Architecture \(PISA\)](#), [Portable Switch Architecture \(PSA\)](#), V1Model, SimpleSumeArchitecture, and [Tofino Native Architecture \(TNA\)](#) [47] [57]. These architectures abstract the target hardware into a forwarding model representation.

[PISA](#) presents a simple representation of the device, so all of these examples can be considered variations of it. The other architectural models follow the same principles with some arbitrary enhancements. The [PISA](#) model will be discussed as it is the most general. However, it should be noted that the V1Model is currently the most popularly used.

[Protocol Independent Switching Architecture](#) is a data plane programming model of [P4](#). It was introduced in [P4₁₄](#) and is currently deprecated. It is applied in a variety of ways to the creation of concrete architectures [57]. As an example, Figure 23 gives a high-level overview of this pipeline.

This architecture depicts three main components. The parser is the first component and handles header fields by specifying their size, order, and other relevant details. This ensures accurate extraction and utilization of the fields in the subsequent stage. The next stage is the match action pipeline, which consists of the multiple sequential tables that define this type of data plane pipeline abstraction. It maps a desired action to a given match by utilizing previously defined headers. The third is the deparser and it does the opposite of the first. It deparses the packet metadata into the packet before it is transmitted on the output link. Beyond the packet, relevant metadata about it is also traversing the pipeline. Examples of this may be packet relevant, like input port and arrival timestamp, or device relevant switch counters, queue depth.

In practice, when a packet arrives at a switch, after being parsed, the packet headers are compared to the entries in the flow table. If a match is found, the corresponding action is taken, and the packet is processed accordingly. If no match is found, the packet is typically forwarded to the controller for further processing or dropped.

3.3.1.4 Software

Each switch must be equipped with a [Basic Input/Output System \(BIOS\)](#), which is similar to the firmware found in microprocessors as it provisions and boots a bare-metal switch. The Open Network Install Environment has emerged as the standard for switch [BIOSes](#) [47]. On top of it, every switch must run a local Switch [Operating System \(OS\)](#), which is responsible for handling calls from the controller and taking appropriate actions on the switch's internal resources. The most widely used Switch [OS](#) is Open Network Linux [61], an open-source project based largely on the Debian distribution of Linux and extended to support hardware unique to switches [47]. Other open-source Switch [OSs](#) include Stratum [62] and SONiC [63]. Vendor-provided software development kits for the on-board switching chips are also important. These kits are analogous to device drivers used in traditional operating systems.

3.3.2 Southbound API

The Southbound [API](#) is the second layer in [SDN](#). As introduced in the previous section, the Southbound [API](#) is a logical interface that connects the controller with all network elements in the data plane [49]. It allows the control plane to push configurations to forwarding elements and the switches to push information to [SDN](#) applications, enabling the controller to have direct control and constant information on the state of the data plane elements. This protocol describes the capabilities of the switches, which defines the communication protocol and methods used to exchange data between the data plane and the controller. The interaction between the control plane and the data plane is therefore defined by this protocol [53].

A well-defined programming interface is crucial for the success of [SDN](#), and its standardization offers several advantages for designing and deploying efficient [SDN](#) solutions [56] [53]. Specifically, it provides flexibility by allowing the control plane to communicate with the data plane through a variety of standardized channels. This enables network operators to choose from a range of controllers and switches, reducing vendor lock-in and expanding their options.

The responsibilities of this interface can be split between Control and Configuration [47] and usually, protocols only fulfill one of these categories. For a long time, OpenFlow was considered the state of the art for Southbound [APIs](#) and as testament to that is still the most widely used today [56]. However, the current state of the art for Southbound interfaces is a combination of P4Runtime, gNMI, and gNOI. P4Runtime and OpenFlow are control interfaces, while the gNBI combination of protocols is responsible for device configuration [47].

These different types of Southbound [APIs](#) also have the huge benefit of sheltering the controller and by extension the application running on top of it from the diversity of network devices that may exist in the network [47].

3.3.2.1 OpenFlow

The previous section already touched on the origins of OpenFlow as a means to justify the industry's distancing from it. McKeown et al. proposed OpenFlow [54] as a means to enable experimentation on campus networks. OpenFlow leveraged the flow tables present in most Ethernet devices, which are used for additional configuration such as firewall setup and subnetting, to allow customization of campus network devices [45].

OpenFlow has been the de-facto standard for the Southbound interface for a long time [44] and remains the most popular Southbound API. This is evidenced by the fact that most commercial switches support the OpenFlow API out-of-the-box [53].

3.3.2.2 P4Runtime

Nowadays, Openflow's development has been dropped by ONF in favor of P4, which means it has been replaced with the more open alternative of the P4Runtime. P4Runtime is a surging data plane API [57] that is closely coupled with the P4 programming language. P4 enables the data plane to be programmed according to network requirements. A static interface would eliminate the advantages of P4 since the controller would not be able to benefit from the custom programmable policies defined in the data plane.

To fully reap the benefits of P4, a dynamic communication interface between the controller and the switch is necessary [47]. To establish this interface, not only must the switch pipeline be known to the controller, but both the switch and the controller must be informed of the P4Runtime contract. P4Runtime is automatically generated along with the client and server-side stubs for the controller and devices, depending on the defined switch architecture [47].

3.3.2.3 gNXI

The term gNXI refers to the combination of gNMI and gNOI. The role of these two is somewhat blurred, but in general, gNMI handles the persistent state, whilst gNOI is responsible for clearing or setting the ephemeral state. In other words, gNMI is used for switch configuration, defining actions for retrieving or manipulating the state of a device via telemetry or configuration data, while gNOI is used for executing commands on a device, such as rebooting, pinging, and accessing other operational variables on the switch rather than directly changing its configuration [47] [64].

In traditional networks, the configuration part of the Southbound api has been called the operations, administration, and maintenance interface which is usually accessed in command-line [47]. The equivalent of gNXI in traditional networks is [Simple Network Management Protocol \(SNMP\)](#). SDN has brought about the possibility for a different approach. [SNMP](#) was originally designed for closed devices and primarily focuses on reading rather than writing onto devices, making it more similar in capabilities to the configuration part of gNXI [47].

The gNMI has three methods. The well-known get and set, similar to their counterparts in [SNMP](#), are used for reading and writing. The third is the subscribe method, which is used to request a stream of telemetry data [47].

OpenConfig [64] is responsible for the development of gNIXI. OpenConfig is an industry-wide standardization effort to drive the industry toward a common set of configuration models. It is a collaborative effort within the networking industry to move toward a more dynamic and programmable approach to multi-vendor network configuration and management. gNIXI is expected to become the standard management protocol due to strong industry support [47].

The underlying mechanism used by gNOI and by gNMI is exactly the same [47]. Actually, P4Runtime and gNIXI all rely on the gRPC framework with protobuf [57]. The intricacies of protobufs and gRPC will not be detailed here, but these technologies are used for the benefit of not having to deal with the daunting list of formatting, reliability, backward compatibility, and security issues that other protocols, including OpenFlow, must deal with [47].

3.3.3 Control plane

Having outlined the structure of the data plane and its major technological developments, it is now time to move from the individual switch level to the controller network view. The control plane is widely considered the most important component of the SDN architecture [45] because it represents the centralized decision-making authority of the network, acting as the “network brain” [53].

The control plane consists of one or more controllers that oversee, manage, and manipulate the data plane through the Southbound API. All forwarding decisions are centralized in the control plane, which is dependent on a global view of the network to make accurate decisions. In this process, device-specific details are abstracted from the control plane [56].

Traditional operating systems provide abstractions of the underlying hardware that simplify the use and management of the underlying resources. As a result, it is possible today to build more complex applications faster, more securely, and more efficiently [53].

In various ways, the SDN layers bear deep resemblance to a regular computer system [49]. The application layer is comparable to the software applications that perform tasks using computing resources, the infrastructure layer is equal to the computer hardware and, most relevant to this case, the computer OS corresponds to the control plane.

For this reason, the controller is commonly referred to as the [Network Operating System \(NOS\)](#). This analogy introduces the same abstractions found in computer operating systems to networks [53], and is often used to better illustrate that the control plane is responsible for representing the network as a single system for network applications [44]. The terms [NOS](#) and controller can and will be used interchangeably from now on.

3.3.3.1 Controller responsibilities

A list of common applications implemented in the [NOS](#) can be collected from different SDN deployments and asserted as the base network service functions [53]. While such a list cannot be considered de facto obligatory due to the SDN root of modularity, it is still relevant to have a comprehensive list for familiarity’s sake.

The controller is mainly responsible for topology management, enforcing the desired packet processing policies and performing other necessary functions to manage and regulate the network [47] [56] [48] [53] [65].

Topology management requires maintaining an up-to-date network view, which involves link discovery and host tracking. The main duty of the controller is to implement the desired packet processing policies. This task involves determining where packets should be forwarded by establishing traffic paths and defining packet flow rules. Furthermore, the controller selects the headers that require modification and installs all of these control commands on each forwarding device.

Lastly, the controller is responsible for performing other necessary functions to manage and regulate the network, such as monitoring device health by collecting status information and other important data from switches and performing maintenance procedures on those same switches.

3.3.3.2 Control centralization

Due to its prominent role, the performance of the controller largely influences overall network performance. The biggest criticism of SDN stems from the belief that centralizing the control plane will lead to scalability and performance issues. However, most practical results suggest that a single controller is capable of handling a substantial number of new flow requests, and therefore should be capable of managing most small to medium sized networks [44] [65].

In medium to high network scenarios, however, the controller can quickly become overwhelmed. A single physical controller represents a performance bottleneck in the network with added security issues [65].

The architecture's logical centralized controller does not imply the existence of a single physical controller. For optimal performance, scalability, and availability, the network controller must be physically distributed [53]. A physically distributed but logically centralized control plane reduces lookup overhead and increases reliability while maintaining a centralized view of the network for applications to write in [44].

The physical distribution of controllers can be classified into flat or hierarchical [65]. In the flat architecture, all controllers maintain a global view of the network, whilst in the hierarchical architecture the network is divided into domains, where a single controller has jurisdiction over a single domain. To achieve a global view, a root controller is created to coordinate the actions of the other controllers and maintain the global view.

Another important advantage of distributed architecture is that a network controller can function similarly to other services by scaling based on workload to ensure availability and saved resources [47].

3.3.3.3 East/Westbound API

Some researchers propose the concept of an API specifically used for the coordination of these distributed scenarios [53] [45] [49] [56] [65]. This API is called the East/Westbound API and it provides the exchange of information between different instances of the controller in order help maintain a consistent network view. This API has no relevant standards. It is not typically considered a part of the overall SDN architecture, although in practice, something similar must always exist in distributed controller scenarios, as can

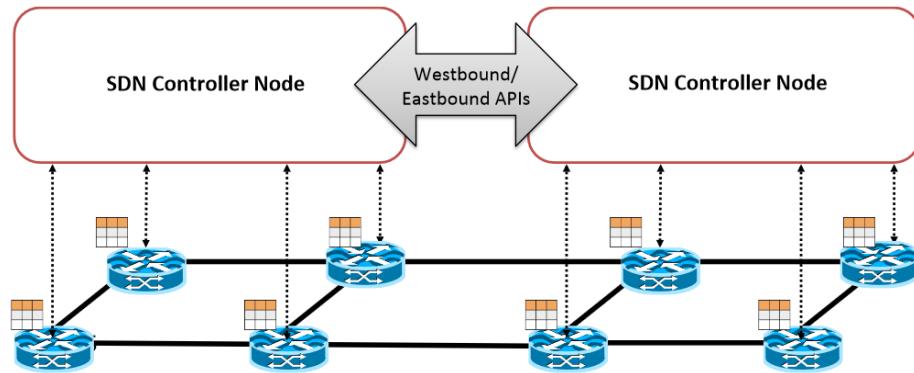


Figure 24: East/Westbound interface between distributed controllers [53]

be seen in Figure 24.

A finer distinction can be made between the Eastbound and Westbound APIs [53]. The Eastbound API mostly refers to communication between SDN controllers and legacy routers, while the Westbound API refers to communication between different SDN controllers, whether they are different frameworks or the same in a distributed architecture.

3.3.3.4 Notable Controller Platforms

Most controllers are licensed as open-source, with few having a proprietary license. They are programmed in a large variety of programming languages, with Java and Python being the most used, and some even contain more than one language for better performance [65]. However, most don't receive frequent updates and maintenance, and lack proper documentation [65]. A lot of different controller platforms exist, and the most notable are [Pythonic Network Operating System \(POX\)](#), Floodlight, [Open Day Light \(ODL\)](#), [Open Network Operating System \(ONOS\)](#), and Ryu.

All controllers can run on any commodity hardware [53]. They are designed to be highly modular, with the controller being customized and configured to the given deployment, only having the subset of required modules [47]. This means that the behavior of most controllers is generally the same, with the only relevant differences being the deployment architecture. Another important point where controller architecture varies is the support for different North, East, West and Southbound APIs, as not all controllers are compatible with all interfaces.

3.3.4 Northbound API

The Northbound API is a logical interface that connects the control and application plane. It is offered to the application developers by the NOS [53] and provides programming abstractions for networks, acting as a bridge between the control and application plane [56].

There is no fixed or widely accepted standard for Northbound APIs [53] [56]. The reason for the lack of standardization of this interface is the huge variation of applications and user requirements, that makes the existence of a single interface for all of them very difficult.

There is therefore a large amount of different APIs that can be used depending on the chosen NOS. The purpose of this large number of different APIs is to sufficiently provide control over all underlying functionalities to the application plane [47].

As a Northbound interface implementation, programmers and most controllers typically use the representational state transfer API [56]. The gNXI interfaces are also popular, with them corresponding exactly to their counterparts in the Southbound API [47].

3.3.5 Management plane

The application layer is the upper-most layer, sitting above the controller. It consists of network applications designed to meet specific user requirements and, through the Northbound API, this layer has easy access to the capabilities of all the lower layers such as the global network view and constant status information, enabling software applications to instruct the network controller on how to implement network control and operation logic [45] [51].

SDN applications can accomplish a variety of goals, from adaptive routing and boundless roaming to enhancing network security and many more [45]. In essence, this layer can request custom control policies for routing, firewalls, load balancers, monitoring, and so on. The controller interprets these policies and determines the actions to be taken by the network devices. The switches then translate the network commands into concrete actions in the forwarding tables [53].

The existence of this layer is sometimes omitted, depending on the interpretation. It can be considered an integral part of the control plane because the two are closely related. The application can be viewed as running on top of the controller as well as on top of the controller [47].

SDVN

The [SDN](#) paradigm has been identified by some researchers and other experts as a powerful and promising tool for addressing some of the challenges posed by [VANETs](#) [66]. As previously discussed, [SDN](#) is a methodology rather than a specific technology, which allows it to be applied to [VANETs](#). This line of thought created a new networking paradigm that is primarily known as [SDVN](#), although the term [Software Defined Vehicular ad hoc Network \(SD-VANET\)](#) is also occasionally used.

The concept of [SDVN](#) was first introduced by Ku et al. in [67], and since then there have been many different proposals of how to implement the [SDN](#) principles in [VANETs](#) [68]. [SDVN](#) has recently received significant support from researchers, which has resulted in significant technical and architectural advances in [SDN](#)-enabled vehicular networks [69].

Similarly to traditional wired networks, non-[SDN](#) systems function reasonably well. Therefore, the implementation of [SDN](#) should not be seen as a solution to make [VANETs](#) work, but rather as a means to enhance them.

4.1 Benefits and Challenges

The implementation of the [SDN](#) paradigm in [VANETs](#) is driven by the desire to bring some of the improvements of [SDN](#) to [VANETs](#) to tackle the challenging requirements of [VANETs](#). This paradigm has the potential to introduce flexibility, programmability, and centralized control to vehicular networks [69]. This section reflects upon the benefits of [SDN](#) and the characteristics of [VANETs](#) to extrapolate the benefits and possible challenges that [SDVN](#) can offer today.

4.1.1 Mandated Interoperability

As previously stated in Section 2.6, [VANETs](#) require high road acceptance for their useful applications to be effective. This implies that all implemented devices must be able to communicate with each other, and to this end the [EU](#) has tried to standardize devices to promote interoperability between different vehicle manufacturers.

In **SDN** networks, operators have complete control over a static network, making it easy to conceal the inner workings of devices from the rest of the Internet and implement them as desired. However, in **SDVN** the same is not possible because a Vehicle **ITS** sub-system must be able to communicate with every single roadside **ITS** station in order to access infrastructure services deployed in the Central **ITS** sub-system. **SDVN** devices must be compliant with existing regulation and communicate with standardized **V2X** protocols in order to be interoperable with other devices.

Such an emphasis on interoperability undermines the powerful control that **SDN** brings. Network operators and developers can't take full advantage of the freedom this technology provides to implement new algorithms and test them in real-world scenarios. Even the effectiveness of the additional control provided by **P4** is reduced by the inability to change network protocols.

This is not to say that the increased control that **SDN** brings is completely useless, as it can still provide some advantages in dealing with certain scenarios, but this **VANET** characteristic largely limits what can be accomplished with **SDN**.

It is also relevant to note one of the major goals of **SDN**, which is to move network control from manufacturers to operators, with the ultimate goal of giving it to network users. In vehicular networks, the identity of the network operator becomes an implementation detail, but it will most likely be the same government or private entities deploying the infrastructure. **SDN** continues to provide these operators with greater network control, but it cannot be given to users for fear of being misused for nefarious purposes.

4.1.2 Network innovation

Another important conclusion to be drawn from the above is that the technologies that will form future **VANETs** can be expected to be even more static than in traditional wired networks. On top of that, current hardware defined devices exacerbate this problem, as any fundamental change to **VANETs** would require the hardware of thousands of roadside systems and millions of vehicle systems to be manually replaced. This presents a logistical nightmare, making it highly unlikely that this scenario would happen except in the rarest of cases, and then only if mandated by the **EU**. This reality is evidenced by the slow and arduous standardization process that **VANETs** have undergone over the last 20 years, which has sought to find a single solution for **VANETs** that would act as a definitive and permanent answer to all of its problems.

SDVN allows new functionalities to be implemented in the same hardware through software updates, facilitating changes to be made throughout the network. This could considerably accelerate the adoption of this technology and facilitate the implementation of any future fix or optimization.

4.1.3 Network management and performance improvements

Even more than in static networks, **SDN** can provide **VANETs** with more flexible and intelligent packet flows, channel allocation, and connectivity. Traditional **VANET** routing algorithms are limited to the individual perception of each node as they are implemented in each individual device. **SDN** is built on a global view to make better decisions based on the combined information from multiple sources, which results in better network performance since informed decisions lead to more optimal outcomes. The automated network

capabilities brought about by [SDN](#) also have the potential to address [VANET](#)'s adversities. Listed below are some of the characteristics of a [VANET](#) and how they can be enhanced through the implementation of [SDN](#).

High mobility In an environment where topology changes occur frequently, it is critical to be capable of adapting to sudden and unexpected events. [SDN](#) provides flexibility, which allows for more dynamic network configuration, enhancing the network's responsiveness to emergencies and changing requirements, while also enabling it to better adapt to changing conditions and needs [67].

Predictable mobility [SDVN](#) can leverage the centralized view of the network, to better exploit the predictable mobility when compared to traditional approaches. The controller can use a vehicle's information to predict future topology changes and set traffic rules accordingly, improving connectivity and overall network performance.

Congestion and scalability Network congestion and the variable vehicle density are issues that can be reduced by leveraging the global view of the network to better optimize the available mediums. Centralizing network control increases cooperation, allowing the optimal path to be calculated, which reduces delays and overhead [66]. Also, situations of high vehicle density can be more easily predicted and the transmission power levels on vehicles can be jointly adjusted based on future vehicle network density, greatly reducing interference [66]. Overall, automated tools promise to make run-time changes to networks based on the unique characteristics of [VANETs](#), helping to achieve a better performing network.

4.1.4 Issues with maintaining a global network view

The advantages of centralized control are plentiful, but the establishment of this centralized view in [VANET](#) is not guaranteed. [VANETs](#) are based on wireless connectivity, which causes the southbound [API](#) to suffer from instability and even complete unavailability [68]. Combined with volatile topologies, this results in increased communication delays in the southbound interface, which can lead to inaccuracies in the controller's view of the network and in the rules established in the switches, casting doubt on the controller's ability to maintain an accurate and up-to-date view of the global topology [70].

Another major concern introduced with [SDVN](#) is the management overhead introduced by the southbound [API](#) required for the controller to manipulate the devices, as switches must be frequently updated to accurately reflect the network's state. This could result in additional traffic and increased network congestion compared to the distributed methods used in traditional networks.

At the same time, it should be noted that it may not be necessary for all instances of the controller to maintain a global view of the entire network. Instead, individual instances of the controller may only need to keep track of a small region of the network, as most messages in [VANETs](#) are only relevant in their immediate surroundings [71].

In summary, the dynamic state of the network combined with poor controller connectivity presents a major concern in [SDVN](#). The process of maintaining an updated view of the global network topology

becomes costly and time-consuming due to potential inaccuracies in the updated information [70]. To address this issue, it is important to ensure robust protection against loss of connectivity and to guarantee the availability of the control plane, which sometimes requires the integration of other technologies such as fog computing with [SDN](#) [70].

4.1.5 Cheaper networks

When considering the deployment of [VANET](#) technology, the cost of infrastructure is a major concern, with up to 86% of the total expected cost being attributed to hardware [16]. Just as in traditional networks, [SDN](#) can reduce costs by reducing complexity and increasing modularity while promoting interoperability.

4.1.6 Network security

The introduction of the [SDN](#) paradigm does not introduce any new significant security or privacy challenge to [VANETs](#). Some researchers [70] mention the centralized controller as a new vulnerability that can be exploited, but in reality current [VANET](#) deployment already relies on services provided by the Central [ITS](#) sub-system.

4.2 SDVN architecture

The [SDVN](#) architecture can be represented based on either of the architectures of the two technologies that define it. Of the two, the most suitable is the architecture of [SDN](#). The [VANET](#) architecture is best represented by the [ITS](#) host architecture, which is built on the [OSI](#) model. The [SDN](#) paradigm aims to move the internet away from closed and complex standardized protocols and into open and simple network components. As such, the [SDVN](#) architecture inherits the format as the general [SDN](#) architecture, retaining the three layers present in it as can be seen in [Figura 25](#).

4.2.1 Data plane

In [SDVN](#), as in [SDN](#), the data plane consists of all forwarding hardware but these components can be divided into two different categories based on whether they are wired or wireless. Mobile hardware refers to the vehicle [ITS](#) station, while stationary hardware refers to the routers and switches found in the remaining [ITS](#) sub-system. Network intelligence is extracted from all [ITS](#) stations, but due to the special characteristics of the mobile nodes all [SDN](#) enabled devices cannot be treated the same.

4.2.2 Control plane

The control plane remains almost unchanged as the centralized logical intelligence of the network, managing device behavior to achieve desired policies. It remains responsible for communicating the high-level applications defined in the application plane to the devices in the data plane, and for transmitting device

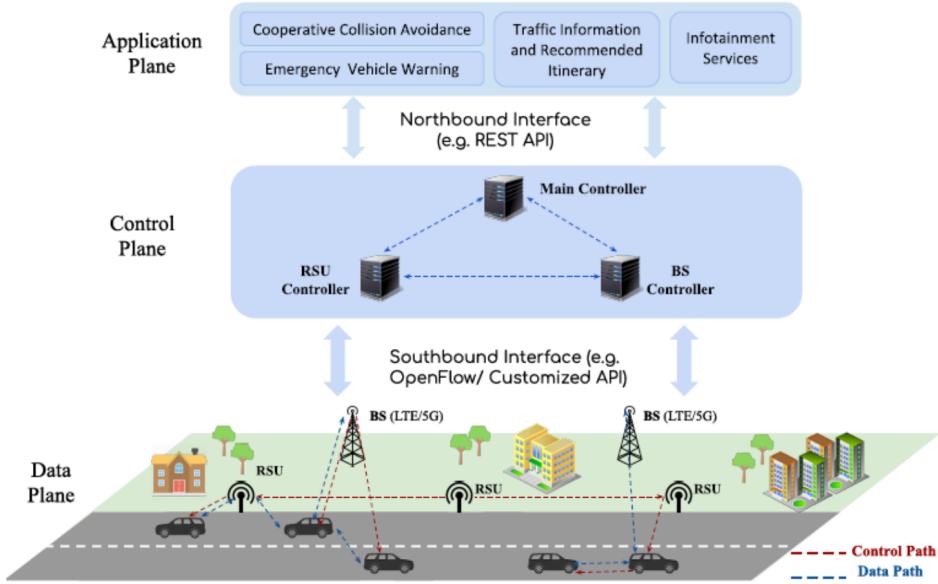


Figure 25: Conceptual view of SDVN architecture [72]

information from the data plane to the application plane. The main difference introduced by SDVN is a consequence of the need to manage devices in the ad hoc space. Due to the two distinct types of components in the data plane, unique and customized networking applications are required to address the unique nature of each scenario. Managing the wired portion of the network can be considered a simple task because the process is the same as SDN, whilst the wireless domain presents unique challenges and requires distinct solutions to manage wireless resources as effectively as wired resources [68].

In SDVN controller distribution becomes a central concern. Vehicular networks have immense coverage given that they include all roads, and due to this vast distribution, a single controller cannot handle the traffic of the entire network. Therefore, it is out of the question for scalability and delay reasons [73]. Instead, the control plane must consist of multiple physical controllers efficiently distributed across the network, coordinating their decision-making efforts to handle the high volume of traffic [70]. Another important motivation for distributing the control plane is to bring the controller closer to the vehicles to reduce communication delays between the data and control planes [74].

It should be noted that a truly distributed controller, where all instances have global information of the network, seems impractical and even counterproductive. To perform effectively, controllers usually only need information about a specific part of the network, which is in part because some messages are only relevant in the context of that particular section of the network [68] and the ones that aren't should be directed to the nearest roadside ITS station. The size and the transition between regions is dependent on implementation.

That being said, there exist many different controller placement distributions proposed in the literature. In practice, the controller can be located in the Central, Roadside, or Vehicle sub-systems. Different architectures place the controllers in a combination of these stations, giving them distinct responsibilities and utilizing different methods to distribute intelligence between them.

Most approaches to controller distribution in the literature [69] [68] [73] converge on a two-tier hierarchical controller. The top-level orchestrator controller, also known as the global or primary controller, is located in a server integrated into the Central ITS sub-system. The lower-level controllers, referred to as local or secondary controller instances, are placed closer to vehicles in the Roadside ITS sub-system in an effort to improve connectivity and reduce delays.

4.2.3 Application plane

The application plane remains unchanged. It consists of network applications designed to fulfill certain networking tasks using the capabilities given by the control plane. Some of these include monitoring, QoS, analytics, recovery, security, routing, load balancing, and management [69]. As stated in Section 2.6, VANET applications can be divided into two main categories, each with different QoS requirements. The most significant of these, the safety applications, are mostly delay sensitive, while the non-safety applications are more bandwidth intensive [66].

4.2.4 Communication interfaces

All communication APIs have no relevant changes. The only noteworthy detail is that by having access to both ITS-G5 and C-V2X for communication and control [73], a controller has two main communication mediums through which, using the southbound API, it can communicate with SDN devices.

4.3 Related Work

Emulation is a convenient and cost-effective method for testing VANET solutions. However, due to the complex and unpredictable nature of real-life environments, relying solely on these tools is insufficient [68]. In order to gain a realistic perspective on VANET, it is necessary to implement it in real-world scenarios so the validity of the expected benefits and drawbacks can be tested. The same is true for SDVN. SDVN is notable for the lack of experimental efforts conducted on actual hardware. In order to enable testing on real scenarios, open-source SDVN tools and frameworks that are compatible with a wide range of hardware are indispensable [68]. This section reviews several related papers that aim to implement VANETs and SDVN in physical hardware using open sourced components, outlining the key details of their approach and architecture.

Raviglione et al. [75] present a demo paper that assembles an open-source platform based on PC Engines' boards and Unex's Wireless Network Interface Controllers (WNICs). The purpose of this paper is to create a testbed for testing applications that communicate in the vehicular environment. Even though this paper does not attempt to implement SDN principles, it is relevant because it provides all the necessary hardware and software components to implement a vehicular testbed. The authors assembled two boards consisting of the embedded PC Engines APU1D board with an AMD G-series dual-core T40E x86 CPU with 64-bit support and 2 GB Dynamic Random-access memory (DRAM). Communication via IEEE 802.11p was achieved using the Unex DHXA-222 Mini PCI Express (mPCIe) card as the WNIC. This chip is based on

the Atheros AR9462 chipset which is supported by the ath9k Linux driver. To enhance storage and memory performance, a [Serial Advanced Technology Attachment \(SATA\)](#) III Transcend MSA370 MCL NAND Flash [Solid State Drive \(SSD\)](#) was installed. The [OS](#) used was OpenWrt release 18.06.1 with Linux kernel 4.14.63. The authors made modifications to the ath9k Linux driver in order to utilize the channels of the 5.8/5.9 GHz frequency band in accordance with [ITS-G5](#) standards. These modifications were then integrated into the OpenC2X project, which was subsequently ported to OpenWrt. The paper also reviews some modifications and implementations made in higher layers, but these are not relevant to the problem addressed in this dissertation.

Sedar et al. [20] developed and validated an experimental, standards-compliant [OBU](#). Their experimental platform is based on an open-source software implementation of the [ETSI ITS](#) protocol stack based on OpenC2X. Its purpose is to facilitate interoperability in communication between various devices and cloud-based services. The [OBU](#) was built using general purpose hardware in the form of a generic laptop, running Ubuntu 18.04. To grant cellular connectivity, it was connected to an [LTE](#) AirPrime EM7565 modem from Sierra Wireless which was in part connected to the 4G cellular network of Vodafone-Spain. The experimental [OBU](#) is also connected to external hardware devices, these being a 4G/5G cellular modem, a [GPS](#)/global navigation satellite system receiver and a connector to receive information from in-vehicle sensors. The article presents an overview of the current state of open-source software implementations of the [ETSI ITS](#) protocol stack. It mentions OpenC2X [76] and Vanetza [77] as the two primary implementations. Although real vehicles were not used in testing, the authors state that their experimental platform can be easily integrated into any vehicle.

Secinti et al. [78] proposed an architectural model that implements [SDN](#) and virtualization principles in order to enable [VANET](#) with Wi-Fi access capability. In this architecture, both the [OBU](#) and the [RSU](#) are implemented using the same type of hardware and software. Both have been implemented using a Raspberry Pi, with the wireless connectivity being provided by the Realtek 5370 Wi-Fi SoC. These switches are implemented using OpenvSwitch v2.3.90 running on OpenWRT. OpenvSwitch is a software switch implementation that is open-source and natively supported by the Linux kernel [79]. Finally, the controller is implemented using OpenDaylight and the southbound [API](#) used is Openflow.

Rito et al. [80] present the deployment and experimentation architecture of the Aveiro Tech City Living Lab in Portugal. The implementation involves a diverse range of devices connected through fiber, radio [ITS-G5](#), and cellular links, utilizing various technologies such as [SDN](#), named data networking, and fog computing. Vehicle and roadside stations are implemented using the PC Engines APU2 board equipped with an [SSD](#), an [IEEE 802.11a/b/g/n mPCIe](#) wireless card and an [LTE CAT-1 mPCIe](#) or 5G m.2 module. Additionally, an external USB dual-band wireless adapter was also installed. The operating system used is not specified, but it is a linux distribution because in order to enable the European version of [IEEE 802.11p](#) in it they used the Linux ath9k driver. This paper implements a myriad of different technologies in vehicular infrastructure, one of them being [SDN](#). It is notable that it does not implement [SDN](#) in vehicle [ITS](#) stations, but only in the backbone of the network. The purpose of this [SDN](#) implementation was to use the increased control in the backbone of the vehicular infrastructure and the vehicle information to predict and execute handovers in advance. The authors developed a custom protocol dubbed OBUIInfo to

provide **CAM** information to the controller. This provides the controller with location, heading, speed, and vehicle type information, which is useful in predicting future handovers ahead of time.

Sadio et al. [81] propose a complete **SDVN** prototype design. The hardware used for the **OBU** was a Raspberry Pi 3 with Cortex-A53 × 64 1.2 GHz and **DRAM** 1 GB. This board has access to WiFi 2.4 GHz **IEEE** 802.11 b/g/n and via a Huawei E8372 **LTE** USB modem to **LTE**. The operating system used was Raspbian Stretch Lite. The authors utilized the Python Twink library to transform the devices into OpenFlow switches. This implementation fails to use the standard for communication in the **VANET** environment, **IEEE** 802.11p.

Paper Reference	Objective	Hardware Used	Software Used	Communication Protocol	Implementation Type	Key Contributions/Notes
Raviglione et al. 2019 [75]	Create a vehicular testbed	PC Engines APU1D, Unex DHXA-222 WNIC	OpenWRT 18.06.1, modified ath9k driver	IEEE 802.11p	VANET	Provided detailed hardware and software to assemble a testbed.
Sedar et al. 2021 [20]	Standards-compliant OBUs	Laptop, Sierra Wireless LTE modem, GPS receiver	Ubuntu 18.04, open-source ETSI ITS stack	LTE , 4G/5G	VANET	Used general-purpose hardware, open-source protocol stack.
Secinti et al. 2017 [78]	SDN -enabled VANET architecture	Raspberry Pi, Realtek 5370 Wi-Fi SoC	OpenWRT, Open-vSwitch v2.3.90, OpenDaylight	Wi-Fi	SDVN	Implemented both OBU and RSU using Raspberry Pi with SDN .
Rito et al. 2023 [80]	Deployment of Aveiro Tech City Living Lab	PC Engines APU2, mP-Cle wireless and LTE cards	Linux (unspecified), ath9k driver	IEEE 802.11p, fiber, LTE , 5G	Hybrid (SDN in backbone only)	Custom protocol (OBUIInfo) for handover prediction in SDN backbone.
Sadio et al. 2020 [81]	Complete SDVN prototype	Raspberry Pi 3, Huawei LTE USB modem	Raspbian Stretch Lite, Python Twink library	Wi-Fi, LTE	SDVN	Used Raspberry Pi for OBU , non-standard protocol for VANET communication.

Table 2: Summary of related literature

In closing, table 2 provides a concise summary the related works presented in this chapter. Main characteristics analysed are the hardware and software used, technologies and communications protocols

supported by the devices, and the key contribution identified on the papers. The table also highlights the ones with SDN support. The devices identified constitute the starting point for the proposed solution to be detailed in next chapters.

4.4 Final considerations

The comprehensive examination of the two fundamental pillars that constitute [SDVN](#) and [SDVN](#) itself has provided insights into the potential usefulness of this technology in real-world scenarios, thereby enabling the identification of an conceptual practical deployment for this technology. As such, in this section, this conceptual view created from the accumulated knowledge presented in this document will be described. Before proceeding to the specifics, it is important to note that this is a theoretical model for an implementation. Consequently, any practical implementations developed in this dissertation will be constrained by the availability and costs of the requisite hardware and the availability and compatibility of the related software, in addition to timing and implementation constraints.

4.4.1 The impact of interoperability

In light of the requirement for interoperability, as delineated in section 4.1.1, an [SDN](#)-compatible OBU deployment is only viable in real-world scenarios if it is capable of communicating with [ETSI](#) conformant devices. This limitation for implementation is only applicable in real-world deployments, as the devices must comply with, or at the very least be compatible and able to collaborate with, the protocols being used in its surrounding in order to be effective. No existing software stack for [SDN](#) is compatible with the [ITS-G5](#) protocol stack. Consequently, the most straightforward approach to create a compatible environment is to utilize [P4](#)'s capacity to reprogram the packet logic of a network. The vast majority of research documents on [SDVN](#) only implement the first generation of [SDN](#), with the few that delve into [P4](#) and its advantages only doing so on a theoretical level, never providing an implementation platform. To the best of our knowledge, and in accordance with the findings of Sarpong et al. [71], no work has been published using [P4](#) as the data plane technology in a [SDVN](#) implementation. It is essencial to note that interoperability diminishes the freedom introduced by [P4](#). However, it does not entirely negate its advantages, because the ability to modify network protocols enables protocol experimentation, thereby allowing existing protocols to be tested, validated, and optimized with greater ease. These observations lead us to conclude that the optimal [SDVN](#) scenario would entail the use of a [P4](#) device that emulates the protocol fields of the [ETSI](#)-defined protocol stack, modified only to accommodate changes to the controller. The aforementioned conclusion, derived from the collected theoretical evidence, suggests that OpenFlow-based solutions may not be the optimal choice. Nevertheless, in order to corroborate this hypothesis, the necessary practical test will be conducted.

4.4.2 Controller placement

The controller prompts the second issue of controller placement. Section 4.2.2 displays the three different possible locations where the controller can be implemented, but most SDN developments only implement controller instances in the wired infrastructure. Due to the issues touched upon in subsection 4.1.1 and 4.1.4 devices must be able to function in tandem with traditional devices and without connection to infrastructure. Ku et al. [67] raised deserved attention to these issues by incorporating recovery mechanisms in their experimental scenario. The authors employed a conventional backup algorithm in the event of a device disconnection with the controller, demonstrating that SDVN must have failure recovery mechanisms to ensure the network works in all scenarios. This work will not explicitly focus on the SDN controller or related issues, such as identifying the best placement or distribution scheme for the SDN controller. This is because it is essential to examine all sub-systems within the ITS domain in order to achieve meaningful results, and this dissertation will focus on the Vehicle ITS sub-system. However, for the aforementioned reasons, it should be considered fundamental for a practical SDVN deployment to contain a local controller able to operate each individual device and take over with something similar to traditional software if the connection to the infrastructure is lost.

Device design and implementation

This section commences with an analysis of the available options, the inherent limitations that were identified, and other significant considerations related to the implementation of this project. Subsequently, the specifications of the system implemented in this dissertation will be described, along with the rationale behind the decisions made during the implementation phase. This process cannot be conducted without first acknowledging the constraints imposed by the unique circumstances of this particular case. The first of these constraints is financial, as the absence of external funding significantly reduces the range of available options. Furthermore, the options for suitable hardware are constrained by their compatibility with the essential free and open-source software that will be utilized in the project. The software in question is of critical importance in meeting the necessary criteria for the system to be considered as vehicular. In the majority of cases of this nature, the issue of performance is also a significant concern. The capabilities of the components of the device are of significant importance when considering the performance of our solution, as they directly impact the device's performance. However, our main goal is to demonstrate the feasibility of such a device, and as such it is necessary to set aside concerns regarding performance in order to reduce the cost. In the event that any of the aforementioned restrictions prove to be applicable, they will be duly noted in the appropriate section.

5.1 Architectural design

This section delineates the design decisions that were implemented in the process of assembling the device, in accordance with the architecture of the [SDN](#) layers.

5.1.1 Data plane

This subsection offers a comprehensive overview of the data plane's fundamental components, encompassing its hardware elements, [OS](#), wireless communication strategies, and virtualized switching programs.

5.1.1.1 Hardware components

The selection of appropriate hardware was guided by the different categories of SDN devices outlined in section 4.2.2. Accordingly, our analysis began with an assessment of the viability of each of the three proposed approaches for the intended implementation of the solution resulting from this project. In light of these criteria and of the existing financial restrictions, general-purpose hardware emerged as the optimal choice for this problem, largely due to its comparatively lower cost. The proposed implementation will be based on the use of general-purpose hardware, with the switching functions being emulated through the use of software. The process of assembling any device requires attention to the individual components that will be incorporated and that comprise the complete system. The most notable components include the motherboard, the CPU, the Random-access memory (RAM), the Network Interface Controller (NIC), the power supply, and the case. It should be noted that a multitude of additional components could be considered, given that, in real world development, an OBU is often equipped with a plethora of other specialized dongles. The selection of the motherboard exerts a determining influence on the subsequent choice of other components; therefore, it became the starting point. Following an exhaustive review of a multitude of products from various corporations in all the methodologies identified, the products from two companies have emerged from the comprehensive research conducted. These solutions have been identified as sufficient candidates for addressing the specific challenges in our case study.

- PC Engines [82]: PC Engines boards are designed for use in academic institutions and for networking purposes. The company provides not motherboards, but complete affordable solutions, which include CPU, RAM, NIC, power supply, and case. Additionally, these solutions rely on active cooling, which facilitates their assembly and reduces operational costs. Another notable feature is the number of ethernet ports and the quantity of Peripheral Component Interconnect (PCI) slots, which are both sufficient for the intended use. Lastly, other projects at this university have utilized boards from this company.
- Raspberry Pi [83]: Raspberry Pi is a very well known series of small single-board computers developed by the Raspberry Pi Foundation in association with Broadcom. As a highly recognizable brand, it offers comprehensive assistance and is a well-established brand with a substantial code and user base. This solution's main strength lies in its modularity, as the hardware can be assembled in accordance with specific requirements, thereby allowing for a high degree of flexibility and adaptability. These components are mostly inexpensive, which presents a valuable upside when considering deployment.

From these two options, and given the established familiarity, the PC Engines option stands as the most appropriate choice. As stated above, PC Engines' boards are sold as integrated solutions, comprising a pre-selected range of components. The decision regarding the specific model to be adopted from this company will consequently have a significant impact on a multitude of other components. The exact model that will be deployed will be discussed in detail at a later stage. In regard to the potential utilization of dongles in a real-world development setting, it is evident that the device would be equipped with a

multitude of sensors intended to gather diverse types of data and enhance the device's capabilities. This work will primarily focus on a proof of concept for the integration of the communication component of VANETs with SDN principles. Consequently, further discussion of the sensors used in vehicles will be omitted. The only noteworthy dongle that we will be employing are bidirectional wireless devices capable of communication at 6GHz.

5.1.1.2 Operating System

Having completed the analysis of the hardware, the appropriate next step is to proceed with the examination of the software. The foundation of any software architecture stack is the BIOS of the system. It is thus fitting to initiate this analysis at this level. In a context where specific switching hardware had been selected, the BIOS of the system would assume a greater significance. In this particular scenario, it would be reasonable to consider the potential use of Open Network Install Environment [84], an open and free standard specifically designed for these use cases. However, given that we will be employing general-purpose hardware, there is no compatibility with such software, and we are therefore effectively compelled to utilize the BIOS that is natively integrated into the motherboard. The OS is installed on top of the BIOS, and thus, this discussion will proceed with an examination of the extant options for the OS. Once more, in the event that we were to utilize switching hardware, Open Network Linux (ONL), Stratum or Software for Open Networking in the Cloud (SONiC) would be decent choices. Stratum, ONL and SONiC are open-source OSs designed for use with switch hardware. Given that we intend to utilize generic hardware, any compatible Linux-based OS is an adequate option.

5.1.1.3 Wireless communication

Our research is focused on achieving wireless communication in vehicular environments. To this end, it must be ensured that the machine is able to communicate using one or both of the IEEE 802.11p and C-V2X standards. It would be very beneficial for the device to be capable of communicating with LTE and 5G networks; given the time constraints, the priority for this device is the ability to communicate in the IEEE 802.11p standard. As one descends through the hierarchical structure of the ETSI standards for vehicular communication, it becomes increasingly challenging for SDN technologies to adapt and conform to existing standards. All layers above the link layer can be readily modified or altered using SDN technologies in a manner that allows them to be considered vehicular. The IEEE 802.11p standard, which represents the lowest level of the stack, is not as straightforward as other layers. Our research revealed three distinct methods for enabling it:

- Purpose made hardware: The first option is to acquire specialized hardware that is capable of communicating at the desired frequency while complying with the necessary standards. This solution is the most straightforward, but also the most expensive. Consequently, for the same reasons as for the rest of the hardware, this option is the most viable solution and will not be considered further.

- Driver patch: Some researchers have solved this problem [85] by employing the use of a capable chipset and modifying the driver code to enable communication in the desired manner. This approach has been previously implemented in other research projects at this institution, and as such the first box received for use in this project is already equipped with this patch. This solution is especially appealing because it is the most affordable. However, it is also the most complex solution, as the implementation of such a feat requires a comprehensive understanding of networking, the IEEE 802.11p standard, drivers, and kernel modules. Even with this knowledge, a chip that is capable of communicating at the desired frequencies and with an open-source device driver for linux systems is required.

The intricate nature of this component renders the development of a solution from scratch an arduous and time-consuming process, which is impractical within the context of this project. Fortunately, as previously stated, there is existing research that can be utilized to achieve this goal. In a similar context, this university has previously employed a patch for an atheros chip. This family of chips operates with open-source drivers, of which the driver in question is named "ath9k." One factor that merits attention is the age of this driver. The ath9k driver is designed for IEEE 802.11n, which was released in 2009, and raises concerns about its long-term reliability. Furthermore, there appears to be no initiative within the community to implement this solution on modern drivers, nor to make the existing implementations official [86].

- Software Defined Radio (SDR): SDR is a relatively recent technology which bears resemblance to SDN. It enables the user to modify and configure radio signals. In comparison to the preceding option, this appears to be a more long-term solution, albeit at a higher price point. The subject of SDR research is beyond the scope of this dissertation, and therefore it will not be a topic of discussion.

Of the three options presented, the most promising is the utilization of an Atheros chip in conjunction with the existing patch. Should this approach prove inadequate, it will be necessary to consider one of the alternative strategies.

5.1.1.4 Virtualized switch program

The selected hardware type mandates the use of software to emulate the typical functionality of a switch. The implementation of this methodology is contingent upon the availability of virtualized switch software. Our research has identified two potential options: Open virtual Switch (OvS) [87] and Behavioral Model version 2 (BMv2) [88]. OvS is an open-source implementation of a distributed virtual multilayer switch, designed with the intention of providing a switching stack for hardware virtualization environments. In contrast, BMv2 is a pure P4 software switch, which is primarily utilized as a testing tool for the P4 language. Both options are available as free software, but it should be noted that OvS is not merely an OpenFlow switch. Consequently, it is significantly more utilized than BMv2 and offers considerably broader support.

In this study, we will evaluate both solutions to ascertain their capabilities and viability within the context of our conceptual framework.

5.1.2 Southbound API

The decision regarding the southbound API is unambiguous and will be made in accordance with the established conventions within this field. Accordingly, tests will be conducted on both OpenFlow and P4Runtime. In regard to the question of the gNIXI stack, its role is not indispensable for demonstrating the research question, and thus it will be set aside for the purposes of this study. While the value of integrating gNIXI in communication is of great importance in real deployments, this study aims to prioritize the establishment of communication, and thus the integration of such systems will be considered secondary.

5.1.3 Control plane

Generally, each controller framework comes with its own properties, features and characteristics. The choice of what controller is best for a network is dependent on performance and compatibility. As previously stated, the issue of performance is not a significant concern in this particular context. In regard to compatibility, the term denotes the ability of the controller framework to operate in conjunction with a multitude of different technologies. This is not a cause for concern since the lower layers utilize open and industry-standard protocols and software, which are universally adopted. It is currently not feasible to integrate SDN software with Vanetza, OpenC2X, or other software that emulates the ITS stack. Additionally, none of the controllers currently available offer support for the standardized VANET protocols, necessitating a comprehensive redesign of the system. When all factors are considered, the choice of control technology is ultimately a matter of preference. Smida et al. [66] compared the performance of four of the most popular controllers in an SDVN environment. These were Pythonic Network Operating System, Floodlight, OpenDaylight and ONOS. The article concludes that the best performing controller framework is ODL, although the results demonstrated that ONOS, Floodlight, and ODL exhibited comparable performance when a minimal number of vehicles were involved. In light of the aforementioned stipulations, any of the aforementioned controllers would be an adequate choice. In the absence of a sufficiently compelling argument to make any of the aforementioned controllers the superior choice, the decision was made to select ONOS [89].

5.1.4 Northbound API and Application plane

The choice of either the Northbound API or the Application plane is inconsequential, as it is entirely contingent upon the specific controller in question. Moreover, the present study will concentrate on the lower layers.

5.2 Box implementation: version 1

To initiate the practical component of this dissertation, a fully operational device was received consisting of a PC Engines apu1d4 motherboard, containing an AMD G-series T40E [CPU](#) and 4 GB of [DRAM](#). The device also possesses three Realtek RTL8111E Gigabit Ethernet channels, a DB9 serial port, and two external USBs. Beyond these features, the device is equipped with two [mPCIe](#) interfaces and one [Mini-Serial Advanced Technology Attachment \(mSATA\)](#) interface. These specifications, along with more specific details, were sourced from the product page and are presented as follows:

- [CPU](#): The device has an AMD G-series T40E, operating at a frequency of 1 GHz. It is comprised of two Bobcat cores, each of which is dual-threaded and capable of supporting 64-bit instructions. Each core has 32K of data, 32K of instructions, and 512K of level 2 cache.
- [RAM](#): 4 GB of DDR3-1066 [DRAM](#)
- Storage: the device may be booted from an [Secure Digital \(SD\)](#) card connected via USB, an external USB drive, or a [mSATA SSD](#). The system's [mSATA](#) interface is equipped with the requisite power connector for the [SSD](#) to function.
- Power supply: the device consumes approximately 6 to 12 W at 12 V DC, with the exact value depending on the [CPU](#) load. The jack is 2.5 mm and center positive.
- Connectivity: The device features three Gigabit Ethernet channels, which are Realtek RTL8111E.
- Input/Output: The device features a DB9 serial port, four USB ports (two external and two internal), three front panel LEDs, and a pushbutton.
- Expansion: The device features two [mPCIe](#) slots, one of which is equipped with a [Subscriber Identity Module \(SIM\)](#) socket.
- Size: The board measures 152.4 x 152.4 mm. The box that encompasses it totals 15.7 cm x 16.8 cm x 2.7 cm, or 16.7 cm x 18.8 cm x 3 cm when all extremities are included.
- Firmware: As part of its low-level system initialization, this device relies on the Coreboot firmware.
- Cooling: A conductive cooling system is employed, whereby the [CPU](#) and south bridge are connected to the enclosure via a 3 mm aluminum heat spreader.

The apparatus was furnished with a 16 GB [SSD](#), a Compex WLE200NX [mPCIe](#) card for WiFi, and another specialized card for [LTE](#), which is not germane to our present discussion. To support the necessary functionality of the aforementioned cards, the device was also fitted with four pigtail cables. The assembled device can be seen on Figure 26.

The device runs on the Debian 12 [OS](#) and has been previously configured to enable communication in the capacity of an [OBU/ITS](#) vehicle station. The inaugural stage of implementation aimed to ascertain



Figure 26: Back and inside view from device version 1

the feasibility of the proposed approach and identify potential avenues for further exploration. It was thus deemed appropriate to commence with an examination of the switching software. The original intention was to install [OvS](#) and [BMv2](#) on the aforementioned device and test their compatibility with the existing software. However, this plan was promptly halted due to the device's lack of the necessary resources for installing both programs, and as such we were compelled to deploy [OvS](#) exclusively on this device. A series of tests was conducted to evaluate its capabilities, utilizing both a local controller and a remote [ONOS](#) controller. The limited capabilities of the devices also motivated us to opt against installing local controllers on each device. This is contrary to our [SDN](#) vision, but it is a consequence of the constraints on resources. Another noteworthy obstacle encountered was the insufficient number of test units, which hindered the assessment of the system's wireless communication capabilities between multiple units with a controller.

5.3 Box implementation: version 2

The dearth of a sufficient number of devices for testing halted our progress. In light of these circumstances, the procurement of one or more supplementary devices became imperative. To that end, work began with the aim to design, and subsequently assemble, two new prototype devices with enhanced specifications. Two recently constructed prototypes were assembled on a PC Engines apu3d4 motherboard. The motherboard in question is equipped with an AMD G-Series GX-412TC Embedded processor, 4 GB of [DRAM](#), three Intel i211AT Gigabit Ethernet channels, a DB9 serial port, and two external USB 3.0 ports. The device is additionally equipped with three [mPCIe](#) slots, one of which is also capable of functioning as a [mSATA](#). These specifications, along with more specific details, were sourced from the product page and are presented as follows:

- **CPU:** The device is equipped with an AMD Embedded G-Series GX-412TC, operating at a frequency of 1 GHz. It comprises quad Jaguar cores with 64-bit processing and AES-NI support. Each core has 32 KB of L1 data cache, 32 KB of L1 instruction cache, and a shared 2 MB L2 cache.
- **RAM:** 4 GB of DDR3-1333 [DRAM](#)
- **Storage:** The device may be booted from an [SD](#) card, an external USB connection, or an [mSATA SSD](#). The system's [mSATA](#) interface is furnished with the necessary power connector for the [SSD](#) to operate, and it can also function as an [mPCIe](#) interface.
- **Power supply:** the device consumes approximately 6 to 12 W at 12 V DC, with the exact value depending on the [CPU](#) load. The jack is 2.5 mm and center positive.
- **Connectivity:** The device features three Gigabit Ethernet channels, which are Intel i211AT.
- **Input/Output:** The device features a DB9 serial port, two USB 3.0 external ports, and four USB 2.0 internal ports. Additionally, it exhibits three front panel LEDs and a pushbutton.

- Expansion: The device is equipped with three [mPCIe](#) slots. Of the aforementioned slots, one is a USB or [mSATA](#) slot that can accommodate a [SIM](#) card. The second is a USB-only slot that can also accept a [SIM](#) card. The final slot is a full [mPCIe](#) slot, devoid of a [SIM](#) card, but equipped with a WiFi module.
- Size: The board measures 152.4 x 152.4 mm. The box that encompasses it totals 15.7 cm x 16.8 cm x 2.7 cm, or 16.7 cm x 16.8 cm x 3 cm when all extremities are included.
- Firmware: As part of its low-level system initialization, this device relies on the Coreboot firmware.
- Cooling: A conductive cooling system is employed, whereby the [CPU](#) is connected to the enclosure via a 3 mm aluminum heat spreader.

The device was outfitted with a Kingston 256GB [SSD](#), a Compex WLE9000VX 7AA WiFi card, and a Quectel EC25-E for [LTE](#). The prototype case under consideration is only capable of accommodating two pigtail cables. In consideration of our intention to utilize WiFi exclusively for the duration of our investigation, this case is deemed sufficient. Nevertheless, should [LTE](#) be required in the future, enhancements to the case will be mandatory. One of the two newly assembled devices can be seen on figure 27.

As the device is devoid of any software, we shall adhere to the aforementioned guidelines that have been previously established. It was thus deemed appropriate to proceed with the Ubuntu 24.04.1 LTS, given that the visual interface of the [OS](#) was considered to be of no consequence. A minor complication manifested as a result of the incompatibility between the selected WiFi card and the existing solution. During the device setup process, it became evident that, although the WiFi card was from the atheros family, it utilized the ath10k driver. The ath10k driver has been developed as a more recent iteration of the ath9k driver. However, at the time of writing, no patch for [IEEE](#) 802.11p has been made available for this driver. This necessitated the pursuit of a different chip that was compatible with the ath9k driver. Fortunately, the Linux wireless community maintained an up-to-date list of products that met the necessary criteria [90]. From the list of available options, the Compex WLE200NX 7A was selected.

5.4 P4/ONOS application implementation

The technological stack previously delineated and to be utilized in the ensuing tests includes [BMv2](#) and [ONOS](#). Consequently, it is necessary p4 code for the data plane devices accompanied by Java code for the controller.

At the outset of the practical component of the dissertation, the original plan was to develop a solution that more closely resembled the [ETSI](#) protocols. However, it was ultimately decided that this approach would not be pursued for two primary reasons. It is first necessary to acknowledge the inherent complexity of the [ETSI](#) protocols. The standards under discussion have been developed over the past two decades by the most esteemed engineers in the field. Given the considerable time and resources required, it is neither realistic nor reasonable to view this as a feasible task in view of the limitations in terms of human

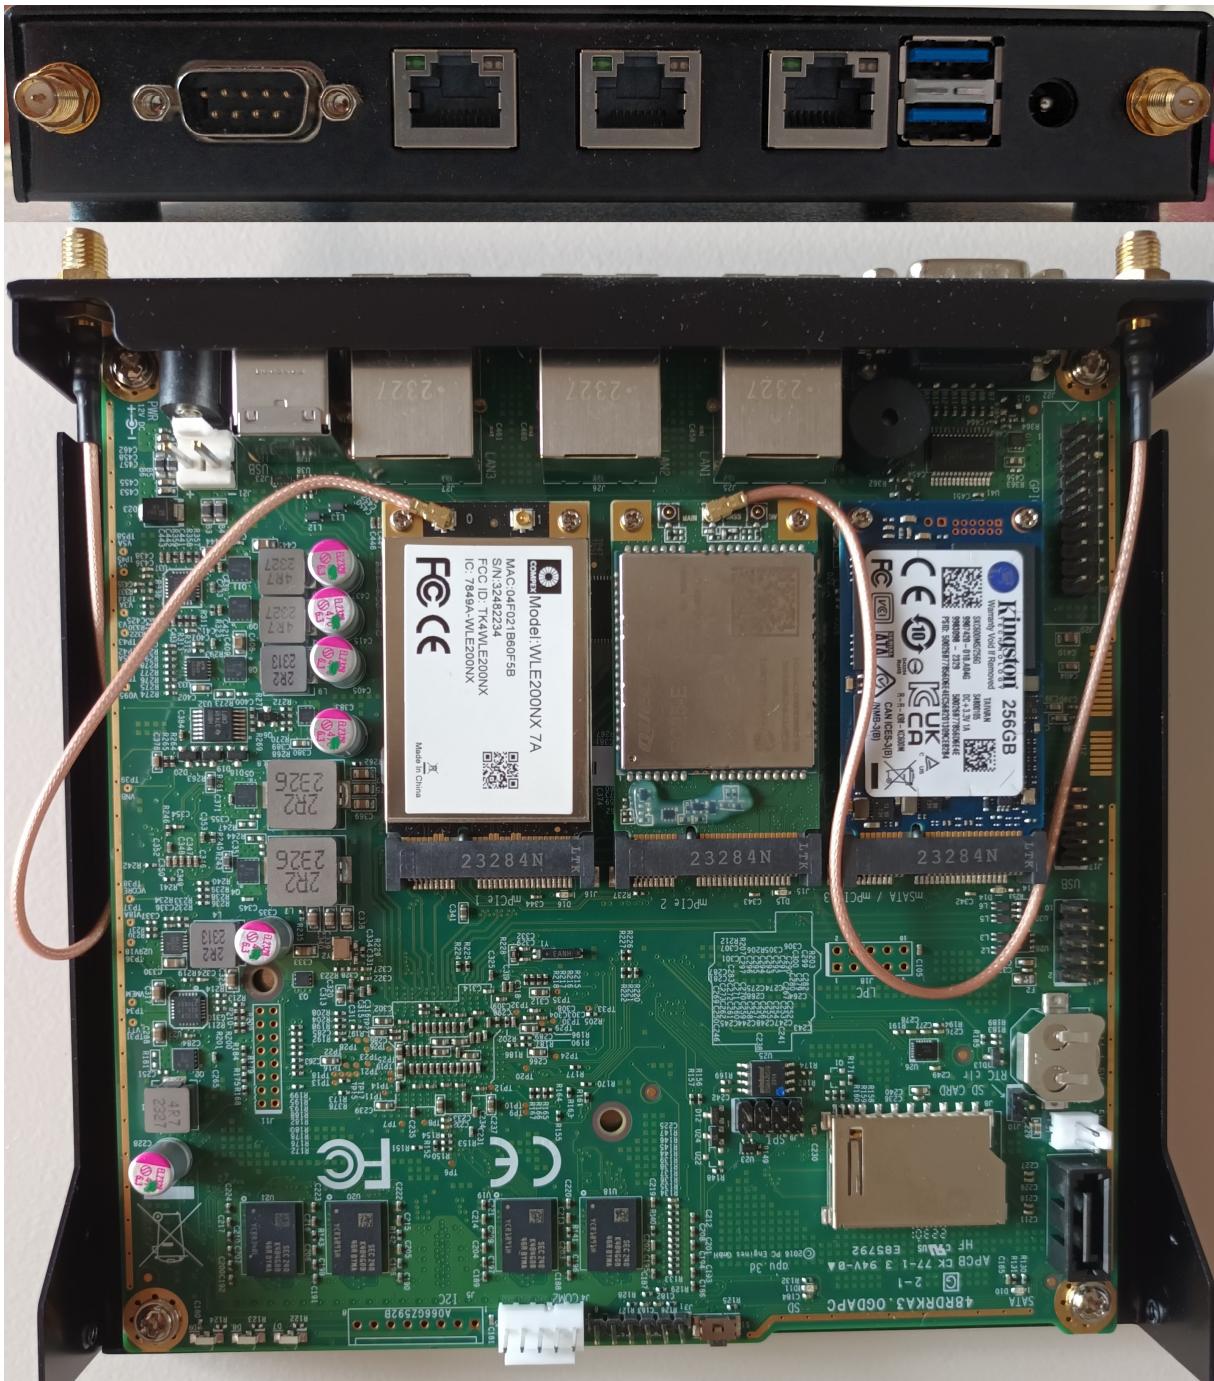


Figure 27: Back and inside view from device version 2

resources and time imposed by the context of this master's dissertation. Secondly, the [ETSI](#) norms establish the indispensable presence of the remaining sub-systems as a mandatory condition for attaining the standardized functionality. These circumstances provided the impetus to pursue an alternative and more appropriate course of action.

In lieu of attempting to create an exact replica of the [ETSI](#) protocol using [SDN](#) technology, it was elected as more appropriate to pursue the development of a proof-of-concept solution. This proposed endpoint has the dual objective of validating the compatibility of [P4](#) with wireless interfaces and [IEEE 802.11p](#), while simultaneously demonstrating the adaptability of [P4](#) functional capabilities within these specific contexts. In order to validate the first proposition, the proposed framework will rely on broadcast communication strategies within the wireless domain. The second part will be illustrated through the use of custom headers. While there are numerous potential applications for custom headers, for the purposes of this study, we sought to specify an appropriate research objective. Ultimately, it was determined that the prevention of the perpetual repetition of broadcast messages was a necessary measure to avoid severe network saturation and therefore an appropriate goal. To this end, we created a "marker" system. The markers are unique identifiers generated by the controller and appended to each message transmitted via the wireless interfaces. This strategy guarantees that no message is broadcast twice by the same device. The final solution was developed into a [P4/ONOS](#) application and the result was published to github [91].

[ONF](#), as part of the Aether Project, offers a practical tutorial to teach the fundamental principles of the next-generation [SDN](#) architecture, which served as the basis for our program. This tutorial was developed for use with Mininet and Stratum, which immediately indicated the necessity for modifications. A more thorough examination of the code revealed the presence of additional issues. It was thus necessary to modify a substantial portion of the code to guarantee its compatibility with the specified environment before it could be utilized. Indeed, a considerable investment of time and effort was necessary to make the program operational.

5.5 Summary

In this chapter, two implementations were presented, detailing design choices and major requirements. The final box version 2 is the current version in use on other projects. For future reference, the alternatives enumerated in this chapter should be considered for the development of version 3.

Functional Tests

In this chapter, the selected hardware prototype is tested, in conjunction with the desired software, and the results presented. The objective of these tests is to ascertain the viability of the software selected for practical deployment of the prototypes. Specifically, this study sought to evaluate the viability of both [OvS](#) and [BMv2](#) in vehicular scenarios. In particular, the aim of the forthcoming testing is to ascertain the degree of compatibility between the aforementioned switching emulation software and a number of other components, including the ath9k patch, the [ONOS](#) controller, and the hardware of the device in question.

6.1 Experimental Setup

In consideration of the practical implications of the goals of this dissertation, the device must meet the criteria for an [OBU](#)/Vehicle [ITS](#) sub-system, in accordance with the [SDN](#) paradigm. For this, the device must be categorized as both [SDN](#) and [VANETs](#). In practical terms, for a device to be considered [SDN](#)-compliant, it is necessary to utilize software to emulate the functions of an [SDN](#) device. Moreover, the minimum requirement for a device to be appropriate for the context of [VANETs](#) is its capability of establishing connectivity in accordance with the [IEEE](#) 802.11p standard. Accordingly, the study employs a three-stage experimental methodology, consisting of installation, integration, and functionality verification. The tests are performed in a series of phases, with each subsequent phase building upon the preceding one and introducing greater degrees of complexity. Their objective is to assess the compatibility of the various elements utilized, namely the hardware, the ath9k patch, and [OvS/BMv2](#). The success of this endeavor is contingent upon the establishment of a connection between devices, which is conducted using the ping tool. The experiments were conducted on the devices described in Section 5.2 and 5.3 with the assistance of an external device, namely a personal computer. Additionally, a switch and six network cables were employed to facilitate communication between the devices and the personal computer, while omnidirectional antennas operating at 6GHz were utilized to enable communication in the requisite frequency range. The software tools utilized in this investigation were as follows:

- Windows 10 is the operating system of the [PC](#);

- Oracle VM Virtualbox 7.0.4 was installed on the host computer and employed to both install and run the [ONOS](#) and p4c;
- [ONOS](#) 2.7.0 was utilized for the controller;
- p4c is the compiler for the [P4](#) language;
- PuTTy Release 0.80 was installed on the [PC](#) and is used to interface with the machines;
- Ubuntu 24.04.1 LTS operating system is used on devices version 2;
- Debian 12 operating system is used on device version 1;
- [OvS](#) and [BMv2](#) are the focus of the current testing.

It is essential to acknowledge that, despite the comprehensive and rigorous nature of these tests, they are not without inherent limitations. Such limitations include the specific conditions of the controller environment and the exclusion of real-world network conditions. The methodology selected was chosen for its capacity to facilitate the systematic identification of critical compatibility issues that must be addressed for successful real-world deployment. Moreover, for the purposes of this study, it was assumed that the devices lacked the requisite resources to install and run [ONOS](#).

6.2 Experiment 1: testing the integration of [OvS](#)

This experiment is intended to assess the compatibility and practicality of integrating [OvS](#) and OpenFlow with [VANET](#) technologies. The initial stages of the experiment were conducted using a single device. Around the midpoint of experimentation, the additional two devices of version 2 arrived and were incorporated into the test setup.

6.2.1 Testing phase 1: standalone device with [OvS](#) and local controller

The opening phase of the experiment sought to examine the functionality of [OvS](#) in a relatively straightforward scenario, with the additional objective of becoming familiar with the technology. Accordingly, the procedure commenced with a standalone examination of [OvS](#), which is possible due to the fact that [OvS](#) is equipped with a local controller that is capable of autonomously managing network traffic. [OvS](#) was installed on the device, and a virtual network was created within the device using [Virtual Routing and Forwardings \(VRFs\)](#) and virtual interfaces.

The setup from this phase can be observed in Figure 28, and is diagrammatically represented in Figure 29. When viewed through the lens of [SDN](#), the configuration of this phase can be conceptualized as illustrated in Figure 30.

Figure 31 presents the configuration of [OvS](#), along with information regarding the device interfaces. Here, all valid interfaces on the device are displayed. This list includes two virtual interfaces, which can be seen to have been added to the [OvS](#) bridge. Figure 32 demonstrates the existing interfaces of the two



Figure 28: Setup from the first phase of the experiments with OvS

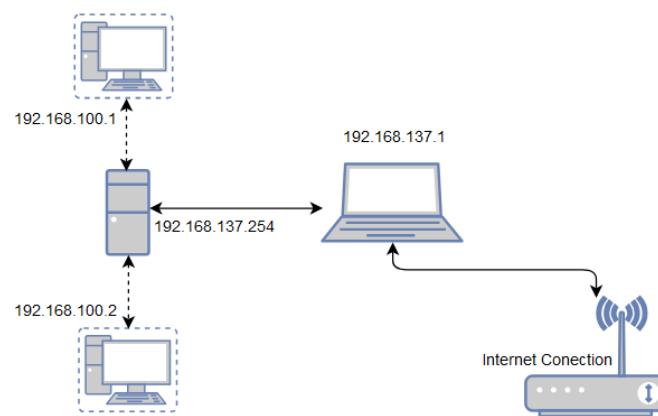


Figure 29: Diagram that represents the setup from the first phase of the experiments with OvS

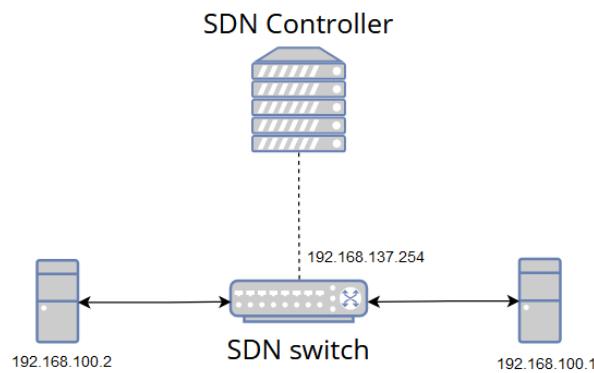


Figure 30: Diagram that depicts the SDN perspective of the setup from the first phase of the experiments with OvS

```
[apu] ~
root@apu ~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inetc6 ::1/128 scope host noprefixroute
        valid_lft forever preferred_lft forever
2: enp1s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 00:0d:b9:52:5e:70 brd ff:ff:ff:ff:ff:ff
    inet 192.168.137.254/24 brd 192.168.137.255 scope global enp1s0
        valid_lft forever preferred_lft forever
    inetc6 fe80::20d:b9ff:fe52:5e70/64 scope link
        valid_lft forever preferred_lft forever
3: enp2s0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc fq_codel state DOWN group default qlen 1000
    link/ether 00:0d:b9:52:5e:71 brd ff:ff:ff:ff:ff:ff
4: enp3s0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc fq_codel state DOWN group default qlen 1000
    link/ether 00:0d:b9:52:5e:72 brd ff:ff:ff:ff:ff:ff
5: wwan0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/ether 7e:79:ce:9e:3:9e brd ff:ff:ff:ff:ff:ff
6: wlp4s0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/ether 04:f0:21:86:56:48 brd ff:ff:ff:ff:ff:ff
7: ovs-system: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/ether 12:aff4:a3:fd:e7 brd ff:ff:ff:ff:ff:ff
8: ovs-br: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/ether 52:ba:31:39:a0:45 brd ff:ff:ff:ff:ff:ff
9: veth4@if10: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master ovs-system state UP group default qlen 1000
    link/ether 42:3a:80:85:56:33 brd ff:ff:ff:ff:ff:ff link-netns VRF2
    inetc6 fe80::403a:80ff:fe85:5633/64 scope link
        valid_lft forever preferred_lft forever
11: veth2@if12: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master ovs-system state UP group default qlen 1000
    link/ether 7a:0e:da:78:f1:eb brd ff:ff:ff:ff:ff:ff link-netns VRF1
    inetc6 fe80::780e:daff:fe78:f1eb/64 scope link
        valid_lft forever preferred_lft forever
root@apu ~# ovs-vsctl show
2bed74a0-d5cc-4f5f-a541-6d3cd80bdf06
  Bridge ovs-br
    Port ovs-br
      Interface ovs-br
        type: internal
    Port veth2
      Interface veth2
    Port veth4
      Interface veth4
  ovs_version: "3.3.90"
root@apu ~#
```

Figure 31: Results from running the commands "ip a" and "ovs-vsctl show" in phase 1

VRFs created in this device, which are shown to be on the same subnet. Following this, a ping is executed between the two devices to demonstrate their ability to communicate.

6.2.2 Testing phase 2: standalone device with OvS and remote controller

Following the success of this initial phase, the decision was made to proceed with the introduction of the remote controller. The controller, which will be utilized for the remainder of the tests, introduced an elevated level of complexity, providing an excellent opportunity to become acquainted with ONOS environments. Consequently, phase two entailed the creation of a Linux-based [Virtual Machine \(VM\)](#) utilizing VirtualBox, with the installation of ONOS on this [VM](#). ONOS was therefore incorporated into the preceding scenario and connected to OvS via OpenFlow.

The setup from this second phase is visually the same as in the first phase, and can therefore be observed in Figure 28. The setup of this phase is diagrammatically represented in Figure 33. When viewed through the lens of SDN, the configuration of this phase can be conceptualized as illustrated in Figure 34.

The configuration of OvS, in conjunction with the demonstration of connectivity between the two VRFs, is illustrated in Figure 35. This particular OvS state is identical to that of the phase 1 of testing.

```

root@apu ~# ip netns exec VRF1 ip a
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN group default qlen 1000
    link/loopback 00:00:00:00:00 brd 00:00:00:00:00:00
12: veth1@if11: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether e6:4f:3d:f6:8f:e5 brd ff:ff:ff:ff:ff:ff link-netnsid 0
        inet 10.10.10.1/24 brd 10.10.10.255 scope global veth1
            valid_lft forever preferred_lft forever
        inet6 fe80::e44f:3dff:fe6:8fe5/64 scope link
            valid_lft forever preferred_lft forever
root@apu ~# ip netns exec VRF2 ip a
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN group default qlen 1000
    link/loopback 00:00:00:00:00 brd 00:00:00:00:00:00
10: veth3@if9: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 9e:eb:b6:8c:25:e2 brd ff:ff:ff:ff:ff:ff link-netnsid 0
        inet 10.10.10.2/24 brd 10.10.10.255 scope global veth3
            valid_lft forever preferred_lft forever
        inet6 fe80::9ceb:b6ff:fe8c:25e2/64 scope link
            valid_lft forever preferred_lft forever
root@apu ~# ip netns exec VRF1 ping 10.10.10.2
PING 10.10.10.2 (10.10.10.2) 56(84) bytes of data.
64 bytes from 10.10.10.2: icmp_seq=1 ttl=64 time=1.38 ms
64 bytes from 10.10.10.2: icmp_seq=2 ttl=64 time=0.196 ms
64 bytes from 10.10.10.2: icmp_seq=3 ttl=64 time=0.194 ms
64 bytes from 10.10.10.2: icmp_seq=4 ttl=64 time=0.198 ms
64 bytes from 10.10.10.2: icmp_seq=5 ttl=64 time=0.197 ms
^C
--- 10.10.10.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4071ms
rtt min/avg/max/mdev = 0.194/0.432/1.375/0.471 ms
root@apu ~# ip netns exec VRF2 ping 10.10.10.1
PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data.
64 bytes from 10.10.10.1: icmp_seq=1 ttl=64 time=0.179 ms
64 bytes from 10.10.10.1: icmp_seq=2 ttl=64 time=0.195 ms
64 bytes from 10.10.10.1: icmp_seq=3 ttl=64 time=0.196 ms
64 bytes from 10.10.10.1: icmp_seq=4 ttl=64 time=0.198 ms
64 bytes from 10.10.10.1: icmp_seq=5 ttl=64 time=0.197 ms
^C
--- 10.10.10.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4103ms
rtt min/avg/max/mdev = 0.179/0.193/0.198/0.007 ms
root@apu ~#

```

Figure 32: Results from running the commands "ip a" on both VRFs and "ping" between them in phase 1

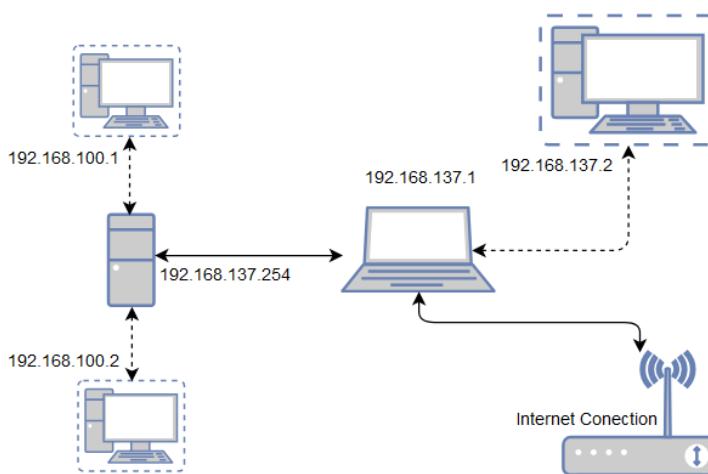


Figure 33: Diagram that represents the setup from the second phase of the experiments with OvS

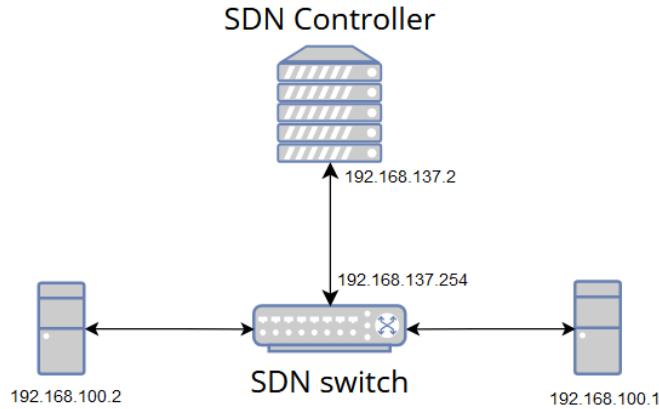


Figure 34: Diagram that depicts the [SDN](#) perspective of the setup from the second phase of the experiments with [OvS](#)

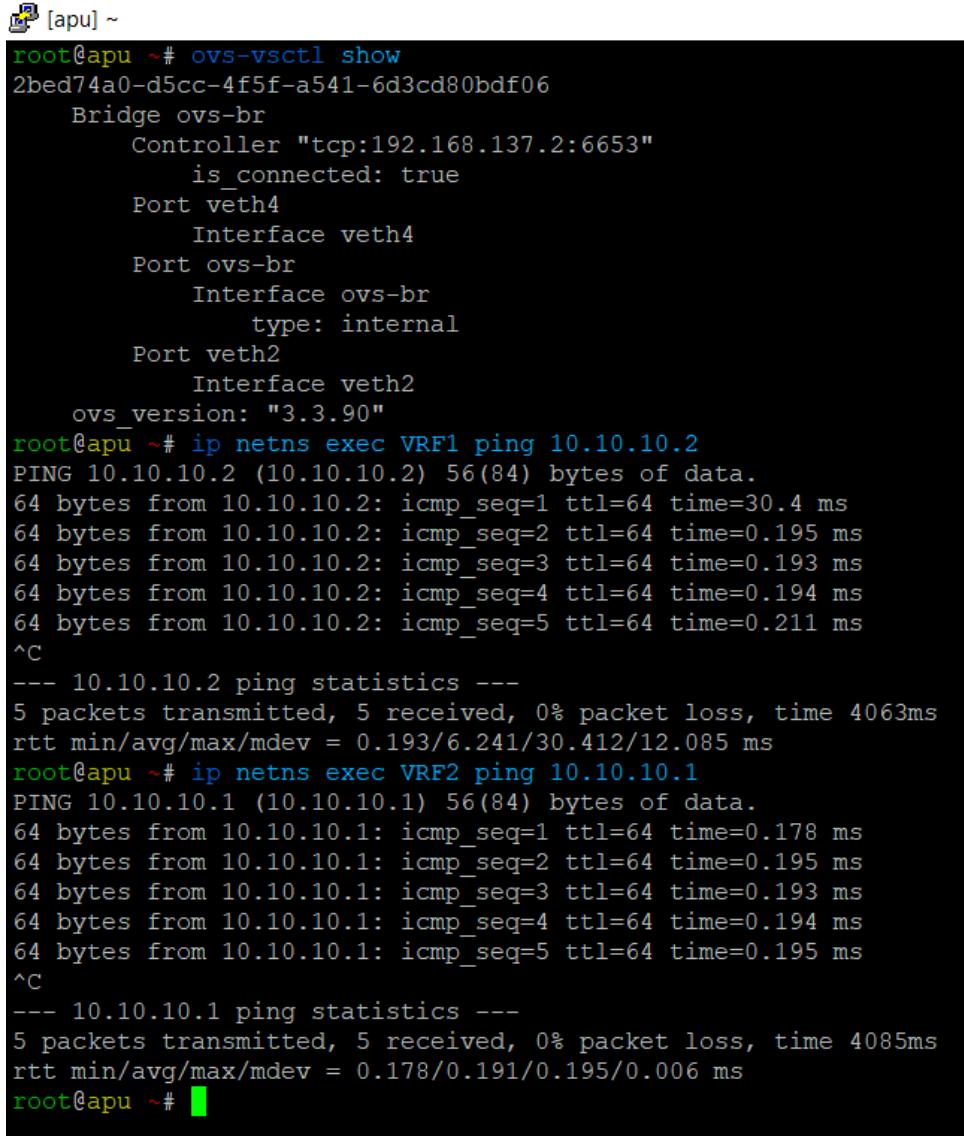
Figure 36 displays a selection of relevant commands that were executed in the [ONOS](#) console. This output demonstrates that [ONOS](#) is connected to one device operating [OvS](#) via OpenFlow version 1.3, and has knowledge of two hosts. Furthermore, we observe the flows that have been automatically integrated by the [ONOS](#) services, thereby providing the switch with the functionality to forward packets.

6.2.3 Testing phase 3: multiple devices with [OvS](#) and remote controller

By the time phase 3 was reached, the two additional devices had been received. Therefore, the test [SDN](#) network was expanded to integrate these two additional devices. [OvS](#) was installed on the two new devices, and the three devices were linked together and connected to the [ONOS](#) controller, and two [VRFs](#) were created, each situated on opposite edges of the network.

The configuration of this phase is illustrated in Figure 37. The setup of this phase is diagrammatically represented in Figure 38 and described in the following itemized list:

- 192.168.137.1 - This IP address identifies the author's personal laptop computer. It serves as the central coordinating hub for the entire network, facilitating the observation and monitoring of all network activity.
- 192.168.137.2 - This address represents a virtual machine that is hosted on the author's laptop via the Oracle VM VirtualBox virtualization platform. The virtual environment is configured with Debian 12 and is primarily intended for the deployment of [ONOS](#).
- 192.168.137.252 - This corresponds to a device version 2. It has a fresh installation of the Ubuntu Server operating system and is running [OvS](#).
- 192.168.137.253 - This device is virtually identical to the device 192.168.137.252.
- 192.168.137.254 - This address represents device version 1. This device has debian 12 and is running [OvS](#).

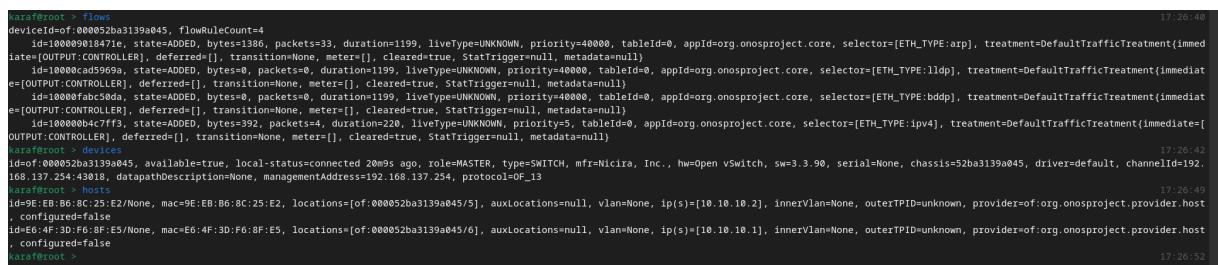


```

[apu] ~
root@apu ~# ovs-vsctl show
2bed74a0-d5cc-4f5f-a541-6d3cd80bdf06
    Bridge ovs-br
        Controller "tcp:192.168.137.2:6653"
            is_connected: true
        Port veth4
            Interface veth4
        Port ovs-br
            Interface ovs-br
                type: internal
        Port veth2
            Interface veth2
    ovs_version: "3.3.90"
root@apu ~# ip netns exec VRF1 ping 10.10.10.2
PING 10.10.10.2 (10.10.10.2) 56(84) bytes of data.
64 bytes from 10.10.10.2: icmp_seq=1 ttl=64 time=30.4 ms
64 bytes from 10.10.10.2: icmp_seq=2 ttl=64 time=0.195 ms
64 bytes from 10.10.10.2: icmp_seq=3 ttl=64 time=0.193 ms
64 bytes from 10.10.10.2: icmp_seq=4 ttl=64 time=0.194 ms
64 bytes from 10.10.10.2: icmp_seq=5 ttl=64 time=0.211 ms
^C
--- 10.10.10.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4063ms
rtt min/avg/max/mdev = 0.193/6.241/30.412/12.085 ms
root@apu ~# ip netns exec VRF2 ping 10.10.10.1
PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data.
64 bytes from 10.10.10.1: icmp_seq=1 ttl=64 time=0.178 ms
64 bytes from 10.10.10.1: icmp_seq=2 ttl=64 time=0.195 ms
64 bytes from 10.10.10.1: icmp_seq=3 ttl=64 time=0.193 ms
64 bytes from 10.10.10.1: icmp_seq=4 ttl=64 time=0.194 ms
64 bytes from 10.10.10.1: icmp_seq=5 ttl=64 time=0.195 ms
^C
--- 10.10.10.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4085ms
rtt min/avg/max/mdev = 0.178/0.191/0.195/0.006 ms
root@apu ~#

```

Figure 35: Results from running the commands "ovs-vsctl show" and "ping" between both VRFs in phase 2



```

onos@root > flows
deviceId=o:000052ba3139a045, flowRuleCount=4
  id=1000000018471e, state=ADDED, bytes=1380, packets=33, duration=1199, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultTrafficTreatment{immediate}
  id=10000000cd5699, state=ADDED, bytes=0, packets=0, duration=0, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:lldp], treatment=DefaultTrafficTreatment{immediate}
  e=[OUTPUT:CONTROLLER], deferred[], transition=None, meter=[], cleared=true, StatTrigger=null, metadata=null
  id=10000fab50db, state=ADDED, bytes=0, packets=0, duration=199, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTrafficTreatment{immediate}
  e=[OUTPUT:CONTROLLER], deferred[], transition=None, meter=[], cleared=true, StatTrigger=null, metadata=null
  id=100000004c7ff3, state=ADDED, bytes=392, packets=4, duration=22, liveType=UNKNOWN, priority=5, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:ipv4], treatment=DefaultTrafficTreatment{immediate}
  e=[OUTPUT:CONTROLLER], deferred[], transition=None, meter=[], cleared=true, StatTrigger=null, metadata=null
onos@root > devices
Id=o:000052ba3139a045, available=true, localStatus=connected 200ms ago, role=MASTER, type=SWITCH, mfr=Nicira, Inc., hw=Open vSwitch, sw=3.3.90, serial=None, chassis=52ba3139a045, driver=default, channelId=192.168.137.254:43018, datapathDescription=None, managementAddress=192.168.137.254, protocol=OF_13
onos@root > hosts
Id=e:9E:EB:B6:8C:25:E2/None, mac=9E:EB:B6:8C:25:E2, locations=[of:000052ba3139a045/5], auxLocations=null, vlan=None, ip(s)=[10.10.10.2], innerVlan=None, outerTPID=unknown, provider=o:org.onosproject.provider.host, configured=false
Id=e:6:4F:3D:F6:8F:E5/None, mac=E6:4F:3D:F6:8F:E5, locations=[of:000052ba3139a045/6], auxLocations=null, vlan=None, ip(s)=[10.10.10.1], innerVlan=None, outerTPID=unknown, provider=o:org.onosproject.provider.host, configured=false
onos@root >

```

Figure 36: ONOS console with information relating to the devices, hosts and flows of the network in phase 2

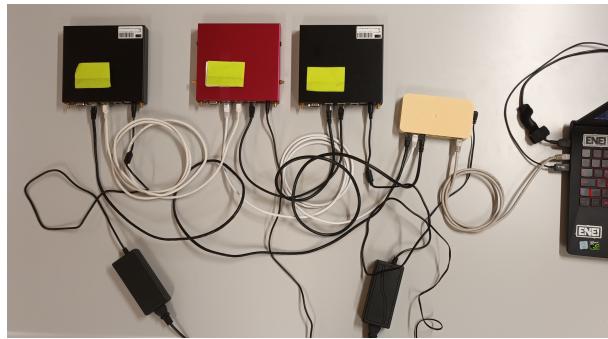


Figure 37: Setup from the third phase of the experiments with OvS

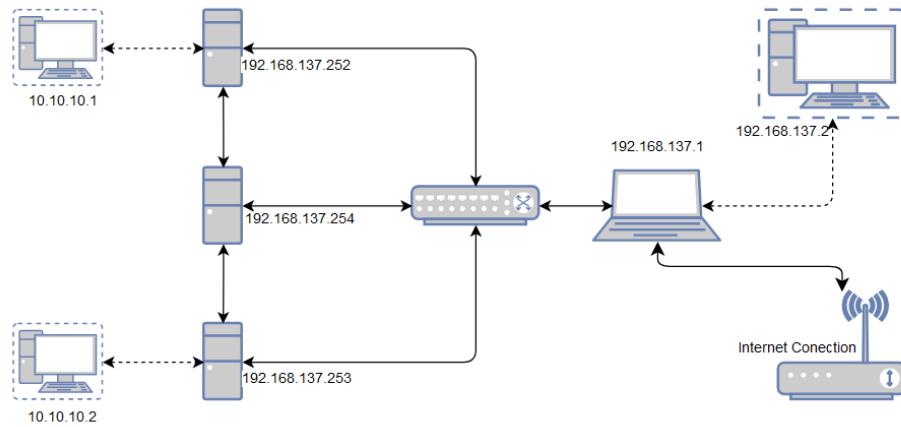


Figure 38: Diagram that represents the setup from the third phase of the experiments with OvS

- 10.10.10.1 - This IP address identifies the virtual interface that has been added to the VRF instance created in the 192.168.137.252 device.
- 10.10.10.2 - Identically to the last one, this IP address identifies the virtual interface that has been added to the VRF instance created in the 192.168.137.253 device.

When viewed through the lens of SDN, the configuration of this phase can be conceptualized as illustrated in Figure 39.

The configuration of the three instances of OvS that are running is Figure 40. This figure presents the configurations of the three instances of OvS installed on each individual device. These three instances of OvS are also acknowledged by the controller, as illustrated in Figure 41. Once again, these devices are connected to the controller via OpenFlow version 1.3, and the controller is aware of two hosts on the network. Figure 42 presents the flows that were injected into each device. These flows are identical across all three devices and are consistent with the flows injected in the previous phase, due to the fact that ONOS has precisely the same services activated. Connectivity between the two VRF is demonstrated in Figure 43. It illustrates the connectivity between devices 192.168.137.252 and 192.168.137.253, manifesting through the VFRs with IPs 10.10.10.1 and 10.10.10.2.

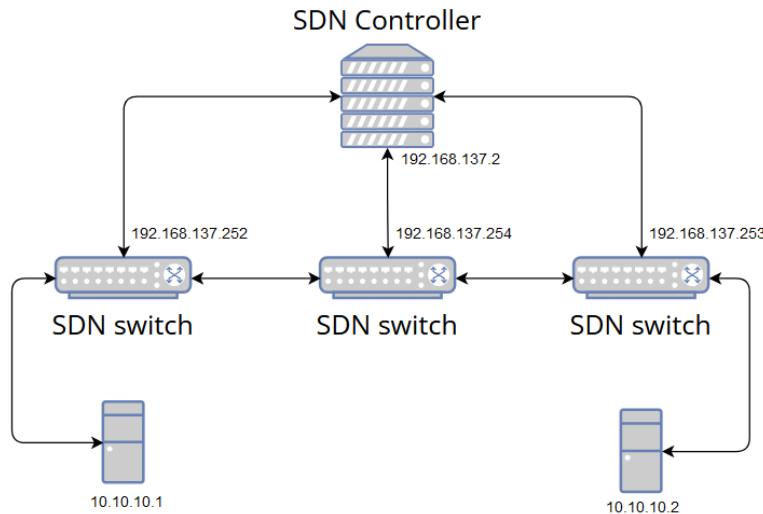


Figure 39: Diagram that depicts the **SDN** perspective of the setup from the third phase of the experiments with **OvS**

[apu] ~	[apu3d4] ~
<pre> root@apu ~# ovs-vsctl show 2bed74a0-d5cc-4f5f-a541-6d3cd80bdf06 Bridge ovs-br Controller "tcp:192.168.137.2:6653" is_connected: true Port ovs-br Interface ovs-br type: internal Port enp2s0 Interface enp2s0 Port enp3s0 Interface enp3s0 ovs_version: "3.3.90" root@apu ~# </pre>	<pre> apu@apu3d4 ~> sudo ovs-vsctl show e5d77087-d9c8-401d-b519-78d7a46cb8ac Bridge ovs-br Controller "tcp:192.168.137.2:6653" is_connected: true Port ovs-br Interface ovs-br type: internal Port veth2 Interface veth2 Port enp2s0 Interface enp2s0 ovs_version: "3.3.0" apu@apu3d4 ~> </pre>
[apu3d4] ~	[apu3d4] ~
	<pre> apu@apu3d4 ~> sudo ovs-vsctl show 6fcf0ddd-2846-47a4-bb3f-40e4f83f87a9 Bridge ovs-br Controller "tcp:192.168.137.2:6653" is_connected: true Port veth2 Interface veth2 Port ovs-br Interface ovs-br type: internal Port enp2s0 Interface enp2s0 ovs_version: "3.3.0" apu@apu3d4 ~> </pre>

Figure 40: Results from running the commands "ovs-vsctl show" in the devices running **OvS** in phase 3

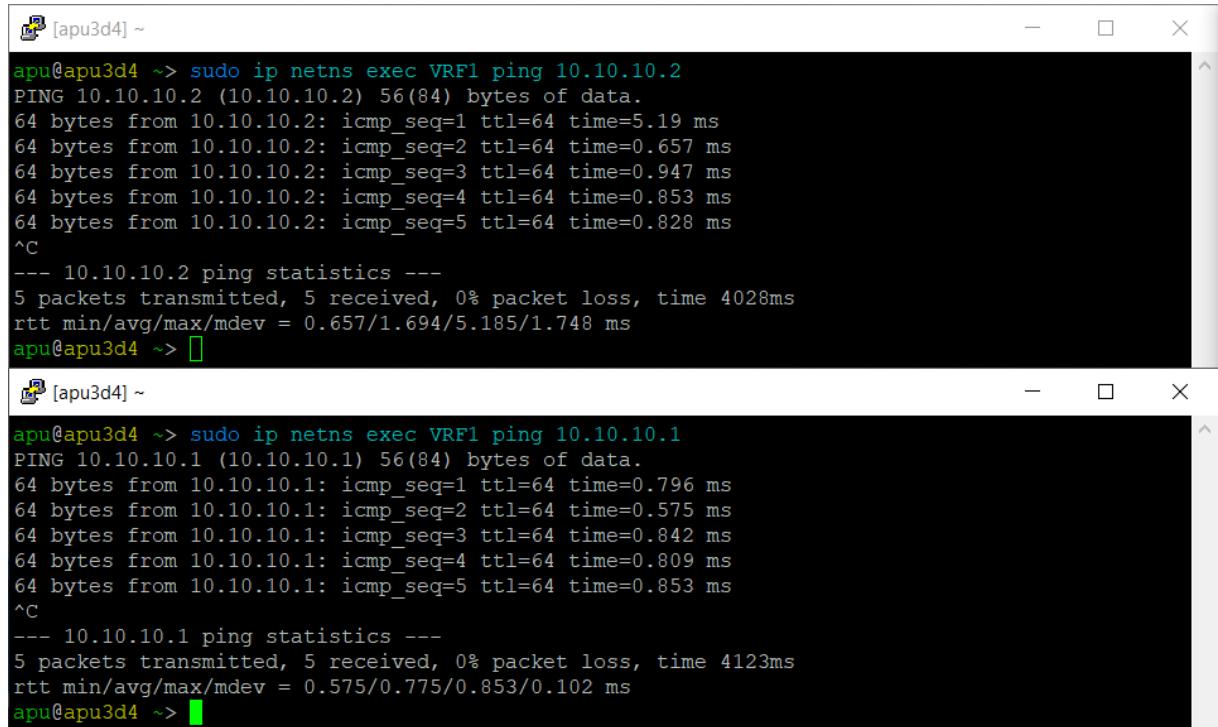
CHAPTER 6. FUNCTIONAL TESTS

```
lara@root > devices
18:45:12
id=0:0000000db9525e71, available=true, localStatus=connected 7m33s ago, role=MASTER, type=SWITCH, mfr=Nicira, Inc., hw=Open vSwitch, sw=3.3.0, serial=None, chassis=db9525e71, driver=default, channelId=192.168.137.254, datapathDescription=None, managementAddress=192.168.137.254, protocol=OF_13
id=0:0000000db962f865, available=true, localStatus=connected 6m59s ago, role=MASTER, type=SWITCH, mfr=Nicira, Inc., hw=Open vSwitch, sw=3.3.0, serial=None, chassis=db962f865, driver=default, channelId=192.168.137.253, datapathDescription=None, managementAddress=192.168.137.253, protocol=OF_13
id=0:0000000db962f845, available=true, localStatus=connected 7m03s ago, role=MASTER, type=SWITCH, mfr=Nicira, Inc., hw=Open vSwitch, sw=3.3.0, serial=None, chassis=db962f845, driver=default, channelId=192.168.137.252, datapathDescription=None, managementAddress=192.168.137.252, protocol=OF_13
lara@root > 18:45:14
id=0:02:7C:0B:1D:F9/None, mac=@02:7C:0B:1D:F9, locations=[of:0000000db962f865/2], auxLocations=null, vlan=None, ip(s)=[10.10.10.2], innerVlan=None, outerTpid=unknown, provider=of:org.onosproject.provider.host, configured=false
id=0:06:70:45:E8:92:2C/None, mac=@6:70:45:E8:92:2C, locations=[of:0000000db962f845/2], auxLocations=null, vlan=None, ip(s)=[10.10.10.1], innerVlan=None, outerTpid=unknown, provider=of:org.onosproject.provider.host, configured=false
lara@root > 18:45:17
```

Figure 41: ONOS console with information relating to the devices and hosts of the network in phase 3

```
lara@root > flows
10:45:39
deviceId=0:0000000db9525e71, flowRuleCount=4
  id=10000748489b5, state=ADDED, bytes=39744, packets=288, duration=475, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:lldp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=1000076c1686, state=ADDED, bytes=240, packets=4, duration=475, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=10000a658d1d5, state=ADDED, bytes=39744, packets=288, duration=475, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=10000c217f300, state=ADDED, bytes=196, packets=2, duration=148, liveType=UNKNOWN, priority=5, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:ipv4], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
deviceId=0:0000000db962f865, flowRuleCount=4
  id=1000000000000000, state=ADDED, bytes=102, packets=12, duration=445, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:lldp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=100007481add8, state=ADDED, bytes=708, packets=16, duration=440, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=10000d643dc, state=ADDED, bytes=1056, packets=142, duration=440, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=10000430876bf, state=ADDED, bytes=196, packets=2, duration=147, liveType=UNKNOWN, priority=5, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:ipv4], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
deviceId=0:0000000db962f845, flowRuleCount=4
  id=1000000303e000, state=ADDED, bytes=19872, packets=144, duration=445, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:lldp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=100005be2f55, state=ADDED, bytes=456, packets=10, duration=445, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=10000fd16ad08, state=ADDED, bytes=19872, packets=144, duration=445, liveType=UNKNOWN, priority=40000, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTrafficTreatment(immediate)
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
  id=10000276ae2f, state=ADDED, bytes=196, packets=2, duration=149, liveType=UNKNOWN, priority=5, tableId=0, appId=org.onosproject.core, selector=[ETH_TYPE:ipv4], treatment=DefaultTrafficTreatment(immediate]
  edata=[OUTPUT:CONTROLLER], deferred[], transition=None, meter[], cleared=true, StatTrigger=null, metadata=null
lara@root > 10:45:40
```

Figure 42: ONOS console with information relating to the flows of the network in phase 3



```
apu@apu3d4 ~> sudo ip netns exec VRF1 ping 10.10.10.2
PING 10.10.10.2 (10.10.10.2) 56(84) bytes of data.
64 bytes from 10.10.10.2: icmp_seq=1 ttl=64 time=5.19 ms
64 bytes from 10.10.10.2: icmp_seq=2 ttl=64 time=0.657 ms
64 bytes from 10.10.10.2: icmp_seq=3 ttl=64 time=0.947 ms
64 bytes from 10.10.10.2: icmp_seq=4 ttl=64 time=0.853 ms
64 bytes from 10.10.10.2: icmp_seq=5 ttl=64 time=0.828 ms
^C
--- 10.10.10.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4028ms
rtt min/avg/max/mdev = 0.657/1.694/5.185/1.748 ms
apu@apu3d4 ~> [ ]
```



```
apu@apu3d4 ~> sudo ip netns exec VRF1 ping 10.10.10.1
PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data.
64 bytes from 10.10.10.1: icmp_seq=1 ttl=64 time=0.796 ms
64 bytes from 10.10.10.1: icmp_seq=2 ttl=64 time=0.575 ms
64 bytes from 10.10.10.1: icmp_seq=3 ttl=64 time=0.842 ms
64 bytes from 10.10.10.1: icmp_seq=4 ttl=64 time=0.809 ms
64 bytes from 10.10.10.1: icmp_seq=5 ttl=64 time=0.853 ms
^C
--- 10.10.10.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4123ms
rtt min/avg/max/mdev = 0.575/0.775/0.853/0.102 ms
apu@apu3d4 ~> [ ]
```

Figure 43: Results from running the commands "ping" in the VRF between the two devices on the edge of the network in phase 3

6.2.4 Analysis and Interpretation

It was intended for the subsequent phase of the investigation to involve an attempt to introduce the ath9k patch and wireless antennas for use in conjunction with the [OvS](#) system. However, Phase 3 was the final phase of this examination. Through additional research it was discovered that [OvS](#) is unable to function with wireless interfaces, which is corroborated by a statement on the official wiki [92] indicating that [OvS](#) is not compatible with wireless interfaces. Furthermore, the testing phase revealed another significant limitation in the capabilities of [OvS](#). Despite its compatibility with OpenFlow, [OvS](#) is constrained by the traditional definition of a switch, thereby limiting its functionality to layer 2. This inherent restriction disables the system from parsing packets from different IP subnets. In conclusion, it was demonstrated that [OvS](#) is not ready to be used within the context of vehicular scenarios.

6.3 Experiment 2: testing the integration of [BMv2](#) and [P4](#)

The objective of this experiment is to evaluate the compatibility and practicality of integrating [BMv2](#) and [P4](#) with [VANET](#) technologies. The experiments were conducted using two 2-box apparatuses.

6.3.1 Testing phase 1: [P4](#) dataplane on devices

Phase 1 of this experiment will focus on assessing the fundamental operational capabilities of [BMv2](#), thereby offering an opportunity for familiarization with this environment. In order to accomplish this goal, the first step is installing [BMv2](#) on the target devices. It became evident that this process was more intricate than initially anticipated, due to the fact the southbound API of [BMv2](#) is, by default, not [P4Runtime](#), but rather Thrift [93]. Following an investigation into the matter, the existence of an option to install supplementary components to [BMv2](#) was identified, thereby enabling communication via the GRPC-based [P4Runtime](#). The modification was implemented, allowing for the establishment of a connection between the devices and [ONOS](#).

The [P4/ONOS](#) application developed in section 5.4 was used. [P4](#) code is compiled by the p4c [94], the code is automatically injected into the devices by [ONOS](#), thereby eliminating the need for manual transfer of the file on each occasion. Once the requisite setup was completed, two [VRFs](#) were created, one on each device. Subsequently, the aforementioned [VRF](#) instances were linked to their corresponding devices via the use of virtual interfaces. Thereafter, [BMv2](#) was initiated, with one virtual and one physical interface allocated to it. Next, [ONOS](#) was activated, along with the requisite [ONOS/P4](#) application and all associated dependencies. At last, the configuration was applied to the controller, thereby informing the controller of the devices and providing it with the necessary information to perform its functions.

The test configuration is illustrated in Figure 44, which was also used as the basis for the diagram seen in Figure 45. A description of said diagram is in the following itemized list:

- 192.168.137.1 - This [IP](#) address identifies the author's personal laptop computer. It serves as the central coordinating hub for the entire network, facilitating the observation and monitoring of all

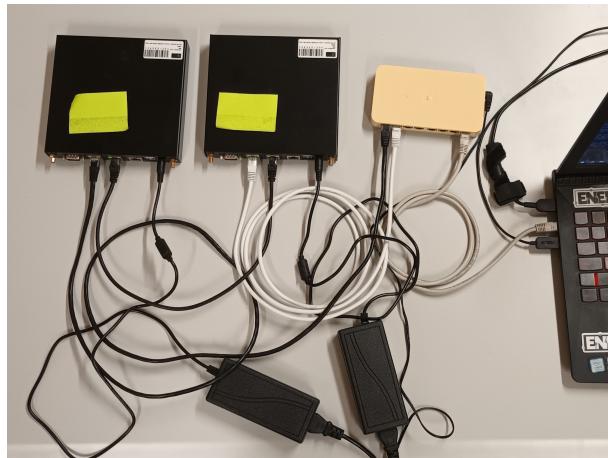


Figure 44: Setup from the first phase of the experiments with BMv2

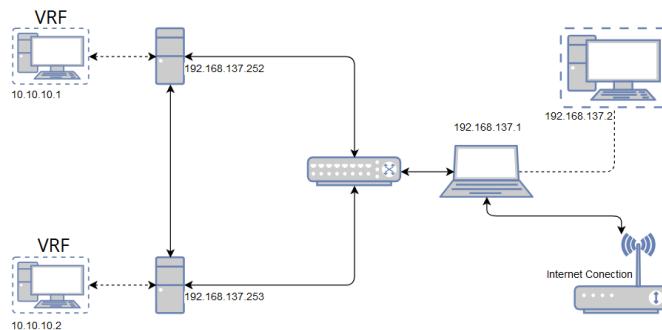


Figure 45: Diagram that represents the setup from the first phase of the experiments with BMv2

network activity.

- 192.168.137.2 - This address represents a virtual machine that is hosted on the author's laptop via the Oracle VM VirtualBox virtualization platform. The virtual environment is configured with Debian 12 and is primarily intended for the deployment of ONOS.
- 192.168.137.252 - This corresponds to a device version 2. It has a fresh installation of the Ubuntu Server operating system and is running BMv2 with grpc.
- 192.168.137.253 - This device is virtually identical to the device 192.168.137.252.
- 10.10.10.1 - This IP address identifies the virtual interface that has been added to the VRF instance created in the 192.168.137.252 device.
- 10.10.10.2 - Identically to the last one, this IP address identifies the virtual interface that has been added to the VRF instance created in the 192.168.137.253 device.

When viewed through the lens of SDN, the physical configuration can be conceptualized as illustrated in Figure 46.

Figure 47 shows BMv2 running in both devices. It displays the initial debug messages that were transmitted by the program in order to create the P4 tables.

6.3. EXPERIMENT 2: TESTING THE INTEGRATION OF BMV2 AND P4

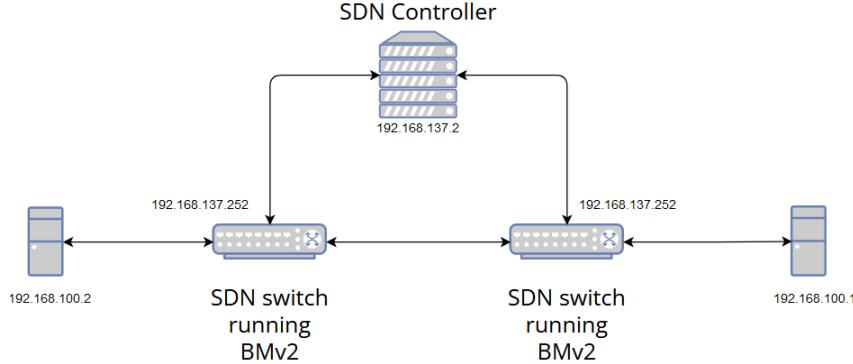


Figure 46: Diagram that depicts the [SDN](#) perspective of the setup from the first phase of the experiments with [BMv2](#) in phase 1

```

apu@apu3d4: ~$ sudo simple_switch_grpc --no-p4 --device-id 1 -i 1@emp2s0 -i 2@veth2 --log-cons
...
[19:12:40.657] [bmv2] [D] [thread 1491] Adding interface emp2s0 as port 1
Adding interface veth2 as port 2
Server listening on 192.168.137.253:9559
[19:13:31.751] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main240': ma
in240 -
[19:13:31.751] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main254': ma
in254 -
[19:13:31.751] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main256': ma
in256 -
[19:13:31.751] [bmv2] [D] [thread 1498] Set default default entry for table 'IngressPipeImpl.
acl_table': NoAction -
[19:13:31.751] [bmv2] [D] [thread 1498] Set default default entry for table 'IngressPipeImpl.
l2_exact_table': NoAction -
[19:13:31.763] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main301': ma
in301 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main319': ma
in319 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main323': ma
in323 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main330': ma
in330 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main336': ma
in336 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main342': ma
in342 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main349': ma
in349 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main358': ma
in358 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main367': ma
in367 -
[19:13:31.764] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main370': ma
in370 -
[19:13:31.765] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main373': ma
in373 -
[19:13:31.765] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main376': ma
in376 -
[19:13:31.765] [bmv2] [D] [thread 1498] Set default default entry for table 'tbl_main378': ma
in378 -
[19:13:31.770] [bmv2] [D] [thread 1498] simple_switch target has been notified of a config sw
ap
[19:13:31.775] [bmv2] [D] [thread 1506] [0.0] [ctx 0] Processing packet received on port 1
[19:13:31.775] [bmv2] [D] [thread 1506] [0.0] [ctx 0] Parser 'parser': start
[19:13:31.775] [bmv2] [D] [thread 1506] [0.0] [ctx 0] Parser 'parser' entering state 'start' v

apu@apu3d4: ~$ sudo simple_switch_grpc --no-p4 --device-id 2 -i 2@veth2 -i 1@emp2s0 --log-cons
...
[19:11:48.100] [bmv2] [D] [thread 1322] Adding interface emp2s0 as port 1
Adding interface veth2 as port 2
Server listening on 192.168.137.253:9559
[19:12:44.903] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main240': ma
in240 -
[19:12:44.903] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main254': ma
in254 -
[19:12:44.903] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main256': ma
in256 -
[19:12:44.903] [bmv2] [D] [thread 1339] Set default default entry for table 'IngressPipeImpl.
acl_table': NoAction -
[19:12:44.903] [bmv2] [D] [thread 1339] Set default default entry for table 'IngressPipeImpl.
l2_exact_table': NoAction -
[19:12:44.914] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main301': ma
in301 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main319': ma
in319 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main323': ma
in323 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main330': ma
in330 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main336': ma
in336 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main342': ma
in342 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main349': ma
in349 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main358': ma
in358 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main367': ma
in367 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main370': ma
in370 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main373': ma
in373 -
[19:12:44.915] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main376': ma
in376 -
[19:12:44.916] [bmv2] [D] [thread 1339] Set default default entry for table 'tbl_main378': ma
in378 -
[19:12:44.922] [bmv2] [D] [thread 1348] simple_switch target has been notified of a config sw
ap
[19:12:44.924] [bmv2] [D] [thread 1348] [0.0] [ctx 0] Processing packet received on port 1
[19:12:44.924] [bmv2] [D] [thread 1348] [0.0] [ctx 0] Parser 'parser': start
[19:12:44.924] [bmv2] [D] [thread 1348] [0.0] [ctx 0] Parser 'parser' entering state 'start' v

```

Figure 47: Consoles running [BMv2](#) on both devices

The list of devices and hosts known to the controller is shown in Figure 48, comprising two devices and two hosts. The devices are interfaced with the controller via P4Runtime. The list of applications activated in [ONOS](#) is presented in Figure 49. This list includes several applications that have been activated to enable the P4runtime to function with [BMv2](#). Furthermore, the final application listed, designated "SDVN Test App", is the [P4/ONOS](#) application developed for this phase. The flows that have been injected into the two devices are illustrated in Figure 50, and as expected, these flows are identical on both devices. It is also possible to identify the table in which each flow is applied. The flows designated for the "acl_table" are inserted by [ONOS](#) native applications. In contrast, the other three rules, injected into the "l2_exact_table", are all inserted by our application. These include a default traffic rule, a broadcast rule and an unicast rule.

CHAPTER 6. FUNCTIONAL TESTS

```
karaf@root > devices
id=device:obu1, available=true, local-status=connected 3m17s ago, role=MASTER, type=SWITCH, mfr=p4.org, hw=master, sw=master, serial=unknown, chassis=0, driver=bmv2:org.onosproject.pipelines.sdvn, gridx=200.0, g
ridy=600.0, locType=grid, managementAddress=grpc://192.168.137.252:95597device_id=1, name=device:obu1, p4deviceId=1, protocol=P4Runtime
id=device:obu2, available=true, local-status=connected 3m17s ago, role=MASTER, type=SWITCH, mfr=p4.org, hw=master, sw=master, serial=unknown, chassis=0, driver=bmv2:org.onosproject.pipelines.sdvn, gridx=800.0, g
ridy=600.0, locType=grid, managementAddress=grpc://192.168.137.253:95597device_id=2, name=device:obu2, p4deviceId=2, protocol=P4Runtime
karaf@root > hosts
id=A2:0B:32:2D:C7:9D:None, mac=A2:0B:32:2D:C7:9D, locations=[device:obu2/2], auxLocations=null, vlan=None, ip(s)=[10.10.10.2], gridx=200.0, gridy=700.0, latitude=null, locType=grid, longitude=null, name=host2-25
3, innerVlan=None, outerTpid=unknown, provider=of:org.onosproject.provider.host, configured=false
id=A6:70:45:E8:92:2C:None, mac=A6:70:45:E8:92:2C, locations=[device:obu1/2], auxLocations=null, vlan=None, ip(s)=[10.10.10.1], gridx=800.0, gridy=700.0, latitude=null, locType=grid, longitude=null, name=host2-25
2, innerVlan=None, outerTpid=unknown, provider=of:org.onosproject.provider.host, configured=false
karaf@root > 19:14:54
19:14:54
```

Figure 48: ONOS console with information relating to the devices and hosts of the network in phase 1

```
karaf@root > apps -a -s
* 13 org.onosproject.protocols.grpc          3.0.0.SNAPSHOT gRPC Protocol Subsystem
* 14 org.onosproject.protocols.gnmi         3.0.0.SNAPSHOT gNMI Protocol Subsystem
* 15 org.onosproject.generaldeviceprovider  3.0.0.SNAPSHOT General Device Provider
* 16 org.onosproject.drivers.gnmi           3.0.0.SNAPSHOT gNMI Drivers
* 19 org.onosproject.optical-model          3.0.0.SNAPSHOT Optical Network Model
* 20 org.onosproject.drivers                3.0.0.SNAPSHOT Default Drivers
* 40 org.onosproject.protocols.p4runtime    3.0.0.SNAPSHOT P4Runtime Protocol Subsystem
* 41 org.onosproject.p4runtime              3.0.0.SNAPSHOT P4Runtime Provider
* 42 org.onosproject.drivers.p4runtime     3.0.0.SNAPSHOT P4Runtime Drivers
* 49 org.onosproject.pipelines.basic       3.0.0.SNAPSHOT Basic Pipelines
* 60 org.onosproject.protocols.gnoi         3.0.0.SNAPSHOT gNOI Protocol Subsystem
* 61 org.onosproject.drivers.gnoi          3.0.0.SNAPSHOT gNOI Drivers
* 62 org.onosproject.drivers.stratum       3.0.0.SNAPSHOT Stratum Drivers
* 74 org.onosproject.hostprovider          3.0.0.SNAPSHOT Host Location Provider
* 75 org.onosproject.lldpprovider         3.0.0.SNAPSHOT LLDP Link Provider
* 76 org.onosproject.openflow-base         3.0.0.SNAPSHOT OpenFlow Base Provider
* 77 org.onosproject.openflow              3.0.0.SNAPSHOT OpenFlow Provider Suite
* 126 org.onosproject.drivers.bmv2         3.0.0.SNAPSHOT BMv2 Drivers
* 170 org.onosproject.gui2                 3.0.0.SNAPSHOT ONOS GUI2
* 180 org.onosproject.sdvn                1.0.0.SNAPSHOT SDVN Test App
karaf@root > 19:14:56
19:14:56
```

Figure 49: ONOS console with information relating to the applications activated on the controller in phase 1

```
karaf@root > flows
deviceId=device:obu1, flowRuleCount=6
  id=1000030540bf6, state=ADDED, bytes=0, packets=0, duration=256, liveType=UNKNOWN, priority=40000, tableId=ingressPipeImpl.acl_table, appid=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTr
afficTreatment([immediate=[ingressPipeImpl.clone_to_cpu(session_id=0x4)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=1000037c7c983, state=ADDED, bytes=126, packets=3, duration=256, liveType=UNKNOWN, priority=40000, tableId=ingressPipeImpl.acl_table, appid=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultT
rafficTreatment([immediate=[ingressPipeImpl.clone_to_cpu(session_id=0x4)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=100009fa9d6a, state=ADDED, bytes=0, packets=0, duration=256, liveType=UNKNOWN, priority=40000, tableId=ingressPipeImpl.acl_table, appid=org.onosproject.core, selector=[ETH_TYPE:lldp], treatment=DefaultTr
afficTreatment([immediate=[ingressPipeImpl.clone_to_cpu(session_id=0x4)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=b40000748b7452, state=ADDED, bytes=0, packets=0, duration=256, liveType=UNKNOWN, priority=10, tableId=ingressPipeImpl.l2_exact_table, appid=org.onosproject.sdvn, selector=[], treatment=DefaultTrafficTreatment
  id=100007e0000000000, state=ADDED, bytes=0, duration=20, packets=2, duration=256, liveType=UNKNOWN, priority=10, tableId=ingressPipeImpl.l2_exact_table, appid=org.onosproject.sdvn, selector=[hdr.ethernet.dst_addr=0xffffffff
ffff], treatment=DefaultTrafficTreatment(immediate=[ingressPipeImpl.set_multicast_group(id=0xff, switch_id.value=0x04)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=b400006f89c88, state=ADDED, bytes=134, packets=13, duration=122, liveType=UNKNOWN, priority=10, tableId=ingressPipeImpl.l2_exact_table, appid=org.onosproject.sdvn, selector=[hdr.ethernet.dst_addr=0xa670
45e092c], treatment=DefaultTrafficTreatment(immediate=[ingressPipeImpl.set_egress_port(port_num=0x2)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
deviceId=device:obu2, flowRuleCount=6
  id=100002772e95a, state=ADDED, bytes=0, duration=256, liveType=UNKNOWN, priority=40000, tableId=ingressPipeImpl.acl_table, appid=org.onosproject.core, selector=[ETH_TYPE:lldp], treatment=DefaultTr
afficTreatment([immediate=[ingressPipeImpl.clone_to_cpu(session_id=0x4)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=100005f040024, state=ADDED, bytes=0, duration=256, liveType=UNKNOWN, priority=40000, tableId=ingressPipeImpl.acl_table, appid=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTr
afficTreatment([immediate=[ingressPipeImpl.clone_to_cpu(session_id=0x5)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=10000c92dc35, state=ADDED, bytes=126, packets=3, duration=256, liveType=UNKNOWN, priority=40000, tableId=ingressPipeImpl.acl_table, appid=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultT
rafficTreatment([immediate=[ingressPipeImpl.clone_to_cpu(session_id=0x5)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=b400007ce1638, state=ADDED, bytes=0, packets=0, duration=256, liveType=UNKNOWN, priority=10, tableId=ingressPipeImpl.l2_exact_table, appid=org.onosproject.sdvn, selector=[hdr.ethernet.dst_addr=0xffffffff
ffff], treatment=DefaultTrafficTreatment(immediate=[ingressPipeImpl.set_multicast_group(id=0xff, switch_id.value=0x05)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null]
  id=b400005eeca2, state=ADDED, bytes=22, packets=4, duration=256, liveType=UNKNOWN, priority=10, tableId=ingressPipeImpl.l2_exact_table, appid=org.onosproject.sdvn, selector=[hdr.ethernet.dst_addr=0xa20b
322dc79d], treatment=DefaultTrafficTreatment(immediate=[ingressPipeImpl.set_egress_port(port_num=0x2)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
karaf@root > 19:15:30
19:15:30
```

Figure 50: ONOS console with information relating to the flows of the network in phase 1

The figure consists of two vertically stacked terminal windows. Both windows have a title bar with a small icon, the text '[apu3d4] ~', and standard window control buttons (minimize, maximize, close). The top window shows the command 'sudo ip netns exec VRFl ping 10.10.10.1' being run. The output shows four ICMP packets being sent to 10.10.10.1 with TTL=64, and the ping statistics: 4 packets transmitted, 0% packet loss, time 3005ms, rtt min/avg/max/mdev = 14.563/15.200/15.752/0.450 ms. The bottom window shows the command 'sudo ip netns exec VRFl ping 10.10.10.2' being run. The output shows four ICMP packets being sent to 10.10.10.2 with TTL=64, and the ping statistics: 4 packets transmitted, 0% packet loss, time 3005ms, rtt min/avg/max/mdev = 15.564/16.079/16.379/0.311 ms.

```

[apu@apu3d4 ~]
apu@apu3d4 ~> sudo ip netns exec VRFl ping 10.10.10.1
PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data.
64 bytes from 10.10.10.1: icmp_seq=1 ttl=64 time=14.6 ms
64 bytes from 10.10.10.1: icmp_seq=2 ttl=64 time=15.0 ms
64 bytes from 10.10.10.1: icmp_seq=3 ttl=64 time=15.8 ms
64 bytes from 10.10.10.1: icmp_seq=4 ttl=64 time=15.5 ms
^C
--- 10.10.10.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3005ms
rtt min/avg/max/mdev = 14.563/15.200/15.752/0.450 ms
apu@apu3d4 ~> [green square]

[apu@apu3d4 ~]
apu@apu3d4 ~> sudo ip netns exec VRFl ping 10.10.10.2
PING 10.10.10.2 (10.10.10.2) 56(84) bytes of data.
64 bytes from 10.10.10.2: icmp_seq=1 ttl=64 time=16.1 ms
64 bytes from 10.10.10.2: icmp_seq=2 ttl=64 time=15.6 ms
64 bytes from 10.10.10.2: icmp_seq=3 ttl=64 time=16.3 ms
64 bytes from 10.10.10.2: icmp_seq=4 ttl=64 time=16.4 ms
^C
--- 10.10.10.2 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3005ms
rtt min/avg/max/mdev = 15.564/16.079/16.379/0.311 ms
apu@apu3d4 ~> [green square]

```

Figure 51: Results from running the commands "ping" in the [VRF](#) between the two devices on the network in phase 1

Figure 51 illustrates the connectivity between devices 192.168.137.252 and 192.168.137.253, manifesting through the [VRFs](#) with [IPs](#) 10.10.10.1 and 10.10.10.2.

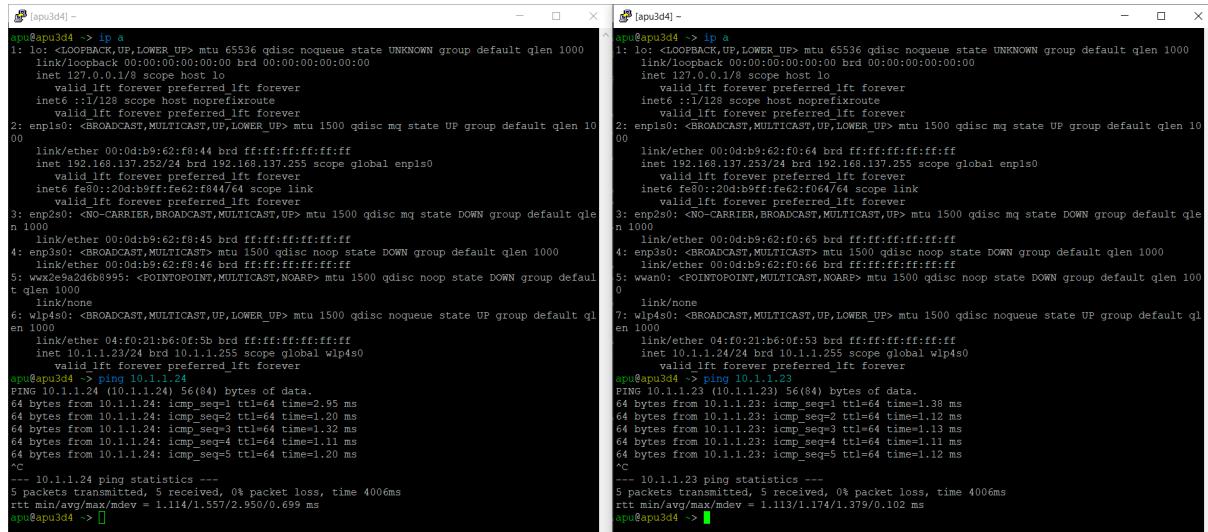
6.3.2 Testing phase 2: P4 control over wireless interfaces

Phase two of the experiments with [BMv2](#) represents a significant advancement over the initial phase, as it introduces the crucial element of wireless communication. Accordingly, the ath9k patch was implemented on the devices to enable communication at the [IEEE](#) 802.11p standard. Figure 52 displays the lists of interfaces on the devices, where it is possible to identify the wireless interfaces called "w1p4s0" on both devices. Furthermore, it demonstrates the connectivity between said devices.

As a result of this, two changes had to be made. The first of these amendments required the activation of the [BMv2](#) with the wireless interface. The second required modification of the [P4/ONOS](#) applications and configuration to accommodate the distinctive characteristics of the wireless antennas. For any message to be successfully received by a wireless interface, that destination Media Access Control address must either be the address of the interface or the broadcast address. It was thus decided to pursue the latter course of action, which entailed changing the addresses of the messages sent via the wireless interface to the broadcast address. The original address is stored within the packet and inserted when required.

The test configuration is illustrated in Figure 53, which was also used as the basis for the diagram seen in Figure 54 and described in the following itemized list:

CHAPTER 6. FUNCTIONAL TESTS



The image shows two terminal windows side-by-side. Both windows are titled 'apu@apu3d4 ~'.

Terminal 1 (Left):

```
apu@apu3d4 ~ > ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
        inet6 ::/128 scope host noprefixroute
            valid_lft forever preferred_lft forever
2: enp1s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
    link/ether 00:0d:b9:62:f8:44 brd ff:ff:ff:ff:ff:ff
    inet 192.168.137.252/24 brd 192.168.137.255 scope global enp1s0
        valid_lft forever preferred_lft forever
    inet6 fe80::20d:b9ff:fe62:f8f4%enp1s0/64 scope link
        valid_lft forever preferred_lft forever
3: enp2s0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc mq state DOWN group default qlen 1000
    link/ether 00:0d:b9:62:f8:45 brd ff:ff:ff:ff:ff:ff
4: enp3s0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/ether 00:0d:b9:62:f8:46 brd ff:ff:ff:ff:ff:ff
5: wwan0: <POINTOPOINT,MULTICAST,NOARP> mtu 1500 qdisc noop state DOWN group default qlen 1000
    link/none
6: wlp4s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 04:f0:21:b6:0f:5b brd ff:ff:ff:ff:ff:ff
    inet 10.1.1.23/24 brd 10.1.1.255 scope global wlp4s0
        valid_lft forever preferred_lft forever
apu@apu3d4 ~ > ping 10.1.1.24
PING 10.1.1.24 (10.1.1.24) 56(84) bytes of data.
64 bytes from 10.1.1.24: icmp_seq=1 ttl=64 time=2.95 ms
64 bytes from 10.1.1.24: icmp_seq=2 ttl=64 time=1.20 ms
64 bytes from 10.1.1.24: icmp_seq=3 ttl=64 time=1.38 ms
64 bytes from 10.1.1.24: icmp_seq=4 ttl=64 time=1.11 ms
64 bytes from 10.1.1.24: icmp_seq=5 ttl=64 time=1.20 ms
```
--- 10.1.1.24 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4006ms
rtt min/avg/max/mdev = 1.114/1.557/2.950/0.699 ms
apu@apu3d4 ~ >
```

**Terminal 2 (Right):**

```
apu@apu3d4 ~ > ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
 link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
 inet 127.0.0.1/8 scope host lo
 valid_lft forever preferred_lft forever
 inet6 ::/128 scope host noprefixroute
 valid_lft forever preferred_lft forever
2: enp1s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP group default qlen 1000
 link/ether 00:0d:b9:62:f0:64 brd ff:ff:ff:ff:ff:ff
 inet 192.168.137.252/24 brd 192.168.137.255 scope global enp1s0
 valid_lft forever preferred_lft forever
 inet6 fe80::20d:b9ff:fe62:f0f4%enp1s0/64 scope link
 valid_lft forever preferred_lft forever
3: enp2s0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc mq state DOWN group default qlen 1000
 link/ether 00:0d:b9:62:f0:65 brd ff:ff:ff:ff:ff:ff
4: enp3s0: <BROADCAST,MULTICAST> mtu 1500 qdisc noop state DOWN group default qlen 1000
 link/ether 00:0d:b9:62:f0:66 brd ff:ff:ff:ff:ff:ff
5: wwan0: <POINTOPOINT,MULTICAST,NOARP> mtu 1500 qdisc noop state DOWN group default qlen 1000
 link/none
6: wlp4s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
 link/ether 04:f0:21:b6:0f:53 brd ff:ff:ff:ff:ff:ff
 inet 10.1.1.23/24 brd 10.1.1.255 scope global wlp4s0
 valid_lft forever preferred_lft forever
apu@apu3d4 ~ > ping 10.1.1.23
PING 10.1.1.23 (10.1.1.23) 56(84) bytes of data.
64 bytes from 10.1.1.23: icmp_seq=1 ttl=64 time=1.38 ms
64 bytes from 10.1.1.23: icmp_seq=2 ttl=64 time=1.12 ms
64 bytes from 10.1.1.23: icmp_seq=3 ttl=64 time=1.11 ms
64 bytes from 10.1.1.23: icmp_seq=4 ttl=64 time=1.11 ms
64 bytes from 10.1.1.23: icmp_seq=5 ttl=64 time=1.12 ms
```
--- 10.1.1.23 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4006ms
rtt min/avg/max/mdev = 1.113/1.174/1.379/0.102 ms
apu@apu3d4 ~ >
```

Figure 52: Results from running the commands "ip a" and "ping" between the two devices on the network using the IP address of the wireless interfaces

- 192.168.137.1 - This IP address identifies the author's personal laptop computer. It serves as the central coordinating hub for the entire network, facilitating the observation and monitoring of all network activity.
- 192.168.137.2 - This address represents a virtual machine that is hosted on the author's laptop via the Oracle VM VirtualBox virtualization platform. The virtual environment is configured with Debian 12 and is primarily intended for the deployment of ONOS.
- 192.168.137.252 - This device corresponds to the version 2 of the devices. The device has a fresh installation of the Ubuntu Server operating system and is running BMv2 with grpc. Furthermore, the device is equipped with an antenna capable of communicating at frequencies up to 6 GHz and has been applied with the ath9k patch to the WiFi driver, thereby enabling communication using IEEE 802.11p.
- 192.168.137.253 - This device is virtually identical to the device 192.168.137.252.
- 10.10.10.1 - This IP address identifies the virtual interface that has been added to the VRF instance created in the 192.168.137.252 device.
- 10.10.10.2 - Identically to the last one, this IP address identifies the virtual interface that has been added to the VRF instance created in the 192.168.137.253 device.

When viewed through the lens of SDN, the physical configuration can be conceptualized as illustrated in Figure 55.

Figure 56 demonstrates that the BMv2 program is operational in both devices, and that the programs have successfully connected to the wireless interfaces.

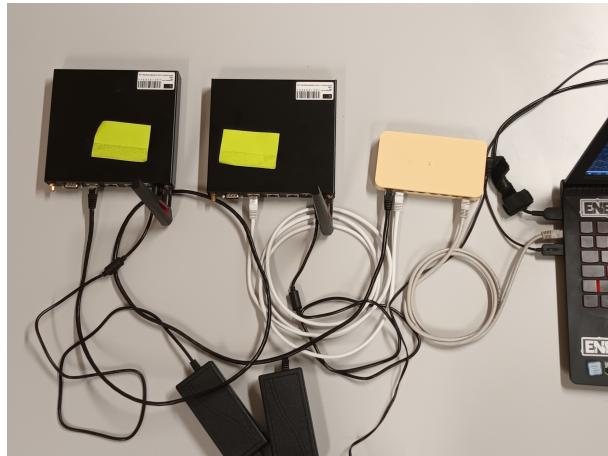


Figure 53: Setup from the second phase of the experiments with BMv2

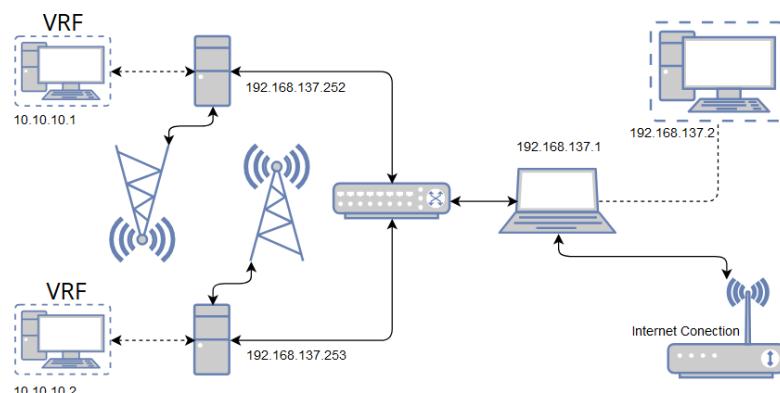


Figure 54: Diagram that represents the setup from the second phase of the experiments with BMv2

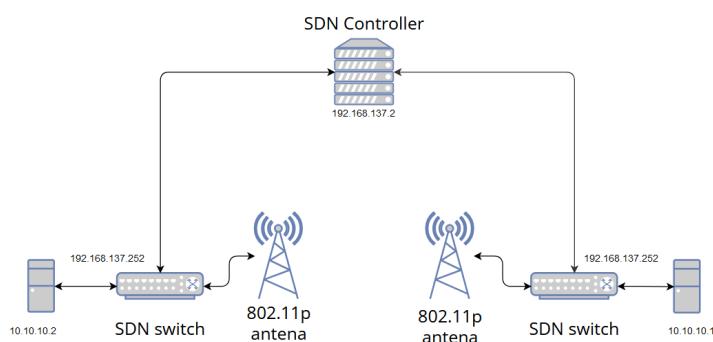


Figure 55: Diagram that depicts the SDN perspective of the setup from the second phase of the experiments with BMv2

CHAPTER 6. FUNCTIONAL TESTS

```
[apu@apu3d4] sudo simple_switch_g ~
apu@apu3d4 ~> sudo simple_switch_grpc --no-p4 --device-id 1 -i 1@wlp4s0 -i 2@veth2 --log-console -- --grpc-server-addr 192.168.137.252:9559 --cpu-port 200
[sudo] password for apu:
Calling target program-options parser
Adding interface wlp4s0 as port 1
[19:29:39.938] [bmv2] [D] [thread 1714] Adding interface wlp4s0 as port 1
Adding interface veth2 as port 2
[19:29:40.146] [bmv2] [D] [thread 1714] Adding interface veth2 as port 2
[19:29:40.210] [bmv2] [E] [thread 1714] Add port operation failed
Server listening on 192.168.137.252:9559

[apu@apu3d4] sudo simple_switch_g ~
apu@apu3d4 ~> sudo simple_switch_grpc --no-p4 --device-id 2 -i 2@veth2 -i 1@wlp4s0 --log-console -- --grpc-server-addr 192.168.137.253:9559 --cpu-port 200
[sudo] password for apu:
Calling target program-options parser
Adding interface wlp4s0 as port 1
[19:28:59.123] [bmv2] [D] [thread 1704] Adding interface wlp4s0 as port 1
Adding interface veth2 as port 2
[19:28:59.320] [bmv2] [D] [thread 1704] Adding interface veth2 as port 2
[19:28:59.382] [bmv2] [E] [thread 1704] Add port operation failed
Server listening on 192.168.137.253:9559
```

Figure 56: Consoles running BMv2 on both devices in phase 2

```
curl@root > devices
id=device:obu1, available=true, local-status=connected 1m57s ago, role=MASTER, type=SWITCH, mfr=p4.org, hw=master, sw=master, serial=unknown, chassis=0, driver=bmv2:org.onosproject.pipelines.sdvn, gridx=200.0, g
ridy=0.0, locType=grid, managementAddress=grpc://192.168.137.252:9559, device_id=1, name=device:obu1, p4DeviceId=1, protocol=p4Runtime
id=device:obu2, available=true, local-status=connected 1m46s ago, role=MASTER, type=SWITCH, mfr=p4.org, hw=master, sw=master, serial=unknown, chassis=0, driver=bmv2:org.onosproject.pipelines.sdvn, gridx=800.0, g
ridy=0.0, locType=grid, managementAddress=grpc://192.168.137.253:9559, device_id=2, name=device:obu2, p4DeviceId=2, protocol=p4Runtime
curl@root > hosts
1d:A2:0B:32:20:D7:90:None, mac=A2:0B:32:20:D7:90, locations=[device:obu2/2], auxLocations=null, vlan=None, ip(s)=[10.10.10.2], gridx=200.0, gridy=700.0, latitude=null, locType=grid, longitude=null, name=host1-25
3, innerVlan=None, outerTID=unknown, provider=of:org.onosproject.provider.host, configured=false
1d:A6:78:45:E8:92:2C:None, mac=A6:78:45:E8:92:2C, locations=[device:obu1/2], auxLocations=null, vlan=None, ip(s)=[10.10.10.1], gridx=800.0, gridy=700.0, latitude=null, locType=grid, longitude=null, name=host2-25
2, innerVlan=None, outerTID=unknown, provider=of:org.onosproject.provider.host, configured=false
curl@root >
```

Figure 57: ONOS console with information relating to the devices and hosts of the network in phase 2

```
curl@root > flows
deviceId=device:obu1, flowRuleCount=6
  id=10000305408f, state=ADDED, bytes=0, packets=0, duration=133, liveType=UNKNOWN, priority=40000, tableId=IngressPipeImpl.ac1_table, appId=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTr
afficTreatment([IngressPipeImpl.clone_to_cpu(session_id=0x404)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=1000037c69ab, state=ADDED, bytes=10, duration=133, liveType=UNKNOWN, priority=40000, tableId=IngressPipeImpl.ac1_table, appId=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultT
rafficTreatment([IngressPipeImpl.clone_to_cpu(session_id=0x404)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=b400007cf895a, state=ADDED, bytes=9, duration=133, liveType=UNKNOWN, priority=40000, tableId=IngressPipeImpl.ac1_table, appId=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultTr
rafficTreatment([IngressPipeImpl.clone_to_cpu(session_id=0x404)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=b40000749b452, state=ADDED, bytes=0, duration=13, liveType=UNKNOWN, priority=10, tableId=IngressPipeImpl.l2_exact_table, appId=org.onosproject.sdn, selector=[], treatment=DefaultTrafficTreat
ment([Immediate], add_switch_id=[port_num=0x1], switch_id.value=0x64), deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=b400007c06038, state=ADDED, bytes=354, packets=8, duration=133, liveType=UNKNOWN, priority=10, tableId=IngressPipeImpl.l2_exact_table, appId=org.onosproject.sdn, selector=[hdr.ethernet.dst_addr=ffffffff
ffff], treatment=DefaultTrafficTreatment([Immediate], set_multicast_group_id=0xff, switch_id.value=0x64), deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=b400005f89c88, state=ADDED, bytes=3414, packets=34, duration=127, liveType=UNKNOWN, priority=10, tableId=IngressPipeImpl.l2_exact_table, appId=org.onosproject.sdn, selector=[hdr.ethernet.dst_addr=0xa070
45e992c2], treatment=DefaultTrafficTreatment([Immediate], [IngressPipeImpl.set_egress_port(port_num=0x2)]], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
deviceId=device:obu2, flowRuleCount=6
  id=100002f72e95a, state=ADDED, bytes=0, packets=0, duration=123, liveType=UNKNOWN, priority=40000, tableId=IngressPipeImpl.ac1_table, appId=org.onosproject.core, selector=[ETH_TYPE:lldp], treatment=DefaultTr
afficTreatment([IngressPipeImpl.clone_to_cpu(session_id=0x655)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=100005f040024, state=ADDED, bytes=0, packets=0, duration=123, liveType=UNKNOWN, priority=40000, tableId=IngressPipeImpl.ac1_table, appId=org.onosproject.core, selector=[ETH_TYPE:bddp], treatment=DefaultTr
afficTreatment([Immediate], [IngressPipeImpl.clone_to_cpu(session_id=0x655)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=10000c92dc95, state=ADDED, bytes=20, duration=123, liveType=UNKNOWN, priority=40000, tableId=IngressPipeImpl.ac1_table, appId=org.onosproject.core, selector=[ETH_TYPE:arp], treatment=DefaultT
rafficTreatment([IngressPipeImpl.clone_to_cpu(session_id=0x655)], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=b400007ce1638, state=ADDED, bytes=0, packets=2, duration=123, liveType=UNKNOWN, priority=10, tableId=IngressPipeImpl.l2_exact_table, appId=org.onosproject.sdn, selector=[], treatment=DefaultTrafficTreat
ment([Immediate], add_switch_id=[port_num=0x1], cleared=false, StatTrigger=null, metadata=null)
  id=b4000058ec2a2, state=ADDED, bytes=102, duration=123, liveType=UNKNOWN, priority=10, tableId=IngressPipeImpl.l2_exact_table, appId=org.onosproject.sdn, selector=[hdr.ethernet.dst_addr=0xffffffff
ffff], treatment=DefaultTrafficTreatment([Immediate], set_multicast_group_id=0xff, switch_id.value=0x605), deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
  id=b4000058444c84, state=ADDED, bytes=303, packets=3, duration=123, liveType=UNKNOWN, priority=10, tableId=IngressPipeImpl.l2_exact_table, appId=org.onosproject.sdn, selector=[hdr.ethernet.dst_addr=0xa20b
322dc79d], treatment=DefaultTrafficTreatment([IngressPipeImpl.set_egress_port(port_num=0x2)]], deferred[], transition=None, meter[], cleared=false, StatTrigger=null, metadata=null)
curl@root >
```

Figure 58: ONOS console with information relating to the flows of the network in phase 2

The list of devices and hosts known to the controller is shown in Figure 57, and the flows that have been injected into the devices are presented in Figure 58. It is evident that both of these are virtually identical to those of the last phase.

Figure 59 illustrates the connectivity between devices 192.168.137.252 and 192.168.137.253, manifesting through the VFRs with IPs 10.10.10.1 and 10.10.10.2.

```

apu@apu3d4 ~> sudo ip netns exec VRF1 ping 10.10.10.2
PING 10.10.10.2 (10.10.10.2) 56(84) bytes of data.
64 bytes from 10.10.10.2: icmp_seq=1 ttl=64 time=17.4 ms
64 bytes from 10.10.10.2: icmp_seq=2 ttl=64 time=16.9 ms
64 bytes from 10.10.10.2: icmp_seq=3 ttl=64 time=17.5 ms
64 bytes from 10.10.10.2: icmp_seq=4 ttl=64 time=17.8 ms
64 bytes from 10.10.10.2: icmp_seq=5 ttl=64 time=16.6 ms
^C
--- 10.10.10.2 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4006ms
rtt min/avg/max/mdev = 16.629/17.259/17.805/0.426 ms
apu@apu3d4 ~> [REDACTED]

apu@apu3d4 ~> sudo ip netns exec VRF1 ping 10.10.10.1
PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data.
64 bytes from 10.10.10.1: icmp_seq=1 ttl=64 time=18.7 ms
64 bytes from 10.10.10.1: icmp_seq=2 ttl=64 time=16.0 ms
64 bytes from 10.10.10.1: icmp_seq=3 ttl=64 time=17.5 ms
64 bytes from 10.10.10.1: icmp_seq=4 ttl=64 time=18.7 ms
64 bytes from 10.10.10.1: icmp_seq=5 ttl=64 time=17.6 ms
^C
--- 10.10.10.1 ping statistics ---
5 packets transmitted, 5 received, 0% packet loss, time 4006ms
rtt min/avg/max/mdev = 16.038/17.714/18.705/0.978 ms
apu@apu3d4 ~> [REDACTED]
    
```

Figure 59: Results from running the commands "ping" in the [VRF](#) of the two devices on the network in phase 2

6.3.3 Analysis and Interpretation

These experiments were constrained by a significant limitation, namely the availability of only two devices with [BMv2](#) installed for use. As previously stated, the original device provided is not suitable for use in this scenario due to insufficient storage capacity to install [BMv2](#). Furthermore, financial constraints precluded the acquisition of another device of the newer version. Obtaining more than two devices was crucial to effectively recreate more intricate testing scenarios that more closely resemble real-world contexts, such as the transmission of a message through a network of wireless devices. However, despite the inability to demonstrate this specific scenario due to a lack of resources, the observed functionality of the program, its code, and the logs from [BMv2](#) and [ONOS](#) provide strong indications that the code performs as intended in the aforementioned scenario.

In testing phase 1 of these tests, it was found that the default southbound API of [BMv2](#) is not P4Runtime, as previously assumed, but rather Thrift. The Thrift API [93] is a component of the Apache Thrift framework, which originated at Facebook. The goal of Thrift is to facilitate cross-language communication and data serialization by providing a language-agnostic interface that can be implemented in a variety of programming languages. In this case, Thrift is used to turn an easy to read file with controller instructions to device specific instruction in the data plane.

This decision by the developers may initially seem incongruous, but it is, in fact, a logical consequence of the broader strategic context. [BMv2](#) is not a production-grade switch. Rather, it is a software tool designed for educational and experimental purposes in conjunction with [P4](#). Thrift offers a more

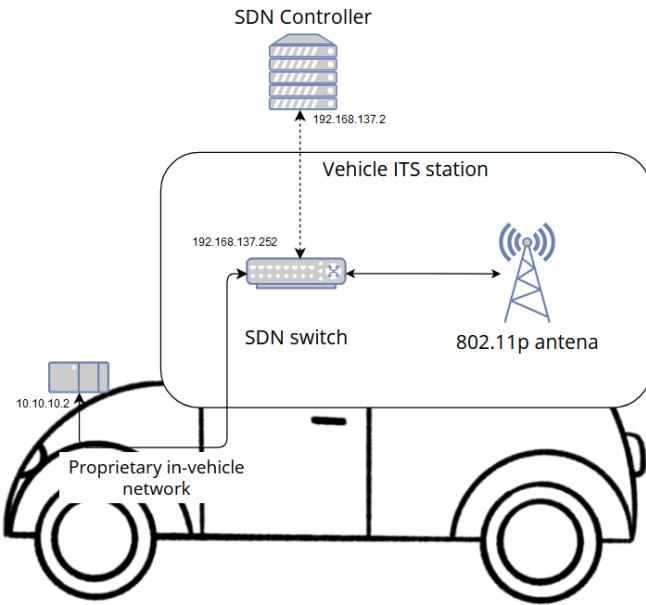


Figure 60: Diagram representing the [SDN](#) network as an Vehicle [ITS](#) sub-system

straightforward mechanism that streamlines the process of inserting rules on network devices. In consideration of this fact, it is clear that [BMv2](#) was not developed with high performance as a primary objective. Accordingly, any practical development should seek to implement a faster solution in the event that performance issues arise.

The transposition of the final scenario to an [SDN](#) network, illustrated in Figure 55, to an Vehicle [ITS](#) sub-system has resulted in the diagram presented in Figure 60. From this diagram, a number of interesting observations and comparisons can be made. In this context, the [SDN](#) switch can be compared to an [ITS-S](#) gateway and an [ITS-S](#) router. It is clear that any comparison should be viewed with caution, as the device under consideration does not align perfectly with any of the aforementioned categories. This comparison is intended to illustrate the device's functionality, rather than to provide a definitive categorization. It is inaccurate to categorize this device as TCP/IP to [ETSI](#) protocol translator, as it lacks the requisite functionality to convert commonly used Internet protocols into [ETSI](#) equivalents. The [SDN](#) controller could be integrated into the vehicle [ITS](#) station as an [ITS](#) application in a manner that would align with the standards and classifications set forth by the [ETSI](#).

As anticipated, the [BMv2](#) proved to be the optimal choice. This experiment was successful in enabling [BMv2](#) to function with an [ONOS](#) controller and the [ath9k](#) patch. The development of this test demonstrates that [P4](#) is capable of functioning with [IEEE](#) 802.11p, which indicates that this technology can be utilized for vehicular deployment. The power and control provided by [P4](#) affords users a vast array of potential applications. According to our tests it is possible to create a [P4](#) based network that complies with [ETSI](#) norms, making it possible for this technology to be used in the real world.

Conclusions and Future work

This chapter presents a synopsis of the work conducted, accompanied by a set of conclusions and a proposal for future research directions.

7.1 Conclusions

In this dissertation, the main objective was to develop an **OBU** in accordance with the principles of **SDN**, thereby aligning this document with the domain of **SDVN**. Throughout the course of this project, a critical examination of the actual implications of this goal was undertaken. The findings of the research revealed that the objective of making an **OBU SDN** compatible is, in essence, equivalent to developing an **SDN** Vehicle **ITS** sub-system. Moreover, the insights gained from the research enabled the presentation of a feasible and realistic trajectory for **SDVN**. This vision is grounded in the recognition that, for **SDVN** to be deployed in the future, interoperability with traditional devices is of paramount importance. In essence, the success of **SDVN** will be contingent upon its capacity to integrate with existing infrastructure. In the practical implementation stage, a series of tests were conducted on the most relevant software for data plane emulation in the field of **SDN** in order to ascertain the most appropriate software for vehicular scenarios. In accordance with the requirements of this dissertation, the queries pertaining to the control plane were set aside in order to prioritize the data plane, thereby facilitating the recreation of the **OBU**. Prototype devices were configured and the **BMv2** and **OvS** software were installed on them. In the absence of any notable requirements, **ONOS** was chosen as the control plane technology. Subsequently, the **ath9k** patch was implemented with the objective of evaluating the compatibility of communication utilizing the intended **IEEE 802.11p** protocol. The results of the tests revealed that **OvS** was incompatible with wireless interfaces, whereas **BMv2** demonstrated no such incompatibility. The tests conducted on **OvS** and **BMv2** served to reinforce the hypothesis that **P4** is an adequate data plane technology. The **P4** language, in conjunction with the controller, was employed to modify the subsequent layers, thereby establishing communication in **IEEE 802.11p** between two devices. It was thus demonstrated that **BMv2** represents the optimal software for utilization in **SDVN**. The present dissertation has been completed in accordance with the stipulated criteria and objectives. All of the proposed goals have been achieved, thereby providing a crucial step in establishing a foundation for the implementation of a complete **SDVN** testing scenario based

on real hardware. Ultimately, the research, work, and subsequent findings presented in this document provide sufficient evidence to demonstrate that the integration of [SDN](#) principles in the field of VANETs is a viable and achievable proposition.

7.2 Future work

The natural progression of this work is to adapt the remaining [ETSI](#) sub-systems to be [SDN](#)-compatible. This process should begin with the Central and Roadside sub-systems, as they represent the most critical components in this context. Upon completion of this phase, the foundation will be established for a comprehensive, standards-compliant [SDN](#) system implementation. This second phase would entail the development of the necessary controller and data plane logic to achieve intercommunication between [SDN](#) and non-[SDN](#) devices in conformance with established standards. The most desirable long-term objective would be the recreation of a [SDN](#) device that is an exact functional replica of a traditional device. Once this has been achieved, the technology can be implemented and evaluated in conjunction with actual road networks, in a controlled but relevant environment. However, the complete system should not be regarded as a final stage of research but rather as the foundation for a new line of inquiry. The potential of [P4](#) can be harnessed to facilitate the testing of a multitude of aspects pertaining to the [ETSI](#) protocols. The first and most crucial task, which is not directly related to the [ETSI](#) protocols, is to identify the optimal controller placement. This, along with related inquiries such as the evaluation of election algorithms for the controller in clusters of vehicles, would represent a vital priority for ensuring the success and advancement of this technology. Finally, this dissertation also lays the groundwork for the redesign and upgrade of the device used on the [OBU](#). The patch currently in use to achieve [IEEE](#) 802.11p functionality may not be a long-term solution. Therefore, it would be prudent to explore a more permanent solution, such as [SDR](#).

Bibliography

- [1] N. Feamster, J. Rexford, and E. Zegura. “The road to sdn”. In: *Queue* 11 (2013-12). doi: [10.1145/2559899.2560327](https://doi.org/10.1145/2559899.2560327) (cit. on pp. 1, 30–35).
- [2] J. Jakubiak and Y. Koucheryavy. “State of the Art and Research Challenges for VANETs”. In: *2008 5th IEEE Consumer Communications and Networking Conference*. 2008 5th IEEE Consumer Communications and Networking Conference. ISSN: 2331-9860. 2008-01, pp. 912–916. doi: [10.1109/ccnc08.2007.212](https://doi.org/10.1109/ccnc08.2007.212) (cit. on pp. 5, 7, 13, 16, 28).
- [3] S. Al-Sultan et al. “A comprehensive survey on vehicular Ad Hoc network”. In: *Journal of Network and Computer Applications* 37 (2014-01-01), pp. 380–392. issn: 1084-8045. doi: [10.1016/j.jnca.2013.02.036](https://doi.org/10.1016/j.jnca.2013.02.036). url: <https://www.sciencedirect.com/science/article/pii/S108480451300074X> (visited on 2022-11-02) (cit. on pp. 5, 7, 13–15, 22, 23, 27, 28).
- [4] W. Liang et al. “Vehicular Ad Hoc Networks: Architectures, Research Issues, Methodologies, Challenges, and Trends”. In: *International Journal of Distributed Sensor Networks* 11.8 (2015-08-01), p. 745303. issn: 1550-1477. doi: [10.1155/2015/745303](https://doi.org/10.1155/2015/745303). url: <http://journals.sagepub.com/doi/10.1155/2015/745303> (visited on 2023-02-22) (cit. on pp. 5, 7, 13, 21, 25, 27).
- [5] H. Dinh Thai et al. “Applications of Repeated Games in Wireless Networks: A Survey”. In: *IEEE Communications Surveys & Tutorials* 17 (2015-01-11). doi: [10.1109/CST.2015.2445789](https://doi.org/10.1109/CST.2015.2445789) (cit. on p. 6).
- [6] *Platform: C-Roads*. url: <https://www.c-roads.eu/platform.html> (visited on 2024-11-25) (cit. on p. 6).
- [7] Y. Toor et al. “Vehicle Ad Hoc networks: applications and related technical issues”. In: *IEEE Communications Surveys & Tutorials* 10.3 (2008). Conference Name: IEEE Communications Surveys & Tutorials, pp. 74–88. issn: 1553-877X. doi: [10.1109/CST.2008.4625806](https://doi.org/10.1109/CST.2008.4625806) (cit. on pp. 7, 8, 13, 16, 17, 21, 26, 28).

BIBLIOGRAPHY

- [8] S. M. Corson and J. P. Macker. *Mobile Ad hoc Networking (MANET): Routing Protocol Performance Issues and Evaluation Considerations*. Request for Comments RFC 2501. Num Pages: 12. Internet Engineering Task Force, 1999-01. doi: [10.17487/RFC2501](https://doi.org/10.17487/RFC2501). url: <https://datatracker.ietf.org/doc/rfc2501> (visited on 2023-09-28) (cit. on p. 8).
- [9] C2C-CC. *CAR 2 CAR Communication Consortium Manifesto - Overview of the C2C-CC System*. 2007-08-28. url: https://www.car-2-car.org/fileadmin/documents/General_Documents/C2C-CC_Manifesto_Aug_2007.pdf (visited on 2023-04-23) (cit. on pp. 9, 13).
- [10] ETSI. *Intelligent Transport Systems (ITS); Communications Architecture*. 2010-09. url: https://www.etsi.org/deliver/etsi_en/302600_302699/302665/01.01.01_60/en_302665v010101p.pdf (visited on 2023-04-27) (cit. on pp. 9–13, 15, 24, 26).
- [11] M. Faezipour et al. “Progress and challenges in intelligent vehicle area networks”. In: *Communications of the ACM* 55.2 (2012), pp. 90–100. issn: 0001-0782. doi: [10.1145/2076450.2076470](https://doi.org/10.1145/2076450.2076470). url: <https://dl.acm.org/doi/10.1145/2076450.2076470> (visited on 2023-09-19) (cit. on p. 13).
- [12] Ş. ŞORIGA. “ITS-G5 AND MOBILE WIMAX PERFORMANCE IN VEHICLE-TO-INFRASTRUCTURE COMMUNICATIONS”. In: *ISSN 1454-234x UPB Scientific Bulletin, Series C: Electrical Engineering*. (2012). url: https://www.scientificbulletin.upb.ro/rev_docs_arhiva/full4d5_755707.pdf (visited on 2023-10-16) (cit. on pp. 13, 19).
- [13] M. S. Anwer and C. Guy. “A Survey of VANET Technologies”. In: *Journal of Emerging Trends in Computing and Information Sciences* (2014) (cit. on pp. 13, 15, 19).
- [14] A. Festag. “Standards for vehicular communication—from IEEE 802.11p to 5G”. In: *e & i Elektrotechnik und Informationstechnik* 132.7 (2015-11-01), pp. 409–416. issn: 1613-7620. doi: [10.1007/s00502-015-0343-0](https://doi.org/10.1007/s00502-015-0343-0). url: <https://doi.org/10.1007/s00502-015-0343-0> (visited on 2023-10-16) (cit. on pp. 13, 22).
- [15] J. Härri and J. Kenney. “Multi-Channel Operations, Coexistence and Spectrum Sharing for Vehicular Communications”. In: ed. by C. Campolo, A. Molinaro, and R. Scopigno. Book Title: *Vehicular ad hoc Networks*. Cham: Springer International Publishing, 2015, pp. 193–218. isbn: 978-3-319-15496-1 978-3-319-15497-8. doi: [10.1007/978-3-319-15497-8_7](https://doi.org/10.1007/978-3-319-15497-8_7). url: https://link.springer.com/10.1007/978-3-319-15497-8_7 (visited on 2023-10-05) (cit. on pp. 13, 18, 19).
- [16] N. Asselin-Miller et al. “Study on the Deployment of C-ITS in Europe: Final Report”. In: 1 (2016-05-02) (cit. on pp. 13, 17, 18, 53, 104).
- [17] R. bibinitperiod Schwarz. “Intelligent Transportation Systems Using IEEE 802.11p”. In: (2019-02-14) (cit. on pp. 13, 16, 17).

- [18] R. Weber, J. Misener, and V. Park. "C-V2X - A Communication Technology for Cooperative, Connected and Automated Mobility". In: *Mobile Communication - Technologies and Applications*; 24. ITG-Symposium. Mobile Communication - Technologies and Applications; 24. ITG-Symposium. 2019-05, pp. 1–6. url: <https://ieeexplore.ieee.org/abstract/document/8731783> (visited on 2023-10-23) (cit. on pp. 13, 20).
- [19] S. Gyawali et al. "Challenges and Solutions for Cellular Based V2X Communications". In: *IEEE Communications Surveys & Tutorials* 23.1 (2021), pp. 222–255. issn: 1553-877X, 2373-745X. doi: [10.1109/COMST.2020.3029723](https://doi.org/10.1109/COMST.2020.3029723). url: <https://ieeexplore.ieee.org/document/9217500/> (visited on 2023-10-16) (cit. on pp. 13, 19, 20).
- [20] R. Sedar et al. "Standards-Compliant Multi-Protocol On-Board Unit for the Evaluation of Connected and Automated Mobility Services in Multi-Vendor Environments". en. In: *Sensors* 21.6 (2021-01). Number: 6 Publisher: Multidisciplinary Digital Publishing Institute, p. 2090. issn: 1424-8220. doi: [10.3390/s21062090](https://doi.org/10.3390/s21062090). url: <https://www.mdpi.com/1424-8220/21/6/2090> (visited on 2022-10-10) (cit. on pp. 13, 56, 57).
- [21] M. N. Tahir, M. Katz, and Z. Javed. "Poster: Connected Vehicles using Short-range (Wi-Fi & IEEE 802.11p) and Long-range Cellular Networks (LTE & 5G)". In: *2021 IEEE 29th International Conference on Network Protocols (ICNP)*. 2021 IEEE 29th International Conference on Network Protocols (ICNP). ISSN: 2643-3303. 2021-11, pp. 1–2. doi: [10.1109/ICNP52444.2021.9651959](https://doi.org/10.1109/ICNP52444.2021.9651959). url: <https://ieeexplore.ieee.org/abstract/document/9651959> (visited on 2023-10-03) (cit. on p. 13).
- [22] ETSI. *Intelligent Transport Systems (ITS); ITS-G5 Access layer specification for Intelligent Transport Systems operating in the 5 GHz frequency band*. 2020-01. url: https://www.etsi.org/deliver/etsi_en/302600_302699/302663/01.03.01_60/en_302663v010301.pdf (visited on 2023-11-08) (cit. on p. 15).
- [23] History of IEEE. url: <https://www.ieee.org/about/ieee-history.html> (visited on 2023-11-14) (cit. on p. 16).
- [24] D. Jiang and L. Delgrossi. "IEEE 802.11p: Towards an International Standard for Wireless Access in Vehicular Environments". In: IEEE Vehicular Technology Conference. 2008-06-14, pp. 2036–2040. doi: [10.1109/VETECS.2008.458](https://doi.org/10.1109/VETECS.2008.458) (cit. on pp. 16, 17, 26).
- [25] A. Festag. "Cooperative intelligent transport systems standards in europe". In: *IEEE Communications Magazine* 52.12 (2014-12). Conference Name: IEEE Communications Magazine, pp. 166–172. issn: 1558-1896. doi: [10.1109/MCOM.2014.6979970](https://doi.org/10.1109/MCOM.2014.6979970) (cit. on pp. 17, 18, 22–24).
- [26] ETSI. *Electromagnetic compatibility and Radio spectrum Matters (ERM); Intelligent Transport Systems (ITS); Part 1: Technical characteristics for pan-European harmonized communications equipment operating in the 5 GHz frequency range and intended for critical road-safety applications; System Reference Document*. 2005-06. url: https://www.etsi.org/deliver/etsi_tr/1

BIBLIOGRAPHY

- [02400_102499/10249201/01.01.01_60/tr_10249201v010101p.pdf](https://ieeexplore.ieee.org/document/10063942) (visited on 2023-11-17) (cit. on p. 18).
- [27] "IEEE Standard for Information Technology–Telecommunications and Information Exchange between Systems Local and Metropolitan Area Networks–Specific Requirements Part 11: Wireless LAN Medium Access Control (MAC) and Physical Layer (PHY) Specifications Amendment 5: Enhancements for Next Generation V2X". In: *IEEE Std 802.11bd-2022 (Amendment to IEEE Std 802.11-2020 as amended by IEEE Std 802.11ax-2021, IEEE Std 802.11ay-2021, IEEE Std 802.11ba-2021, IEEE Std 802.11-2020/Cor 1-2022, and IEEE Std 802.11az-2022)* (2023-03). Conference Name: IEEE Std 802.11bd-2022 (Amendment to IEEE Std 802.11-2020 as amended by IEEE Std 802.11ax-2021, IEEE Std 802.11ay-2021, IEEE Std 802.11ba-2021, IEEE Std 802.11-2020/Cor 1-2022, and IEEE Std 802.11az-2022), pp. 1–144. doi: [10.1109/IEEESTD.2023.10063942](https://doi.org/10.1109/IEEESTD.2023.10063942). url: <https://ieeexplore.ieee.org/document/10063942> (visited on 2023-10-25) (cit. on p. 18).
- [28] 3GPP – *The Mobile Broadband Standard*. 3GPP. url: <https://www.3gpp.org/> (visited on 2023-11-16) (cit. on p. 19).
- [29] 5GAA. 5GAA. url: <https://5gaa.org/> (visited on 2023-11-16) (cit. on p. 20).
- [30] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; GeoNetworking; Part 1: Requirements*. 2014-04. url: https://www.etsi.org/deliver/etsi_en/302600_302699/30263601/01.02.01_60/en_30263601v010201p.pdf (visited on 2023-11-27) (cit. on p. 21).
- [31] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; GeoNetworking; Part 2: Scenarios*. 2013-11. url: https://www.etsi.org/deliver/etsi_en/302600_302699/30263602/01.02.01_60/en_30263602v010201p.pdf (visited on 2023-11-27) (cit. on p. 21).
- [32] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; GeoNetworking; Part 3: Network Architecture*. 2014-12. url: https://www.etsi.org/deliver/etsi_en/302600_302699/30263603/01.02.01_60/en_30263603v010201p.pdf (visited on 2023-10-06) (cit. on pp. 21, 24, 26).
- [33] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; GeoNetworking; Part 4: Geographical addressing and forwarding for point-to-point and point-to-multipoint communications; Sub-part 1: Media-Independent Functionality*. 2020-01. url: https://www.etsi.org/deliver/etsi_en/302600_302699/3026360401/01.04.01_60/en_3026360401v010401p.pdf (visited on 2023-11-27) (cit. on p. 21).
- [34] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; GeoNetworking; Part 5: Transport Protocols; Sub-part 1: Basic Transport Protocol*. 2019-05. url: https://www.etsi.org/deliver/etsi_en/302600_302699/3026360501/02.02.01_60/en_3026360501v020201p.pdf (visited on 2023-11-27) (cit. on p. 21).

- [35] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; GeoNetworking; Part 6: Internet Integration; Sub-part 1: Transmission of IPv6 Packets over GeoNetworking Protocols*. 2014-05. url: https://www.etsi.org/deliver/etsi_en/302600_302699/3026360601/01.02.01_60/en_3026360601v010201p.pdf (visited on 2023-11-27) (cit. on p. 21).
- [36] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; Basic Set of Applications; Part 2: Specification of Cooperative Awareness Basic Service*. 2019-04. url: https://www.etsi.org/deliver/etsi_en/302600_302699/30263702/01.04.01_60/en_30263702v010401p.pdf (visited on 2023-11-27) (cit. on p. 22).
- [37] ETSI. *Intelligent Transport Systems (ITS); Vehicular Communications; Basic Set of Applications; Part 3: Specifications of Decentralized Environmental Notification Basic Service*. 2019-04. url: https://www.etsi.org/deliver/etsi_en/302600_302699/30263703/01.03.01_60/en_30263703v010301p.pdf (visited on 2023-11-27) (cit. on p. 22).
- [38] A. K. Malhi, S. Batra, and H. S. Pannu. "Security of vehicular ad-hoc networks: A comprehensive survey". In: *Computers & Security* 89 (2020-02), p. 101664. issn: 01674048. doi: [10.1016/j.cose.2019.101664](https://doi.org/10.1016/j.cose.2019.101664). url: <https://linkinghub.elsevier.com/retrieve/pii/S0167404818312872> (visited on 2023-11-28) (cit. on pp. 25, 26).
- [39] H. Hasrouny et al. "VANet security challenges and solutions: A survey". In: *Vehicular Communications* 7 (2017-01-01), pp. 7–20. issn: 2214-2096. doi: [10.1016/j.vehcom.2017.01.002](https://doi.org/10.1016/j.vehcom.2017.01.002). url: <https://www.sciencedirect.com/science/article/pii/S2214209616301231> (visited on 2023-12-15) (cit. on p. 25).
- [40] ETSI. *Intelligent Transport Systems (ITS); Security; ITS communications security architecture and security management*. 2018-04. url: https://www.etsi.org/deliver/etsi_ts/102900_102999/102940/01.03.01_60/ts_102940v010301p.pdf (visited on 2023-12-05) (cit. on p. 27).
- [41] ETSI. *Intelligent Transport Systems (ITS); Security; ITS communications security architecture and security management; Release 2*. 2021-07. url: https://www.etsi.org/deliver/etsi_ts/102900_102999/102940/02.01.01_60/ts_102940v020101p.pdf (visited on 2023-11-28) (cit. on p. 27).
- [42] M. J. N. Mahi et al. "A Review on VANET Research: Perspective of Recent Emerging Technologies". In: *IEEE Access* 10 (2022). Conference Name: IEEE Access, pp. 65760–65783. issn: 2169-3536. doi: [10.1109/ACCESS.2022.3183605](https://doi.org/10.1109/ACCESS.2022.3183605). url: <https://ieeexplore.ieee.org/abstract/document/9797696> (visited on 2023-12-19) (cit. on p. 28).
- [43] M. Lu et al. "Pan-European deployment of C-ITS: the way forward". In: (2019) (cit. on pp. 28, 29).
- [44] B. A. A. Nunes et al. "A Survey of Software-Defined Networking: Past, Present, and Future of Programmable Networks". In: *IEEE Communications Surveys & Tutorials* 16.3 (2014). Conference Name: IEEE Communications Surveys & Tutorials, pp. 1617–1634. issn: 1553-877X. doi: [10.1109/SURV.2014.012214.00180](https://doi.org/10.1109/SURV.2014.012214.00180) (cit. on pp. 30, 31, 33–36, 45–47).

BIBLIOGRAPHY

- [45] W. Xia et al. “A Survey on Software-Defined Networking”. In: *IEEE Communications Surveys & Tutorials* 17.1 (2015). Conference Name: IEEE Communications Surveys & Tutorials, pp. 27–51. issn: 1553-877X. doi: [10.1109/COMST.2014.2330903](https://doi.org/10.1109/COMST.2014.2330903) (cit. on pp. 30, 33–37, 40, 45–47, 49).
- [46] Open Networking Foundation. en-US. url: <https://opennetworking.org/> (visited on 2024-01-12) (cit. on p. 30).
- [47] L. L. Peterson et al. *Software-Defined Networks: A Systems Approach*. English. Wroclaw: Systems Approach LLC, 2021-01. isbn: 978-1-73647-210-1 (cit. on pp. 30–32, 35–38, 40–49).
- [48] R. Bifulco and G. Retvari. “A Survey on the Programmable Data Plane: Abstractions, Architectures, and Open Problems”. en. In: *2018 IEEE 19th International Conference on High Performance Switching and Routing (HPSR)*. Bucharest, Romania: IEEE, 2018-06, pp. 1–7. isbn: 978-1-5386-7801-5. doi: [10.1109/HPSR.2018.8850761](https://doi.org/10.1109/HPSR.2018.8850761). url: <https://ieeexplore.ieee.org/document/8850761/> (visited on 2023-12-22) (cit. on pp. 30, 31, 33, 35–38, 40, 42, 47).
- [49] A. S. Thyagaturu et al. “Software Defined Optical Networks (SDONs): A Comprehensive Survey”. In: *IEEE Communications Surveys & Tutorials* 18.4 (2016). Conference Name: IEEE Communications Surveys & Tutorials, pp. 2738–2786. issn: 1553-877X. doi: [10.1109/COMST.2016.2586999](https://doi.org/10.1109/COMST.2016.2586999) (cit. on pp. 31, 32, 34, 36, 44, 46, 47).
- [50] A. Alowa. “Scalable Reliable Controller Placement in Software Defined Networking”. In: (2020-09). url: <https://core.ac.uk/reader/355870830> (visited on 2024-03-19) (cit. on p. 32).
- [51] A. Liatifis et al. “Advancing SDN from OpenFlow to P4: A Survey”. In: *ACM Computing Surveys* 55.9 (2023-01), 186:1–186:37. issn: 0360-0300. doi: [10.1145/3556973](https://doi.org/10.1145/3556973). url: <https://dl.acm.org/doi/10.1145/3556973> (visited on 2024-01-11) (cit. on pp. 32, 38, 49).
- [52] R. C. Sofia and J. Soldatos. “Shaping the Future of IoT with Edge Intelligence; How Edge Computing Enables the Next Generation of IoT Applications”. en. In: (2024) (cit. on p. 32).
- [53] D. Kreutz et al. “Software-Defined Networking: A Comprehensive Survey”. In: *Proceedings of the IEEE* 103.1 (2015-01). Conference Name: Proceedings of the IEEE, pp. 14–76. issn: 1558-2256. doi: [10.1109/JPROC.2014.2371999](https://doi.org/10.1109/JPROC.2014.2371999) (cit. on pp. 32–37, 39, 42, 44–49).
- [54] N. McKeown et al. “OpenFlow: enabling innovation in campus networks”. In: *ACM SIGCOMM Computer Communication Review* 38.2 (2008-03), pp. 69–74. issn: 0146-4833. doi: [10.1145/1355734.1355746](https://doi.org/10.1145/1355734.1355746). url: <https://doi.org/10.1145/1355734.1355746> (visited on 2023-02-10) (cit. on pp. 33, 37, 43, 45).
- [55] S. Li et al. “Protocol Oblivious Forwarding (POF): Software-Defined Networking with Enhanced Programmability”. In: *IEEE Network* 31.2 (2017-03). Conference Name: IEEE Network, pp. 58–66. issn: 1558-156X. doi: [10.1109/MNET.2017.1600030NM](https://doi.org/10.1109/MNET.2017.1600030NM) (cit. on pp. 36, 39).

- [56] Z. Latif et al. "A comprehensive survey of interface protocols for software defined networks". en. In: *Journal of Network and Computer Applications* 156 (2020-04), p. 102563. issn: 1084-8045. doi: [10.1016/j.jnca.2020.102563](https://doi.org/10.1016/j.jnca.2020.102563). url: <https://www.sciencedirect.com/science/article/pii/S1084804520300370> (visited on 2022-11-07) (cit. on pp. 37, 44, 46–49).
- [57] F. Hauser et al. "A survey on data plane programming with P4: Fundamentals, advances, and applied research". In: *Journal of Network and Computer Applications* 212 (2021-08), p. 103561. issn: 1084-8045. doi: [10.1016/j.jnca.2022.103561](https://doi.org/10.1016/j.jnca.2022.103561). url: <https://www.sciencedirect.com/science/article/pii/S1084804522002028> (visited on 2024-01-26) (cit. on pp. 38–41, 43, 45, 46).
- [58] P. Bosshart et al. "P4: programming protocol-independent packet processors". In: *ACM SIGCOMM Computer Communication Review* 44.3 (2014-07), pp. 87–95. issn: 0146-4833. doi: [10.1145/2656877.2656890](https://doi.org/10.1145/2656877.2656890). url: <https://doi.org/10.1145/2656877.2656890> (visited on 2022-11-30) (cit. on pp. 39, 40).
- [59] Open Compute Project. en. url: <https://www.opencompute.org> (visited on 2024-03-15) (cit. on p. 40).
- [60] Indigo - Confluence. url: <https://floodlight.atlassian.net/wiki/spaces/Indigo/overview?homepageId=2392077> (visited on 2024-03-15) (cit. on p. 41).
- [61] Open Network Linux. url: <http://opennetlinux.org/> (visited on 2024-03-18) (cit. on p. 44).
- [62] Stratum. en-US. url: <https://opennetworking.org/stratum/> (visited on 2024-03-18) (cit. on p. 44).
- [63] Sonic Foundation – Linux Foundation Project. en-US. url: <https://sonicfoundation.dev/> (visited on 2024-03-18) (cit. on p. 44).
- [64] OpenConfig. url: <https://openconfig.net/> (visited on 2024-03-15) (cit. on pp. 45, 46).
- [65] L. Zhu et al. "SDN Controllers: A Comprehensive Analysis and Performance Evaluation Study". In: *ACM Computing Surveys* 53.6 (2020), 133:1–133:40. issn: 0360-0300. doi: [10.1145/3421764](https://doi.org/10.1145/3421764). url: <https://doi.org/10.1145/3421764> (visited on 2023-02-04) (cit. on pp. 47, 48).
- [66] K. Smida et al. "Efficient SDN Controller for Safety Applications in SDN-Based Vehicular Networks: POX, Floodlight, ONOS or OpenDaylight?" In: *2020 IEEE Eighth International Conference on Communications and Networking (ComNet)*. ISSN: 2473-7585. 2020-10, pp. 1–6. doi: [10.1109/ComNet47917.2020.9306095](https://ieeexplore.ieee.org/abstract/document/9306095). url: <https://ieeexplore.ieee.org/abstract/document/9306095> (visited on 2024-02-19) (cit. on pp. 50, 52, 55, 64).
- [67] I. Ku et al. "Towards software-defined VANET: Architecture and services". In: *2014 13th Annual Mediterranean Ad Hoc Networking Workshop (MED-HOC-NET)*. 2014-06, pp. 103–110. doi: [10.1109/MedHocNet.2014.6849111](https://doi.org/10.1109/MedHocNet.2014.6849111) (cit. on pp. 50, 52, 59).

BIBLIOGRAPHY

- [68] N. Cardona et al. "Software-Defined Vehicular Networking: Opportunities and Challenges". In: *IEEE Access* 8 (2020). Conference Name: IEEE Access, pp. 219971–219995. issn: 2169-3536. doi: [10.1109/ACCESS.2020.3042717](https://doi.org/10.1109/ACCESS.2020.3042717) (cit. on pp. 50, 52, 54, 55).
- [69] J. Bhatia et al. "Software defined vehicular networks: A comprehensive review". en. In: *International Journal of Communication Systems* 32.12 (2019), e4005. issn: 1099-1131. doi: [10.1002/dac.4005](https://doi.org/10.1002/dac.4005). url: <https://onlinelibrary.wiley.com/doi/abs/10.1002/dac.4005> (visited on 2022-10-11) (cit. on pp. 50, 55).
- [70] W. Ben Jaballah, M. Conti, and C. Lal. "Security and design requirements for software-defined VANETs". In: *Computer Networks* 169 (2020-03), p. 107099. issn: 1389-1286. doi: [10.1016/j.comnet.2020.107099](https://doi.org/10.1016/j.comnet.2020.107099). url: <https://www.sciencedirect.com/science/article/pii/S1389128619306553> (visited on 2023-12-15) (cit. on pp. 52–54).
- [71] R. Sarpong, B. Sousa, and F. Araujo. "The potential of SDNs and multihoming in VANETs: A Comprehensive Survey". en. In: (2023-11) (cit. on pp. 52, 58).
- [72] S. Toufga et al. "Towards Dynamic Controller Placement in Software Defined Vehicular Networks". In: *Sensors* 20 (2020-03), p. 1701. doi: [10.3390/s20061701](https://doi.org/10.3390/s20061701) (cit. on p. 54).
- [73] S. Toufga et al. "OpenFlow based Topology Discovery Service in Software Defined Vehicular Networks: limitations and future approaches". In: *2018 IEEE Vehicular Networking Conference (VNC)*. ISSN: 2157-9865. 2018-12, pp. 1–4. doi: [10.1109/VNC.2018.8628349](https://doi.org/10.1109/VNC.2018.8628349) (cit. on pp. 54, 55).
- [74] L. Nkenyereye et al. "Software-Defined Network-Based Vehicular Networks: A Position Paper on Their Modeling and Implementation". en. In: *Sensors* 19.17 (2019-01). Number: 17 Publisher: Multidisciplinary Digital Publishing Institute, p. 3788. issn: 1424-8220. doi: [10.3390/s19173788](https://doi.org/10.3390/s19173788). url: <https://www.mdpi.com/1424-8220/19/17/3788> (visited on 2022-10-10) (cit. on p. 54).
- [75] F. Raviglione, M. Malinverno, and C. Casetti. "Open source testbed for vehicular communication". In: *Proceedings of the Twentieth ACM International Symposium on Mobile Ad Hoc Networking and Computing*. MobiHoc '19. New York, NY, USA: Association for Computing Machinery, 2019-07, pp. 405–406. isbn: 978-1-4503-6764-6. doi: [10.1145/3323679.3326623](https://doi.org/10.1145/3323679.3326623). url: <https://dl.acm.org/doi/10.1145/3323679.3326623> (visited on 2024-03-06) (cit. on pp. 55, 57).
- [76] S. Laux et al. "Demo: OpenC2X – An open source experimental and prototyping platform supporting ETSI ITS-G5". In: *2016 IEEE Vehicular Networking Conference (VNC)*. 2016 IEEE Vehicular Networking Conference (VNC). Columbus, OH, USA: IEEE, 2016-12, pp. 1–2. isbn: 978-1-5090-5197-7. doi: [10.1109/VNC.2016.7835955](https://doi.org/10.1109/VNC.2016.7835955). url: <http://ieeexplore.ieee.org/document/7835955/> (visited on 2025-04-22) (cit. on p. 56).
- [77] Vanetza - Your open-source ETSI C-ITS protocol stack. url: <https://www.vanetza.org/> (visited on 2025-04-22) (cit. on p. 56).

- [78] G. Secinti et al. “Software Defined Architecture for VANET: A Testbed Implementation with Wireless Access Management”. In: *IEEE Communications Magazine* 55.7 (2017-07). Conference Name: IEEE Communications Magazine, pp. 135–141. issn: 1558-1896. doi: [10.1109/MCOM.2017.1601186](https://doi.org/10.1109/MCOM.2017.1601186) (cit. on pp. 56, 57).
- [79] *Open vSwitch*. url: <https://www.openvswitch.org/> (visited on 2024-10-24) (cit. on p. 56).
- [80] P. Rito et al. “Aveiro Tech City Living Lab: A Communication, Sensing, and Computing Platform for City Environments”. In: *IEEE Internet of Things Journal* 10.15 (2023-08). Conference Name: IEEE Internet of Things Journal, pp. 13489–13510. issn: 2327-4662. doi: [10.1109/JIOT.2023.3262627](https://doi.org/10.1109/JIOT.2023.3262627). url: <https://ieeexplore.ieee.org/abstract/document/10083171> (visited on 2024-02-21) (cit. on pp. 56, 57).
- [81] O. Sadio, I. Ngom, and C. Lishou. “Design and Prototyping of a Software Defined Vehicular Networking”. In: *IEEE Transactions on Vehicular Technology* 69.1 (2020-01). Conference Name: IEEE Transactions on Vehicular Technology, pp. 842–850. issn: 1939-9359. doi: [10.1109/TVT.2019.2950426](https://doi.org/10.1109/TVT.2019.2950426) (cit. on p. 57).
- [82] *PC Engines apu2 system boards*. url: <https://www.pcengines.ch/apu2.htm> (visited on 2024-11-25) (cit. on p. 61).
- [83] R. P. Ltd. *Buy a Raspberry Pi 5*. en-GB. url: <https://www.raspberrypi.com/products/raspberry-pi-5/> (visited on 2024-11-25) (cit. on p. 61).
- [84] *Deploy AMD*. url: <https://opennetlinux.org/doc-deploy-amd.html> (visited on 2024-11-25) (cit. on p. 62).
- [85] <https://gitlab.com/hpi-potsdam/osm/g5-on-linux/11p-on-linux#id14>. url: <https://gitlab.com/hpi-potsdam/osm/g5-on-linux/11p-on-linux#id14> (visited on 2024-11-25) (cit. on p. 63).
- [86] *ath10k support (#9) · Issues · HPI Potsdam / Operating Systems and Middleware Group / ETSI ITS-G5 on Linux / IEEE 802.11p on Linux · GitLab*. en. 2023-10. url: <https://gitlab.com/hpi-potsdam/osm/g5-on-linux/11p-on-linux/-/issues/9> (visited on 2024-11-25) (cit. on p. 63).
- [87] *Open vSwitch*. url: <https://www.openvswitch.org/> (visited on 2024-11-25) (cit. on p. 63).
- [88] *p4lang/behavioral-model: The reference P4 software switch*. url: <https://github.com/p4lang/behavioral-model> (visited on 2024-11-25) (cit. on p. 63).
- [89] *opennetworkinglab/onos: Open Network Operating System*. url: <https://github.com/opennetworkinglab/onos> (visited on 2024-11-25) (cit. on p. 64).
- [90] *External products sold which can use ath9k – Linux Wireless documentation*. url: <https://wireless.docs.kernel.org/en/latest/en/users/drivers/ath9k/products/external.html> (visited on 2024-11-25) (cit. on p. 68).

BIBLIOGRAPHY

- [91] *Baco-66/onos-apps: ONOS application for SDVN scenarios, based on the NGSDN Tutorial (Advanced) repository by the Open Networking Lab.* url: <https://github.com/Baco-66/onos-apps> (visited on 2024-11-25) (cit. on p. 70).
- [92] *Common Configuration Issues – Open vSwitch 3.4.90 documentation.* url: <https://docs.openvswitch.org/en/latest/faq/issues/> (visited on 2024-11-25) (cit. on p. 81).
- [93] *Apache Thrift - Home.* url: <https://thrift.apache.org/> (visited on 2024-11-25) (cit. on pp. 81, 89).
- [94] *p4lang/p4c: P4_16 reference compiler.* url: <https://github.com/p4lang/p4c> (visited on 2024-11-25) (cit. on p. 81).
- [95] J. M. Lourenço. *The NOVAthesis L^AT_EX Template User's Manual.* NOVA University Lisbon. 2021. url: <https://github.com/joaomlourenco/novathesis/raw/main/template.pdf> (cit. on p. 92).



Bundles of services

ANNEX I. BUNDLES OF SERVICES

Service bundle	C-ITS Services	Rationale
Bundle 1 Day 1, V2V, ITS-G5	<ul style="list-style-type: none"> Emergency brake light Emergency vehicle approaching Slow or stationary vehicle(s) Traffic jam ahead warning Hazardous location notification 	<ul style="list-style-type: none"> Day 1 safety-based V2V services based on ITS-G5 communication, likely to be deployed to vehicles supported by US legislation
Bundle 2 Day 1, V2I, mainly applicable to motorways	<ul style="list-style-type: none"> In-vehicle signage In-vehicle speed limits Probe vehicle data Shockwave damping Road works warning Weather conditions 	<ul style="list-style-type: none"> Day 1 V2I, services that deliver most benefit to motorways. Some services listed here may also be applicable to other road types
Bundle 3 Day 1, V2I, mainly applicable to urban areas	<ul style="list-style-type: none"> Green Light Optimal Speed Advisory (GLOSA) / Time To Green (TTG) Signal violation/Intersection safety Traffic signal priority request by designated vehicles 	<ul style="list-style-type: none"> Day 1 V2I, services expected to only be applicable in urban areas. Therefore, these services are in a separate bundle to those in Bundle 2
Bundle 4 Day 1.5, V2I, Parking Information	<ul style="list-style-type: none"> Off street parking information On street parking management and information Park & Ride information Information on AFV fuelling & charging stations 	<ul style="list-style-type: none"> C-ITS services intended to provide information regarding parking (and refuelling) to drivers
Bundle 5 Day 1.5, V2I, Traffic and other information	<ul style="list-style-type: none"> Traffic information and smart routing 	<ul style="list-style-type: none"> C-ITS services intended to provide traffic information to drivers
Bundle 6 Day 1.5, Freight specific services	<ul style="list-style-type: none"> Loading zone management Zone access control management 	<ul style="list-style-type: none"> Zone management services
Bundle 7 Day 1.5, V2X (mainly applicable to urban areas), likely to be ITS-G5	<ul style="list-style-type: none"> Vulnerable road user protection (pedestrians and cyclists) 	<ul style="list-style-type: none"> V2X service expected to be post day 1. Communication method is likely to be ITS-G5. Main benefits are likely to be seen in urban areas.
Bundle 8 Day 1.5, V2V, likely to be ITS-G5	<ul style="list-style-type: none"> Cooperative collision risk warning Motorcycle approaching indication 	<ul style="list-style-type: none"> Post day 1 V2V services that are likely to be based on ITSG5. As for Day 1 services, V2V and V2I services are in separate service bundles.
Bundle 9 Day 1.5, V2I	<ul style="list-style-type: none"> Wrong way driving 	<ul style="list-style-type: none"> Post day 1 V2I service. As for Day 1 services, V2V and V2I services are in separate service bundles.

Table 3: C-ITS service bundles for scenario building[16]